



非常天空

# Principles of Computer Systems

*Hardware/Software interface*

楼学庆

浙江大学计算机学院

<http://10.214.47.99/>

[Email:hzlou@163.com](mailto:hzlou@163.com)



玉泉校区曹光彪东楼507室



08:51:14

浙江大学计算机学院

# 联系方式

- 网站:

- <http://10.214.47.99>

- 邮箱:

- [hzlou@163.com](mailto:hzlou@163.com) (不收作业)

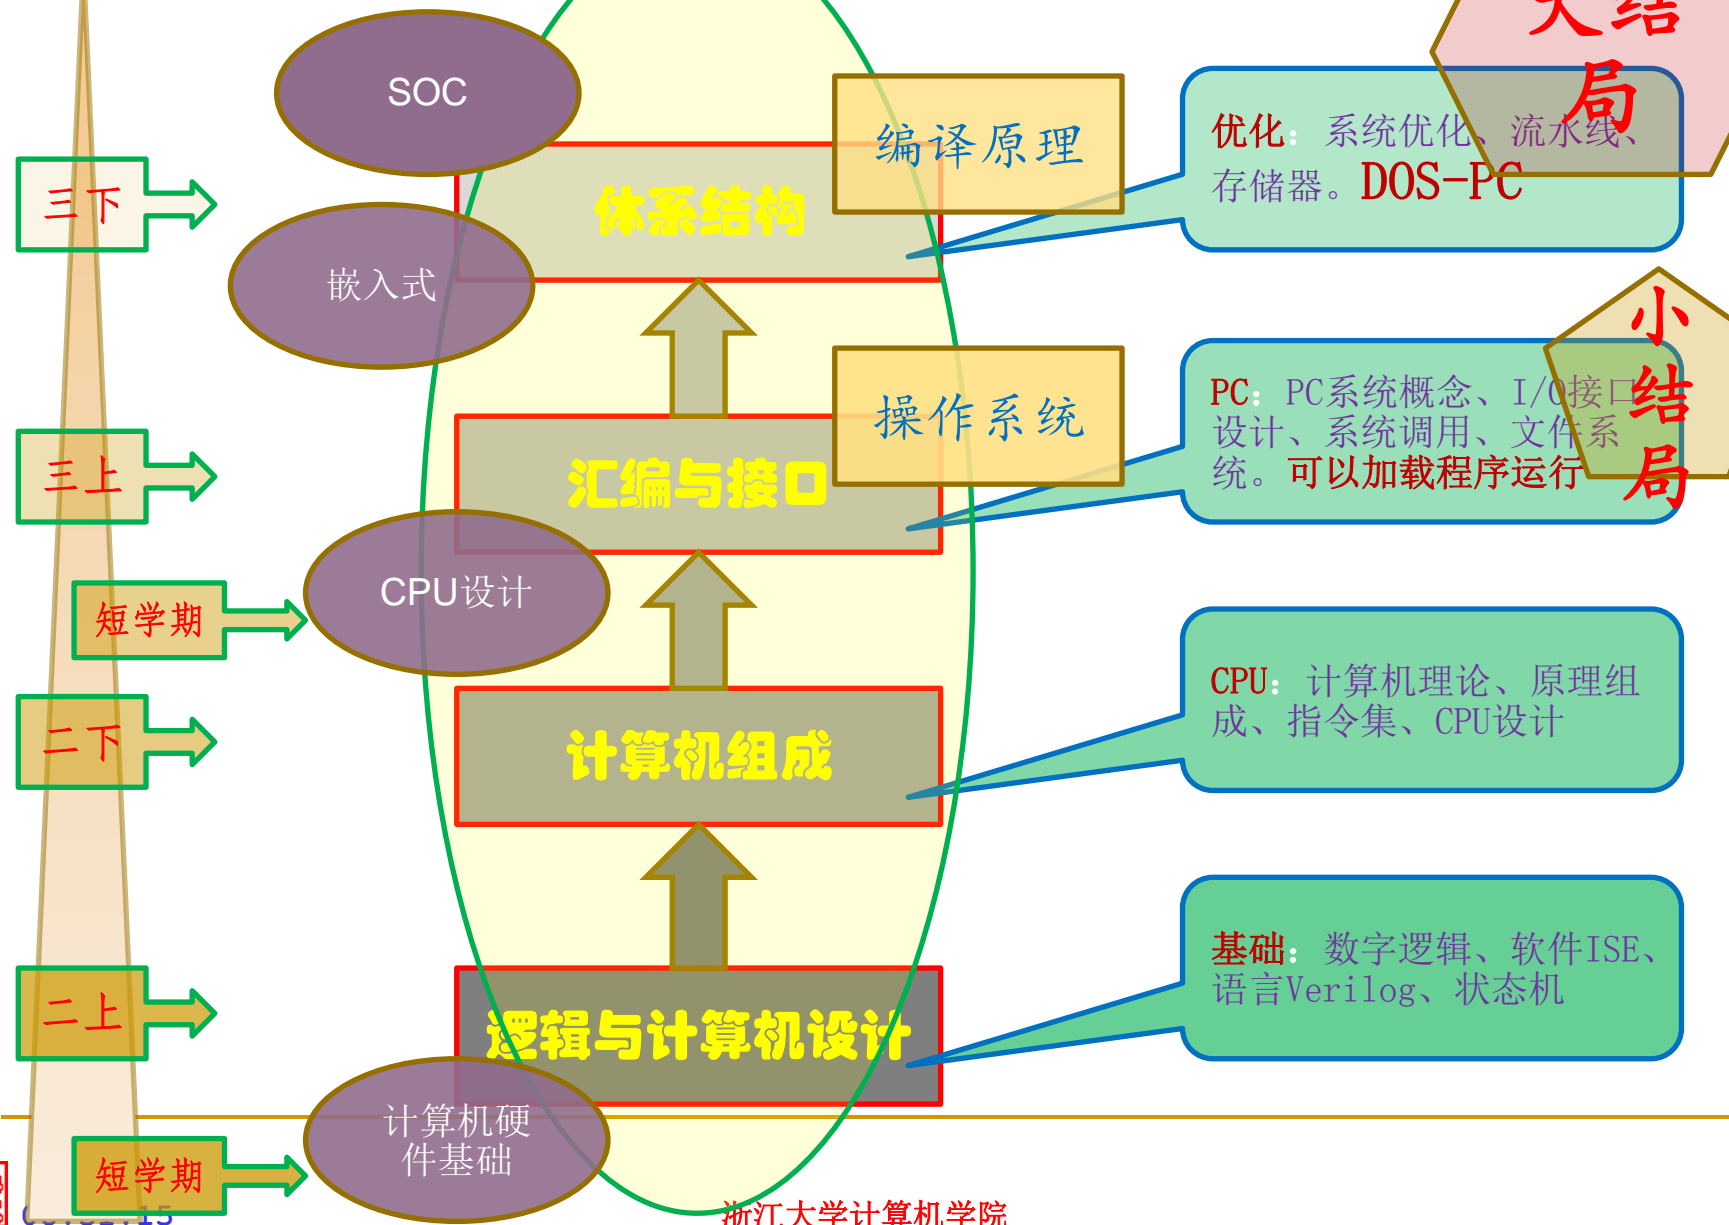




# 计算机系统能力培养系列课程体系

时间轴

主干线



# Course outline


- Name:

## Computer Systems

- Students:

Undergraduate students  
some others department


- Score : 4.5
- Hours/week: 3.5-2
- Total: 88 hours



# Textbook



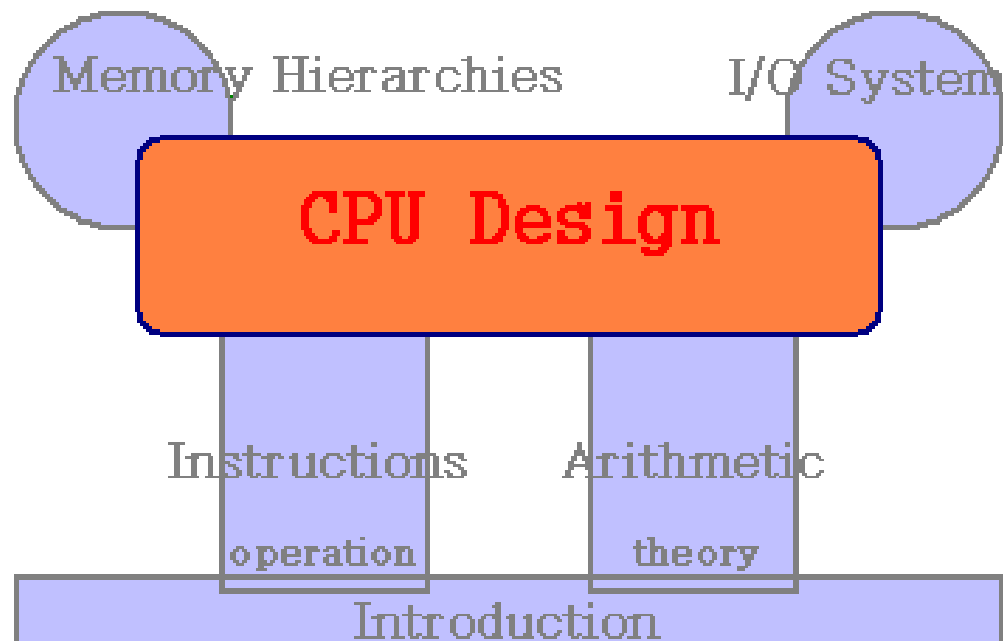
2nd  
Edition



3rd  
Edition

# CPU design

- Topics: Datapath & Control









# Processor


- We're ready to look at an implementation of the MIPS
- Simplified to contain only:
  - memory-reference instructions: `lw, sw`
  - arithmetic-logical instructions: `add, sub, and, or, slt`
  - control flow instructions: `beq, j`
- Generic Implementation:
  - use the program counter (PC) to supply instruction address
  - get the instruction from memory
  - read registers
  - use the instruction to decide exactly what to do






# State Elements

- Unclocked vs. Clocked
- Clocks used in synchronous logic
  - when should an element that contains state be updated?



# Our Implementation

- An edge triggered methodology
- Typical execution:
  - read contents of some state elements,
  - send values through some combinational logic
  - write results to one or more state elements


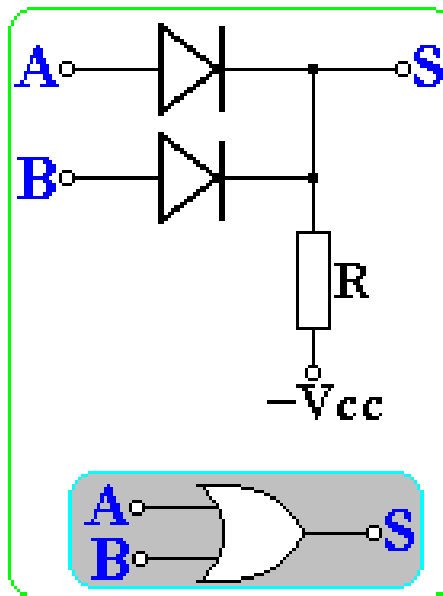
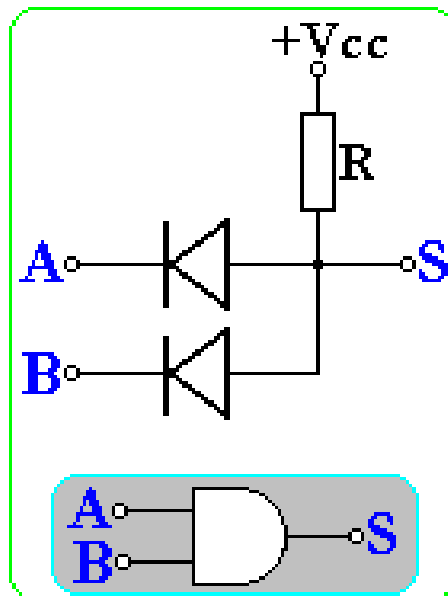


# Basic elements

## Basic logic gate.

- AND:  $S=A \cdot B$
- OR :  $S=A+B$
- NOT:  $S=\sim A$

A	B	$A \cdot B$	$A+B$	$\sim A$
0	0	0	0	1
0	1	0	1	1
1	0	0	1	0
1	1	1	1	0






# MUX

- 多路开关：有若干个输入，由控制端决定那一路输入将输出。

□ 经常具有多路并行

C	S
0	A
1	B



# Adder

## ■ Half adder:


A	0	0	1	1
B	0	1	0	1
S	0	1	1	0

## ■ Full adder: $S=A+B$

□ Input: A, B, C<sub>0</sub>. Output: Sum,




□  $S = A \oplus B \oplus C$

A	0	0	1	1	0	1	1
B	0	1	0	1	0	1	1
C	0	0	0	0	1	1	1
S	0	1	1	0	1	0	1
C <sub>+1</sub>	0	0	0	1	0	1	1




# 串行进位加法器

- 4位串行进位加法器：  $S=A+B$ 
  - 各位串行进行。
  - 同时根据结果产生各种标记状态。



# Adder

## ■ 32-bit Adder







# ALU

## ■ 算术逻辑运算器ALU：即运算

□ 5 Operations

□ "Set on less than":  
if  $A < B$  then Result=1;  
else Result=0.

Operation	Function
000	And
001	Or
010	Add
110	Sub
111	Slt


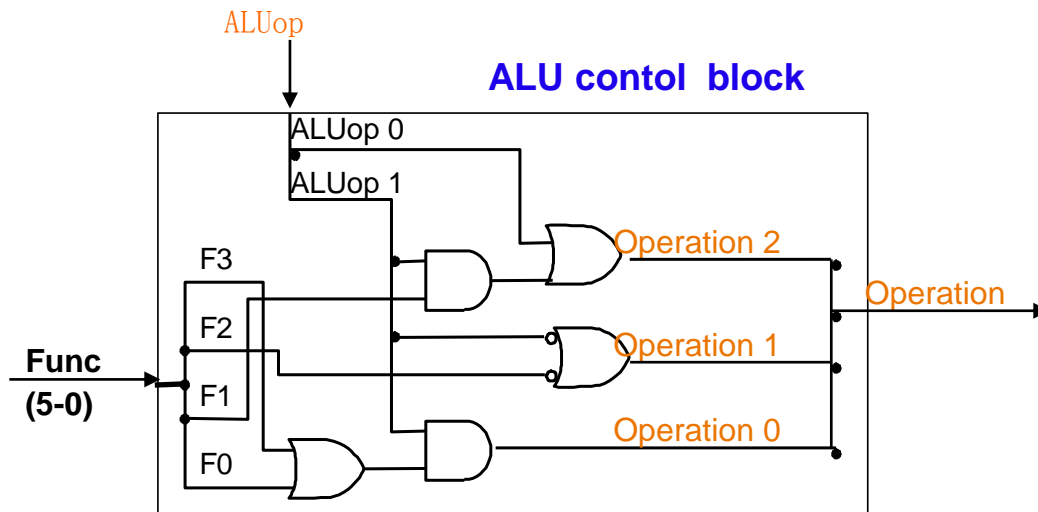


# ALU Control

## ■ ALU Control:

- **ALU由ALUcontrol控制:**  
**ALUcontrol由ALUop和指令的低6位 (5-0) 联合产生ALU控制码。**
- 这样的好处在于分级控制, 对于用的最多的加、减操作, 系统只需给出两位的**ALUop**即可。

ALUop	Operation	Function
10	000	And
10	001	Or
00	010	Add
01	110	Sub
10	111	Slt



# ALU Control


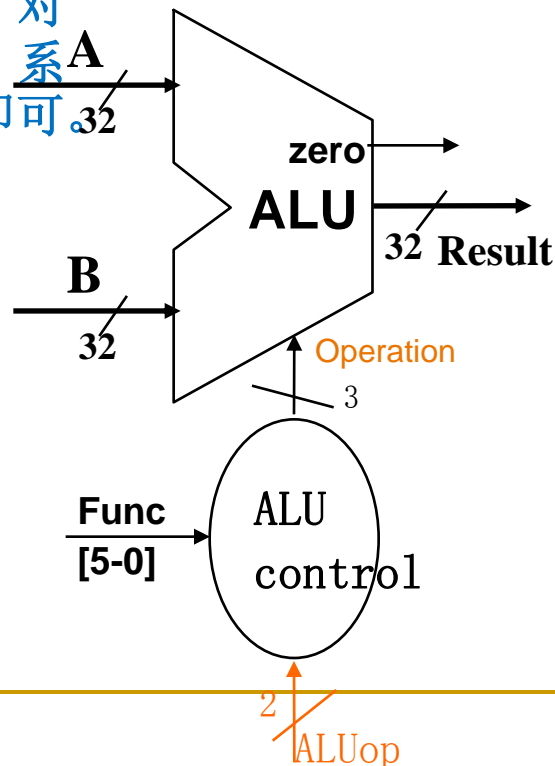
## ■ ALUop与Func联合对ALU的控制

ALUOp		Func field						Operation	ALU Function
ALUOp1	ALUOp0	F5	F4	F3	F2	F1	F0		
0	0	X	X	X	X	X	X	010	Add
1	1	X	X	X	X	X	X	110	Sub
1	X	X	X	0	0	0	0	010	Add
1	X	X	X	0	0	1	0	110	Sub
1	X	X	X	0	1	0	0	000	And
1	X	X	X	0	1	0	1	001	Or
1	X	X	X	1	0	1	0	111	SetonLT

# ALU Control

## ■ ALU控制:

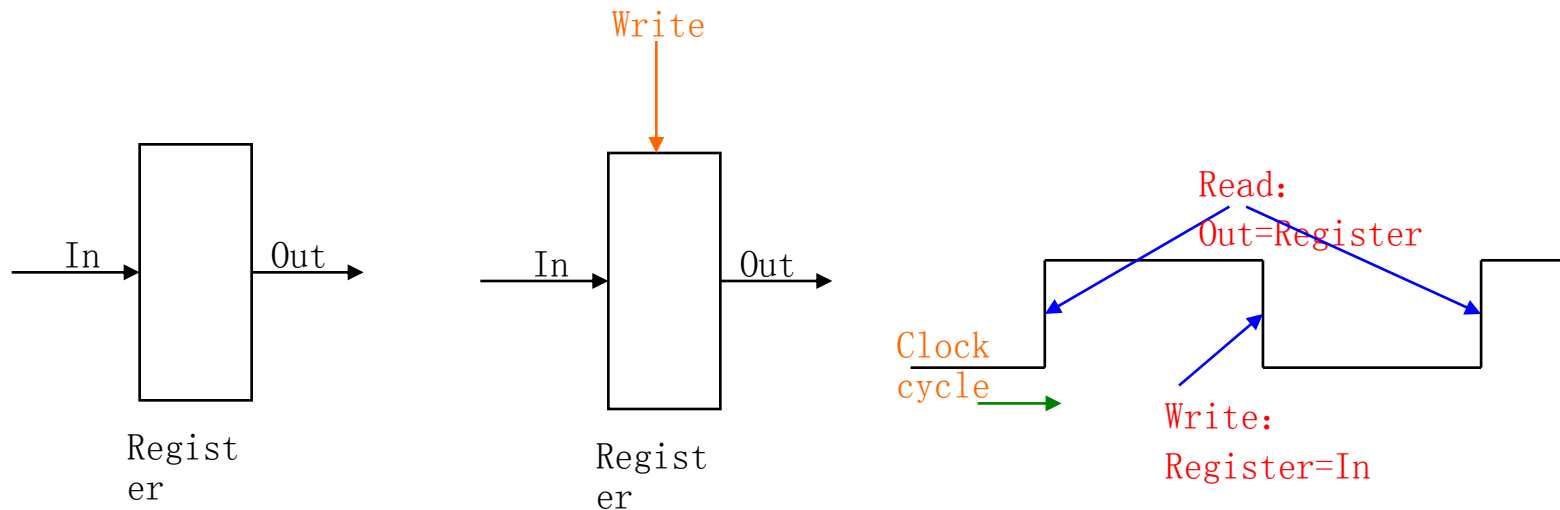
- ALU由ALUcontrol控制:  
ALUcontrol由ALUop和指令的低6位(5-0)联合产生ALU控制码。
- 这样的好处在于分级控制, 对于用的最多的加、减操作, 系统只需给出两位的ALUop即可



# REGISTER

## ■ Register


- State element.
- Can be controlled by **Write** signal.



# Memory

## ■ 存储器:

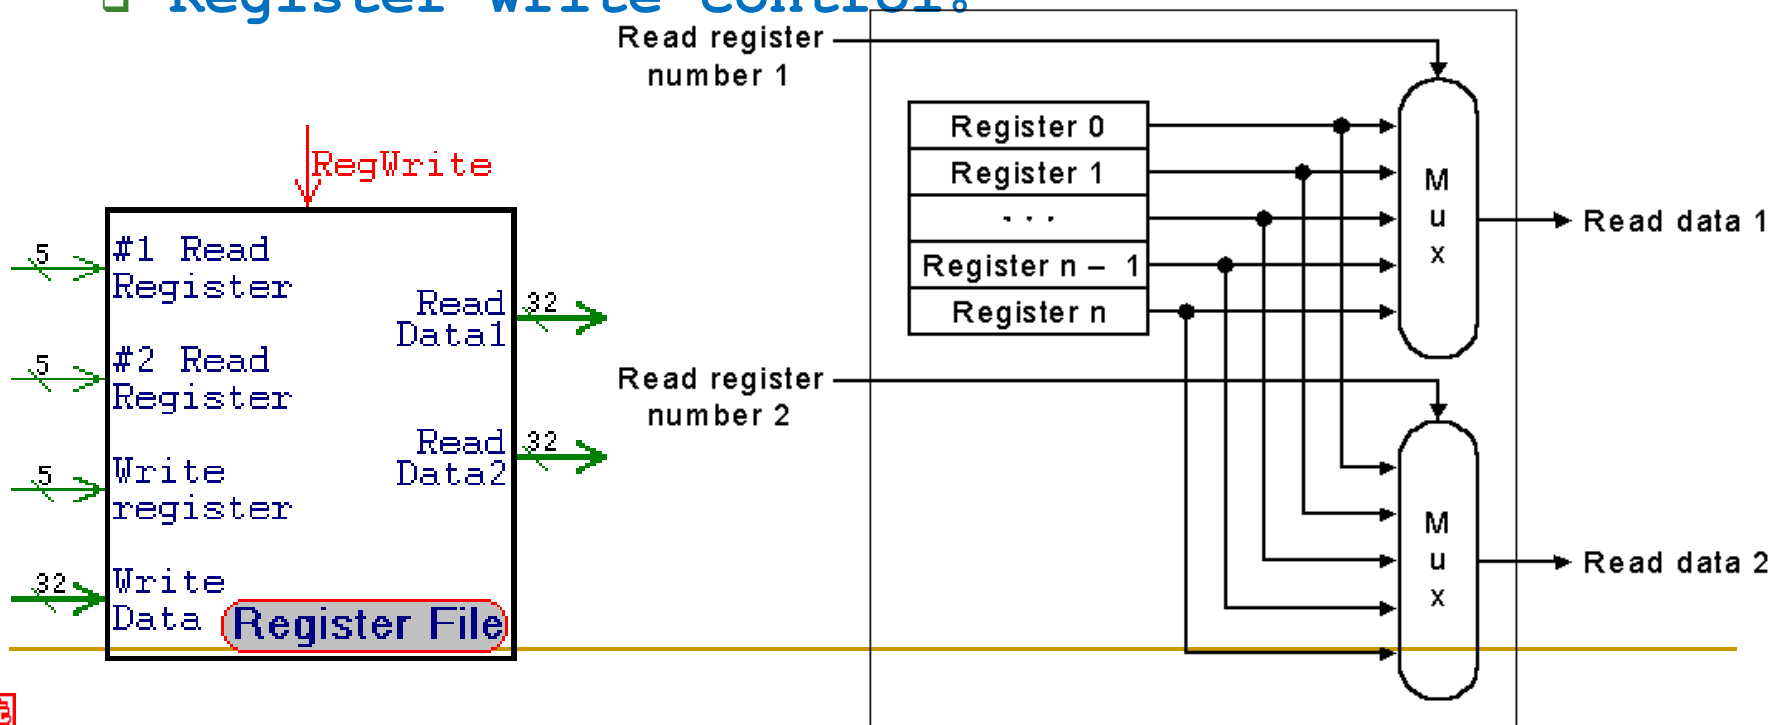
- 可分为指令存储器与数据存储器;
- 指令存储器设为只读; 输入指令地址, 输出指令。
- 数据存储器可以读写, 由MemRead和MemWrite控制。按地址读出数据输入, 或将写数据写入地址所指存储器单元。



# Register File

## ■ Register File:

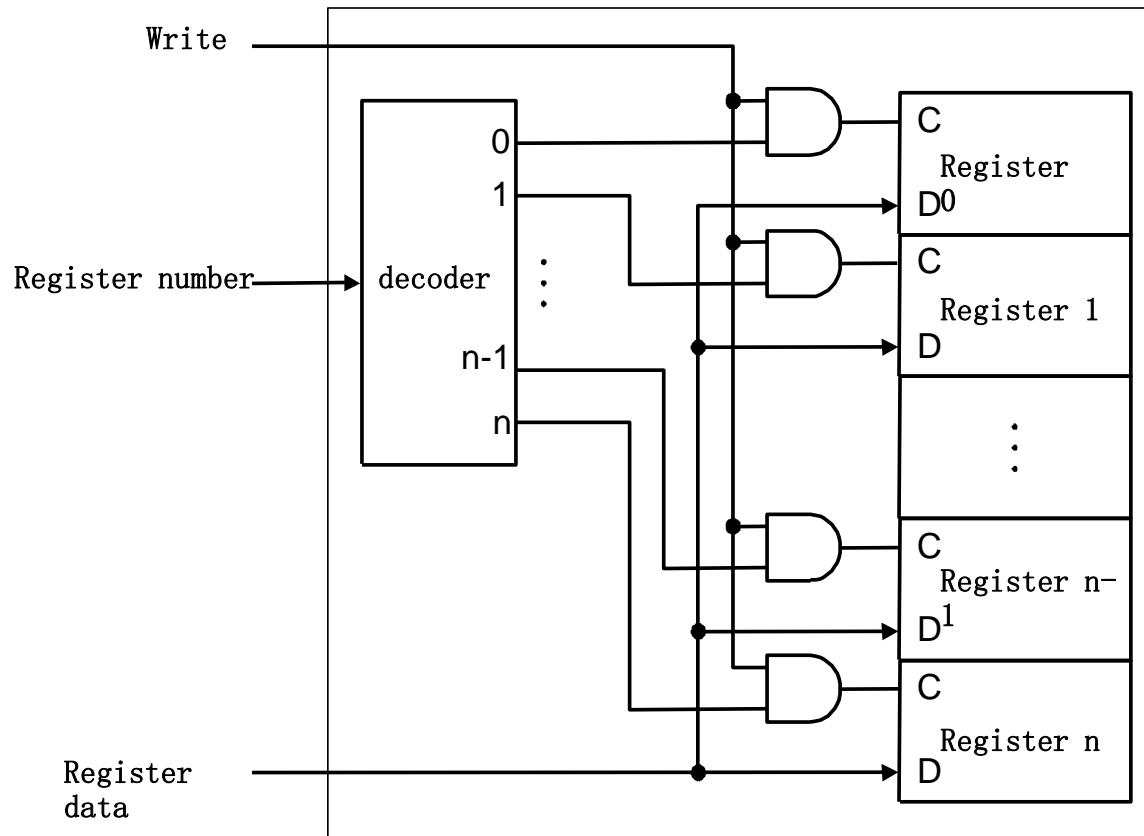
- 32 32-bit Registers;
- Input: 2 32-bit;
- Output: 32-bit data, 32-bit register number;
- Register write control.





# Register File

- 写寄存器: we still use the real clock to determine when to write



# 32个Register的功用

- 寄存器组有32个32位寄存器，其功用分配如下表：

Name	Register number	Usage
\$zero	0	the constant value 0
\$at	1	reserved for the assembler
\$v0-\$v1	2-3	values for results and expression evaluation
\$a0-\$a3	4-7	arguments
\$t0-\$t7	8-15	temporaries
\$s0-\$s7	16-23	saved
\$t8-\$t9	24-25	more temporaries
\$k0-\$k1	26-27	reserved for the operating system
\$gp	28	global pointer
\$sp	29	stack pointer
\$fp	30	frame pointer
\$ra	31	return address

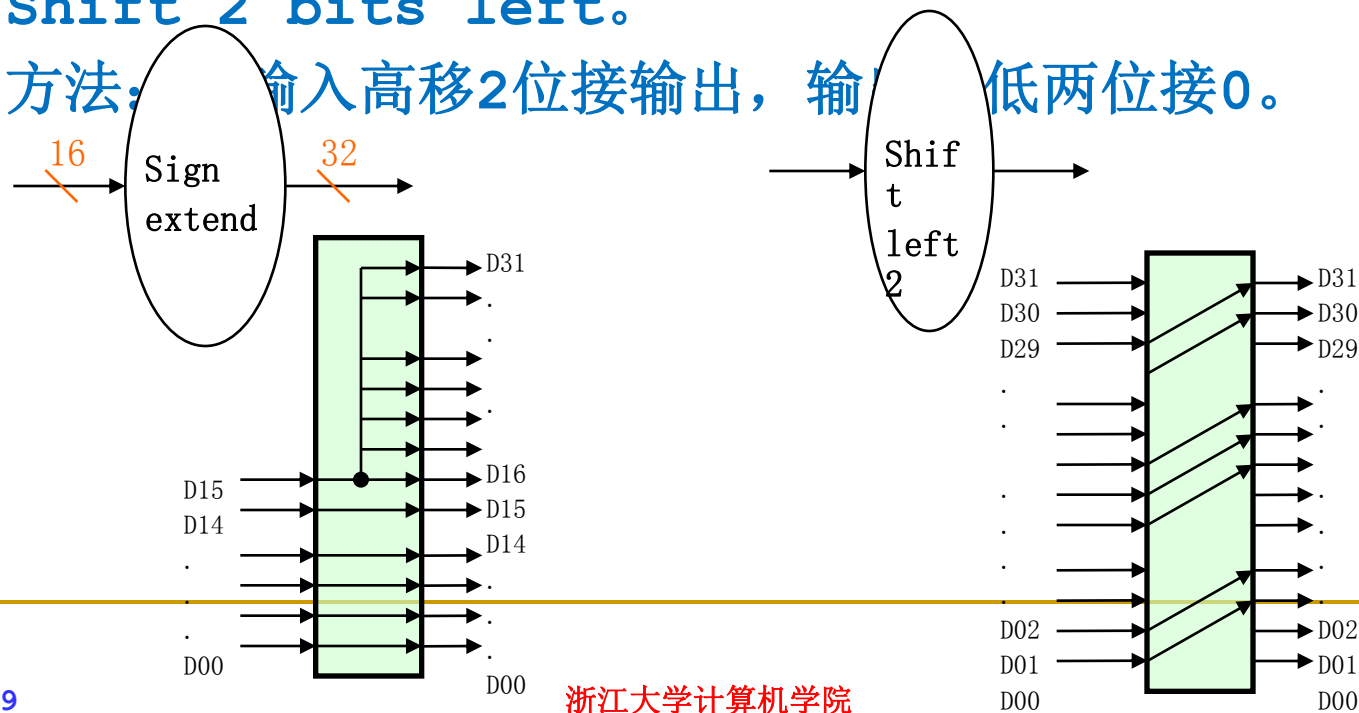
# The other elements

## ■ Sign Extend:

- 将16位的补码表示的有符号立即数，扩展为32位。
- 方法：只需重复符号位即可。

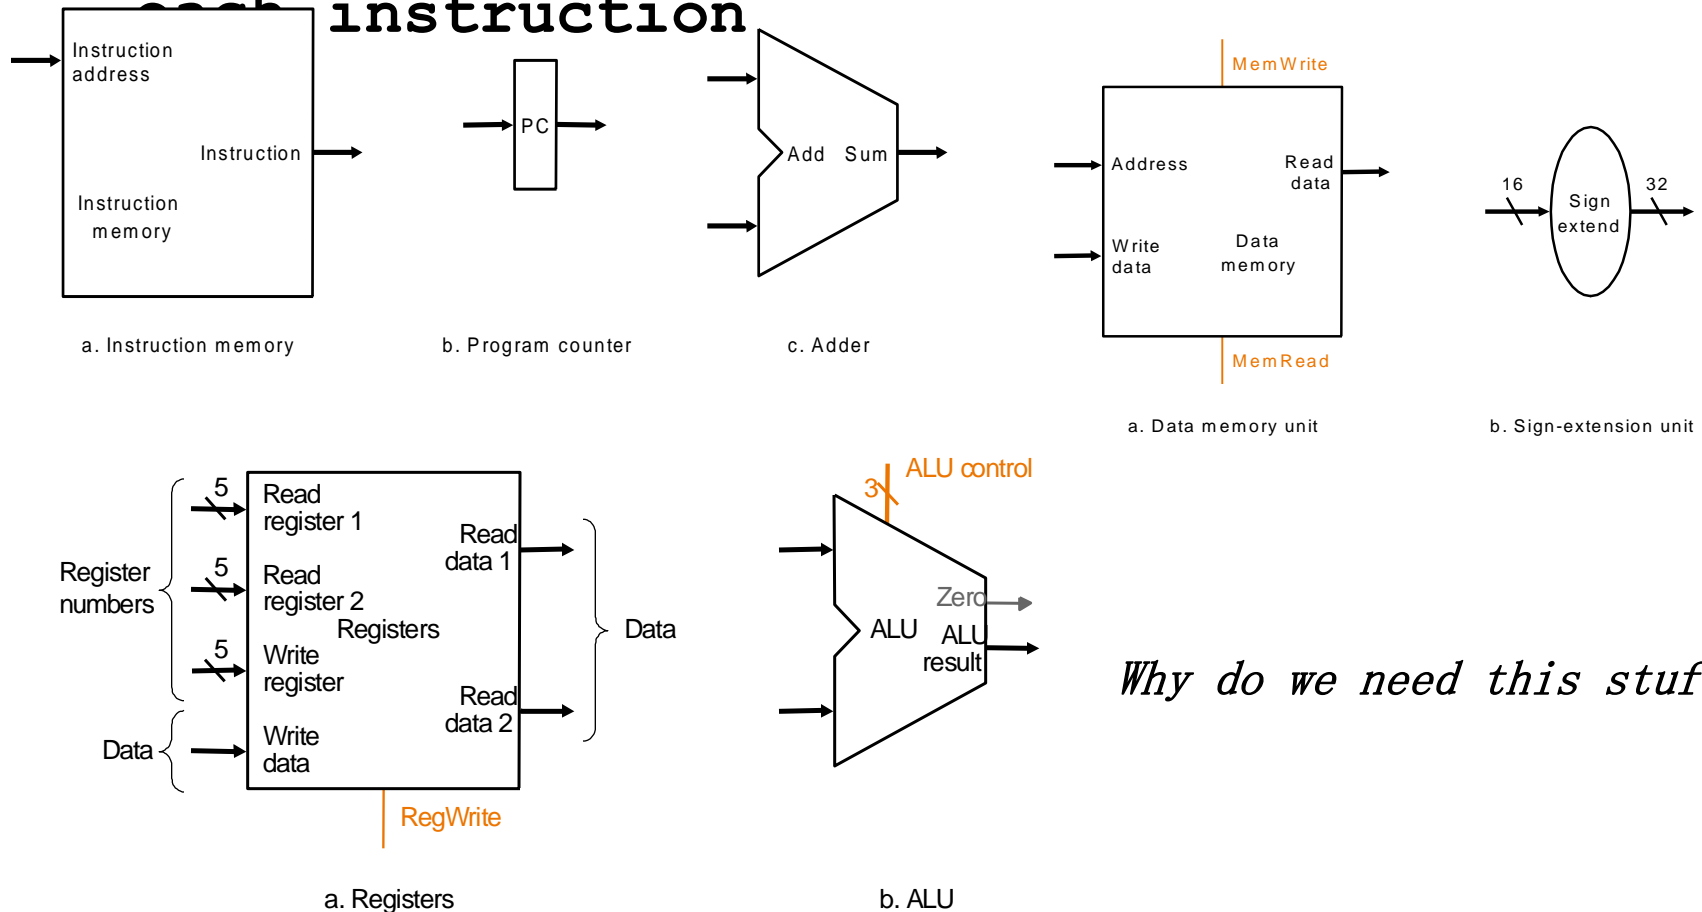
## ■ Shift:

- Shift 2 bits left.
- 方法：输入高移2位接输出，输入低两位接0。



# Simple Implementation

- Include the functional units we need for each instruction

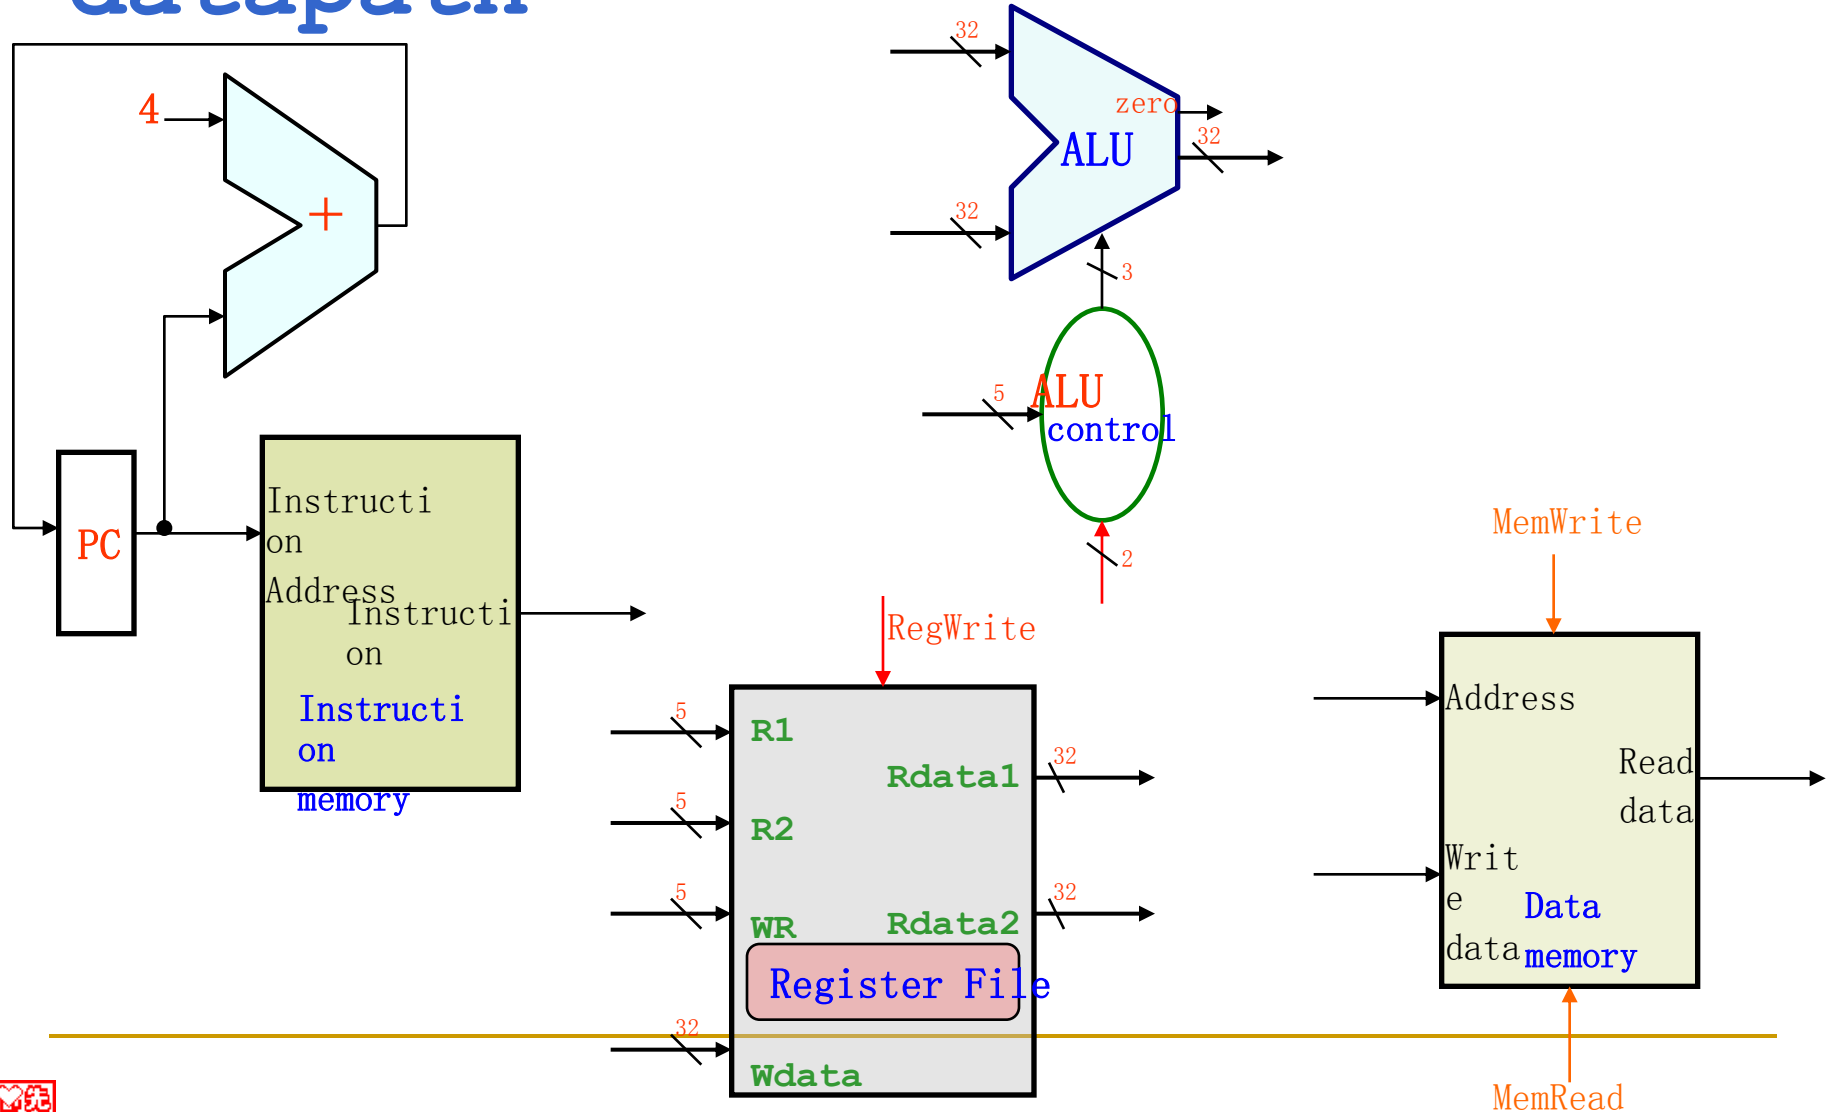


*Why do we need this stuff?*

# Section 2


## datapath

# Building a



# More Details

## ■ Abstract / Simplified View:

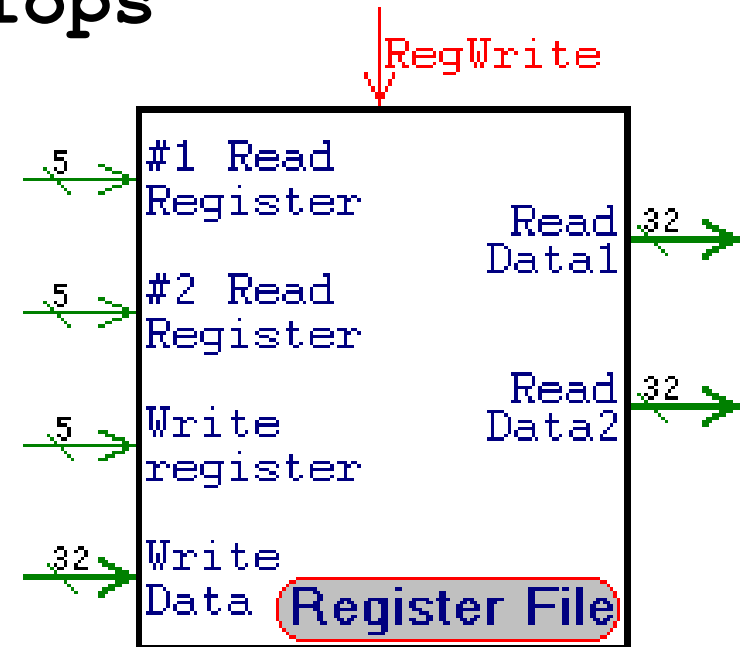
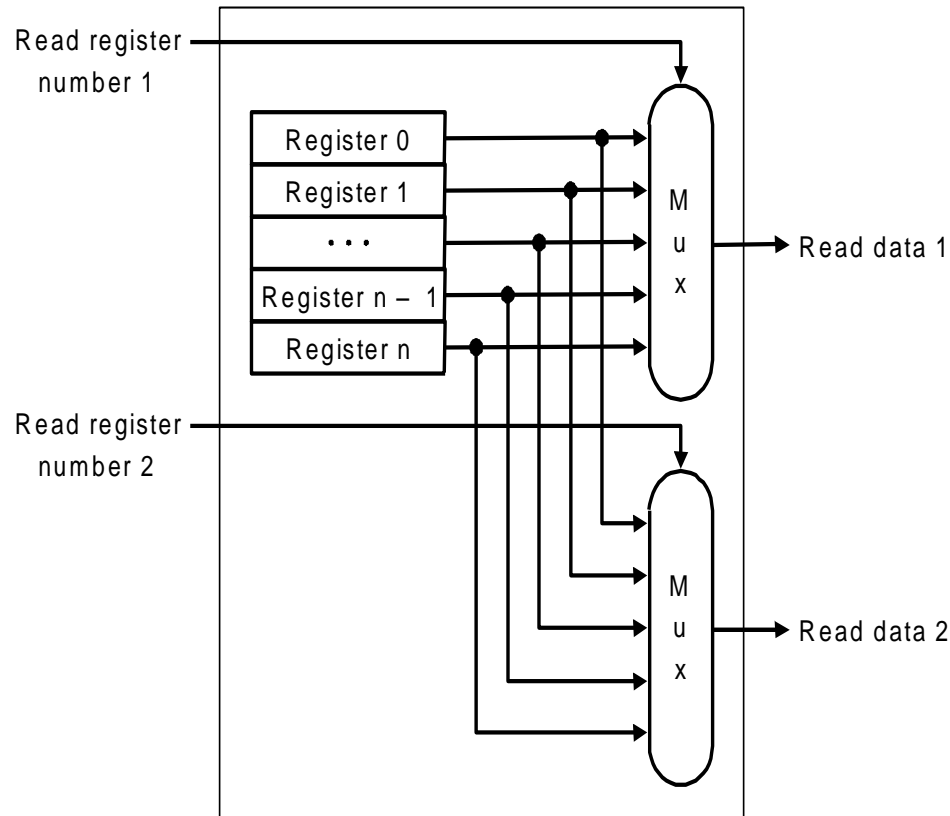


Two types of functional units:

- elements that operate on data values (combinational)
- elements that contain state (sequential)

# Register File

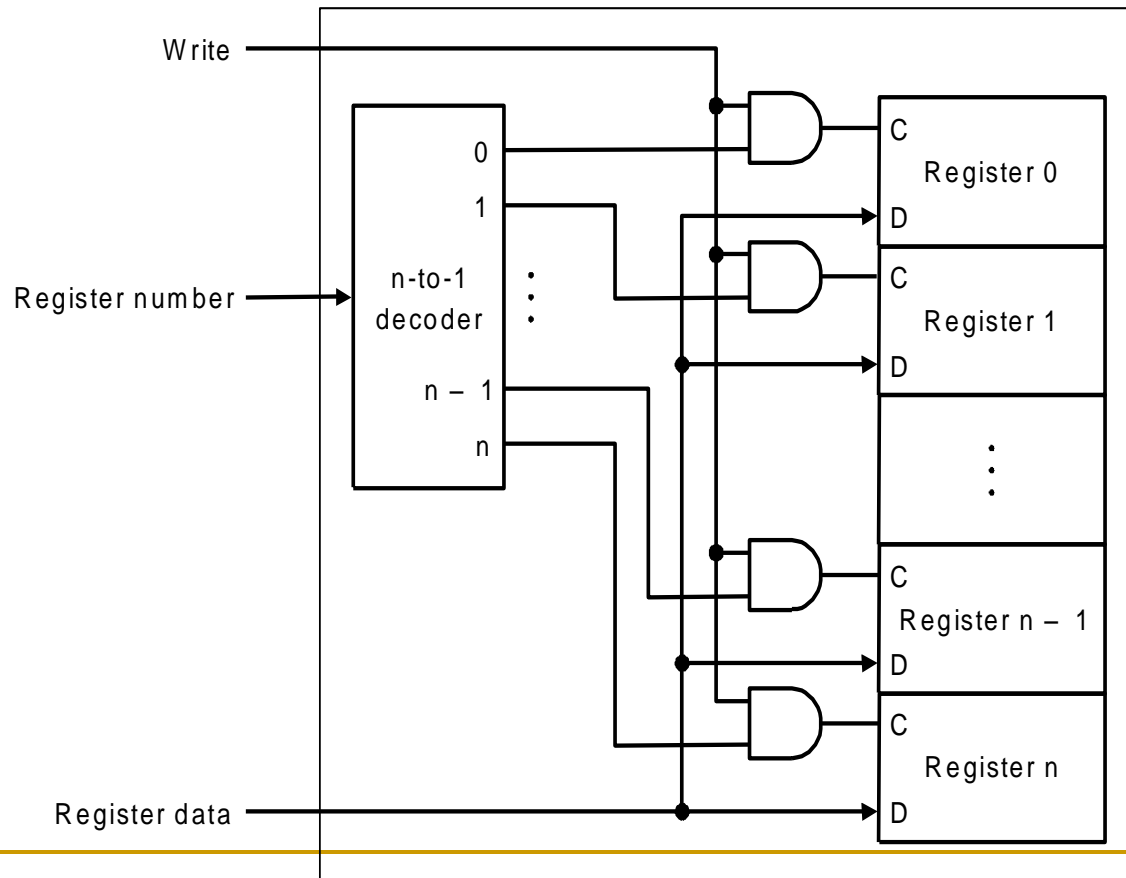
## ■ Built using D flip-flops



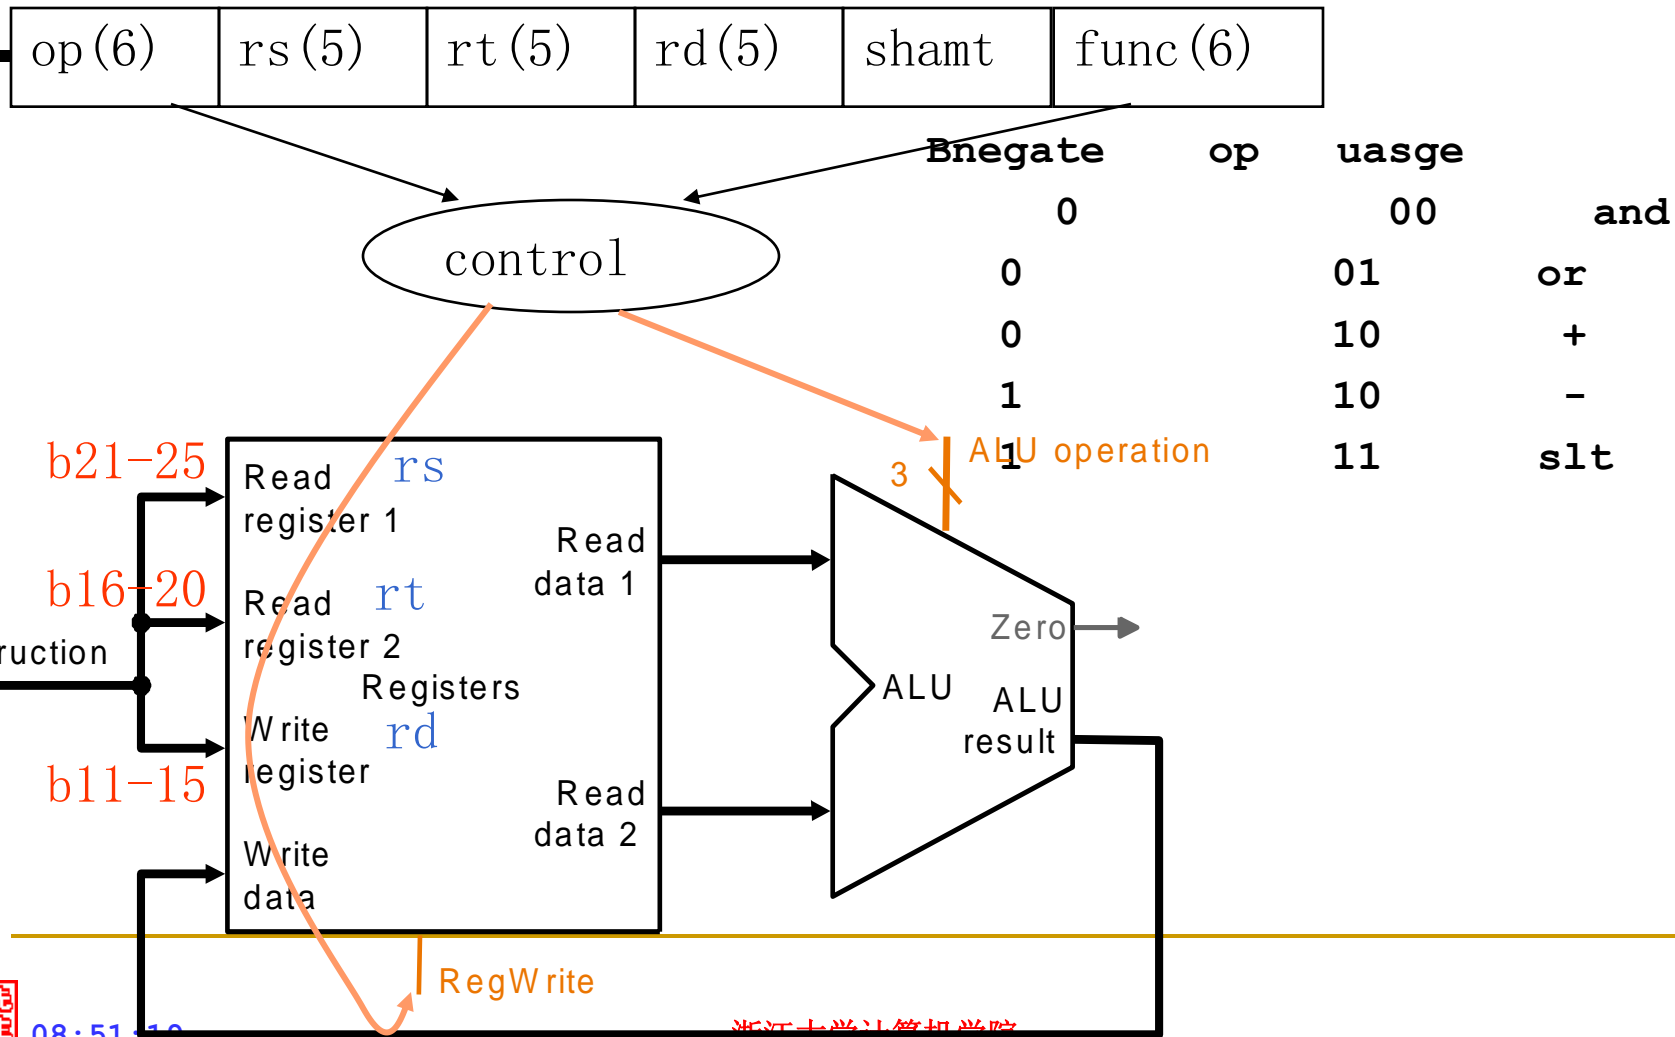


# Register File

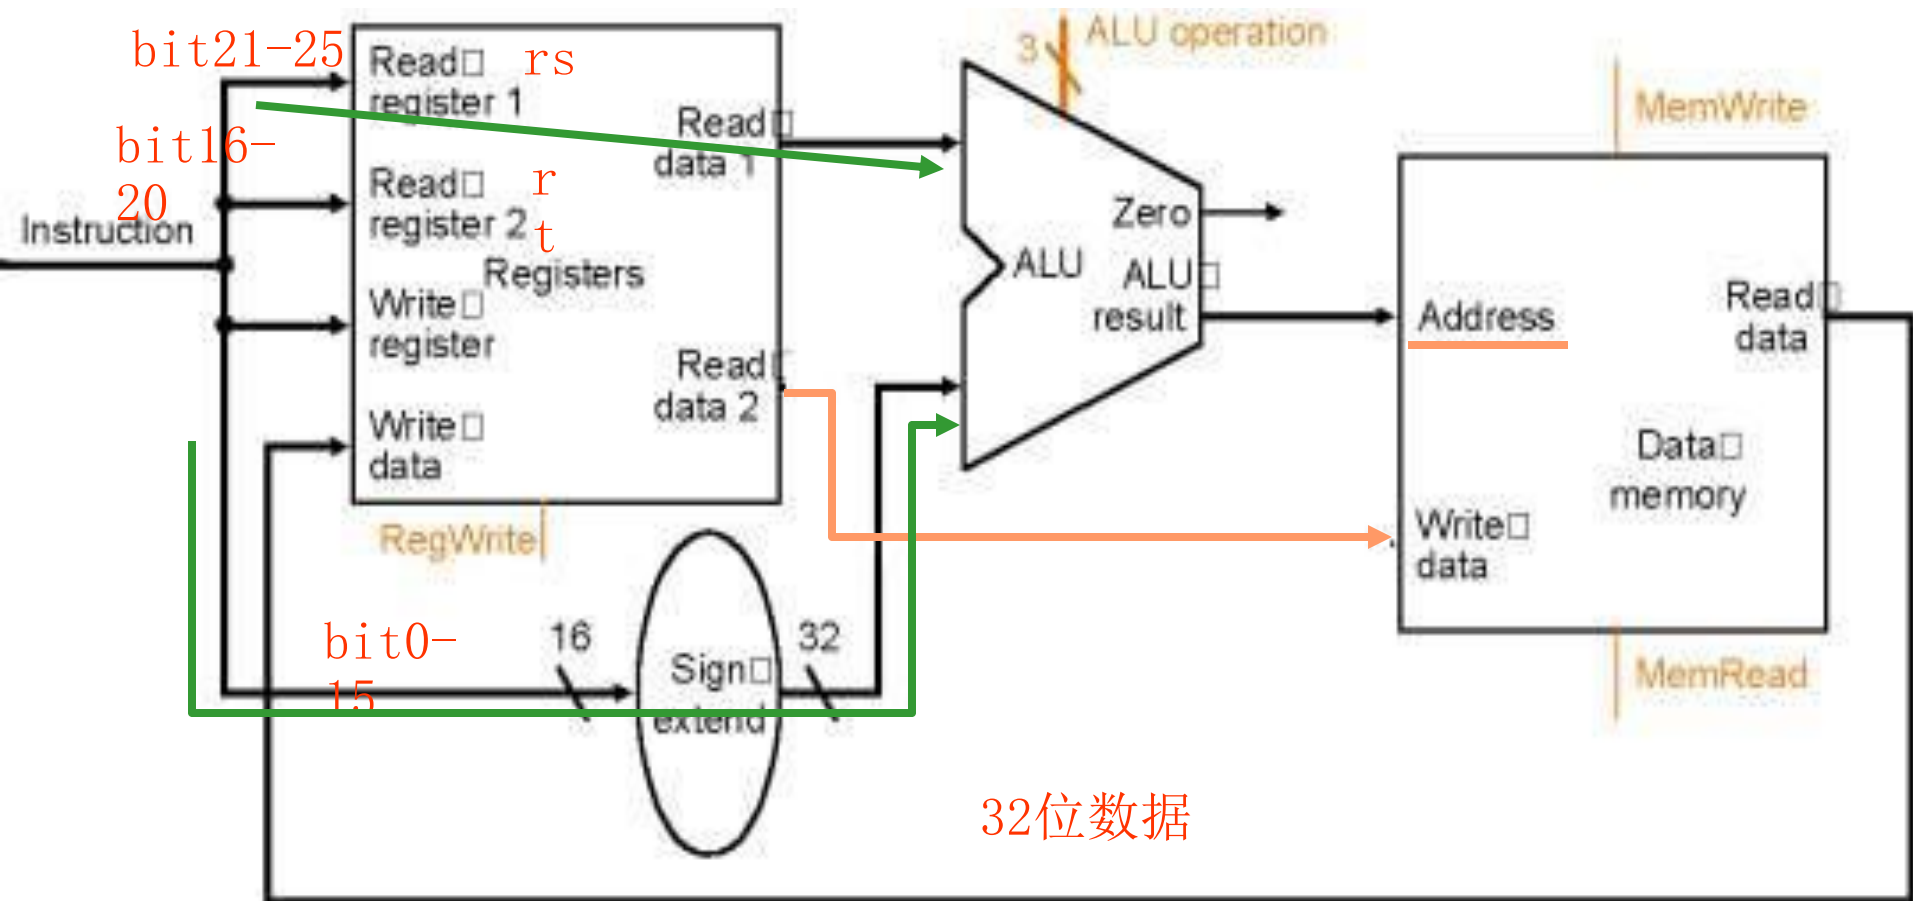
- Note: we still use the real clock to determine when to write



# Implemente the R type instruction format:



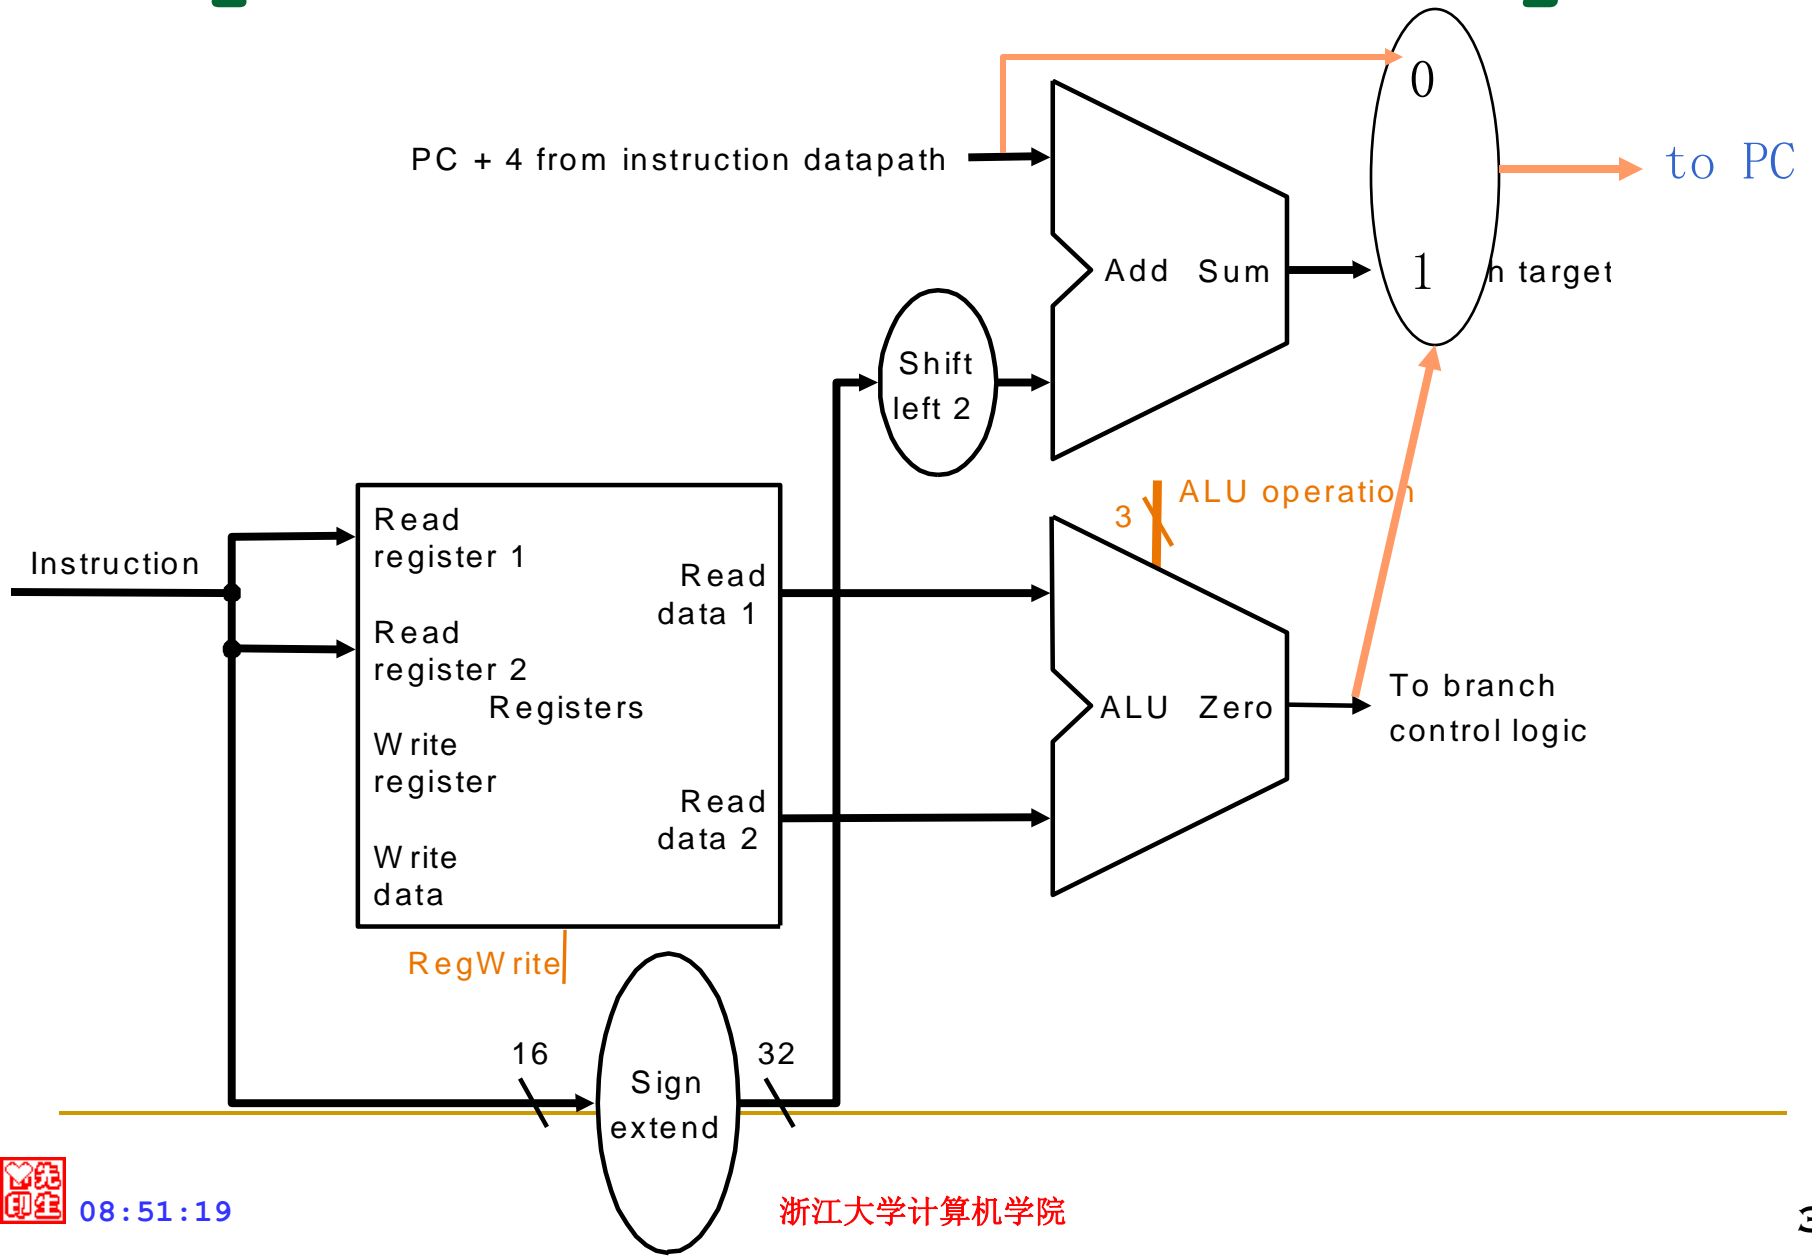
# Implemente the I type



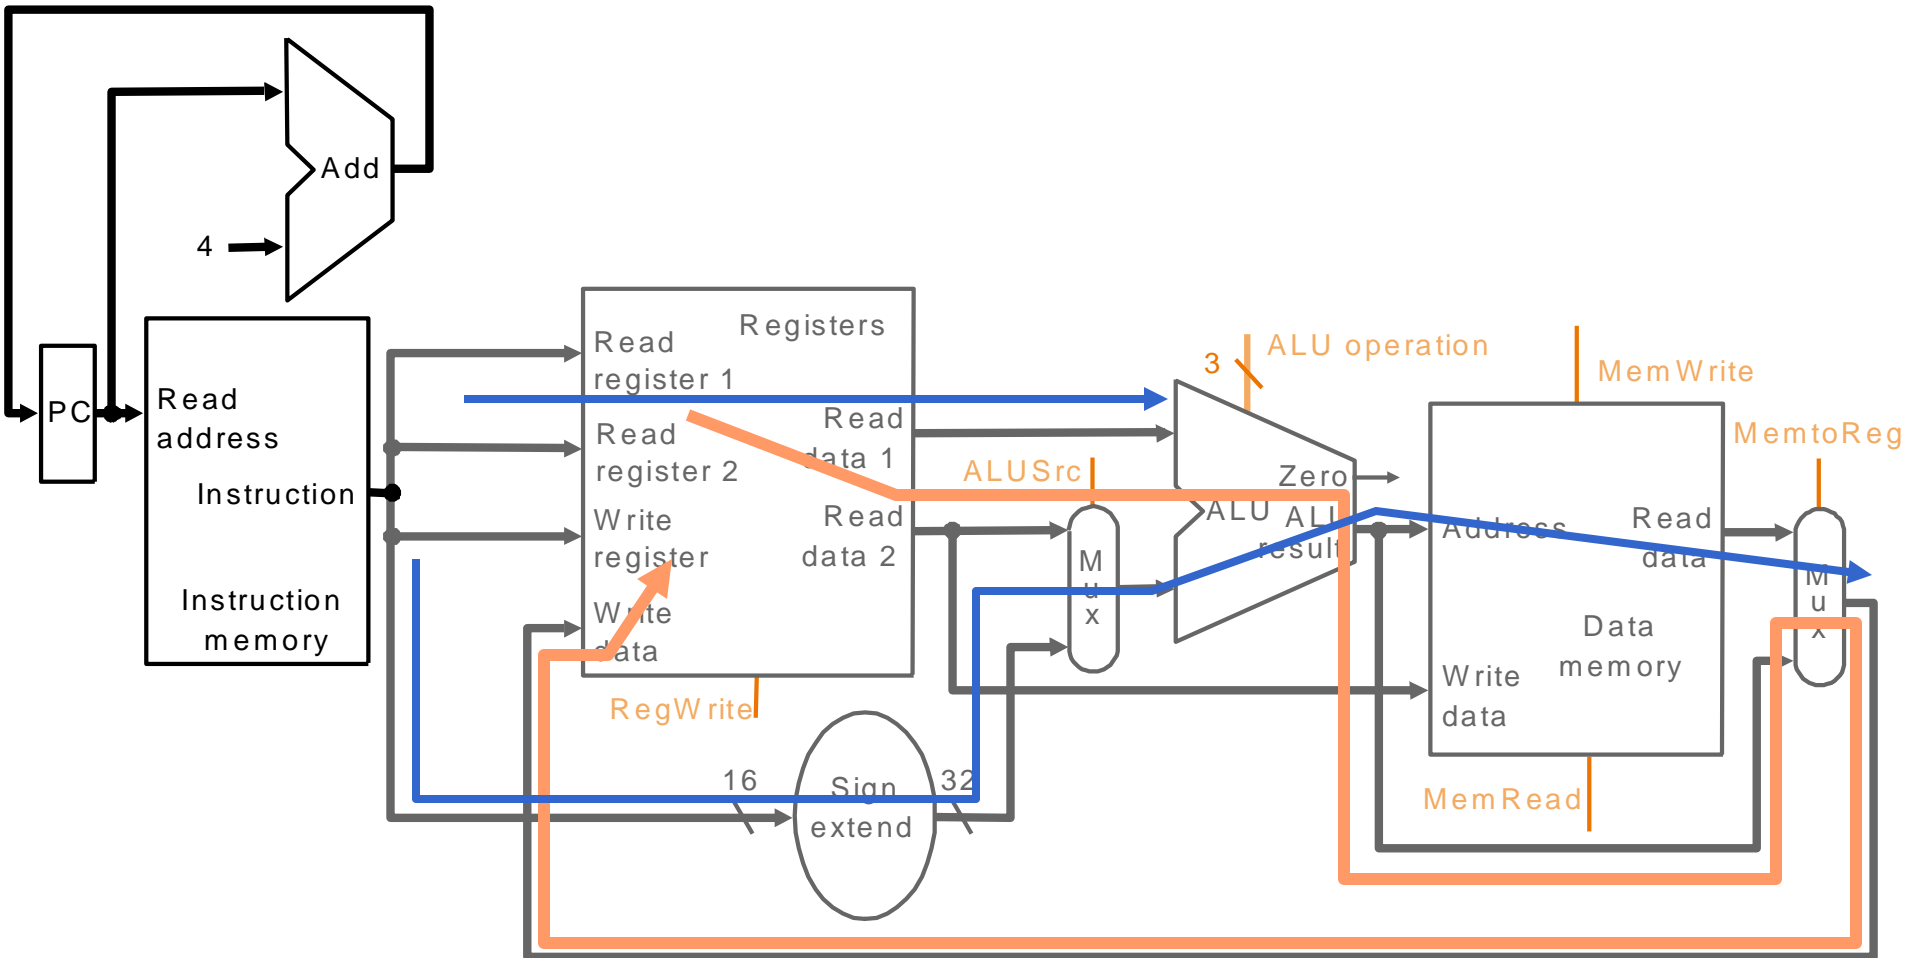
sw \$t0, 200(\$s2)

若\$s2=1000, 则将\$t0存入1200为地址的内存单元的一个字


# Implementation of *beq*



# Combine R-type and I

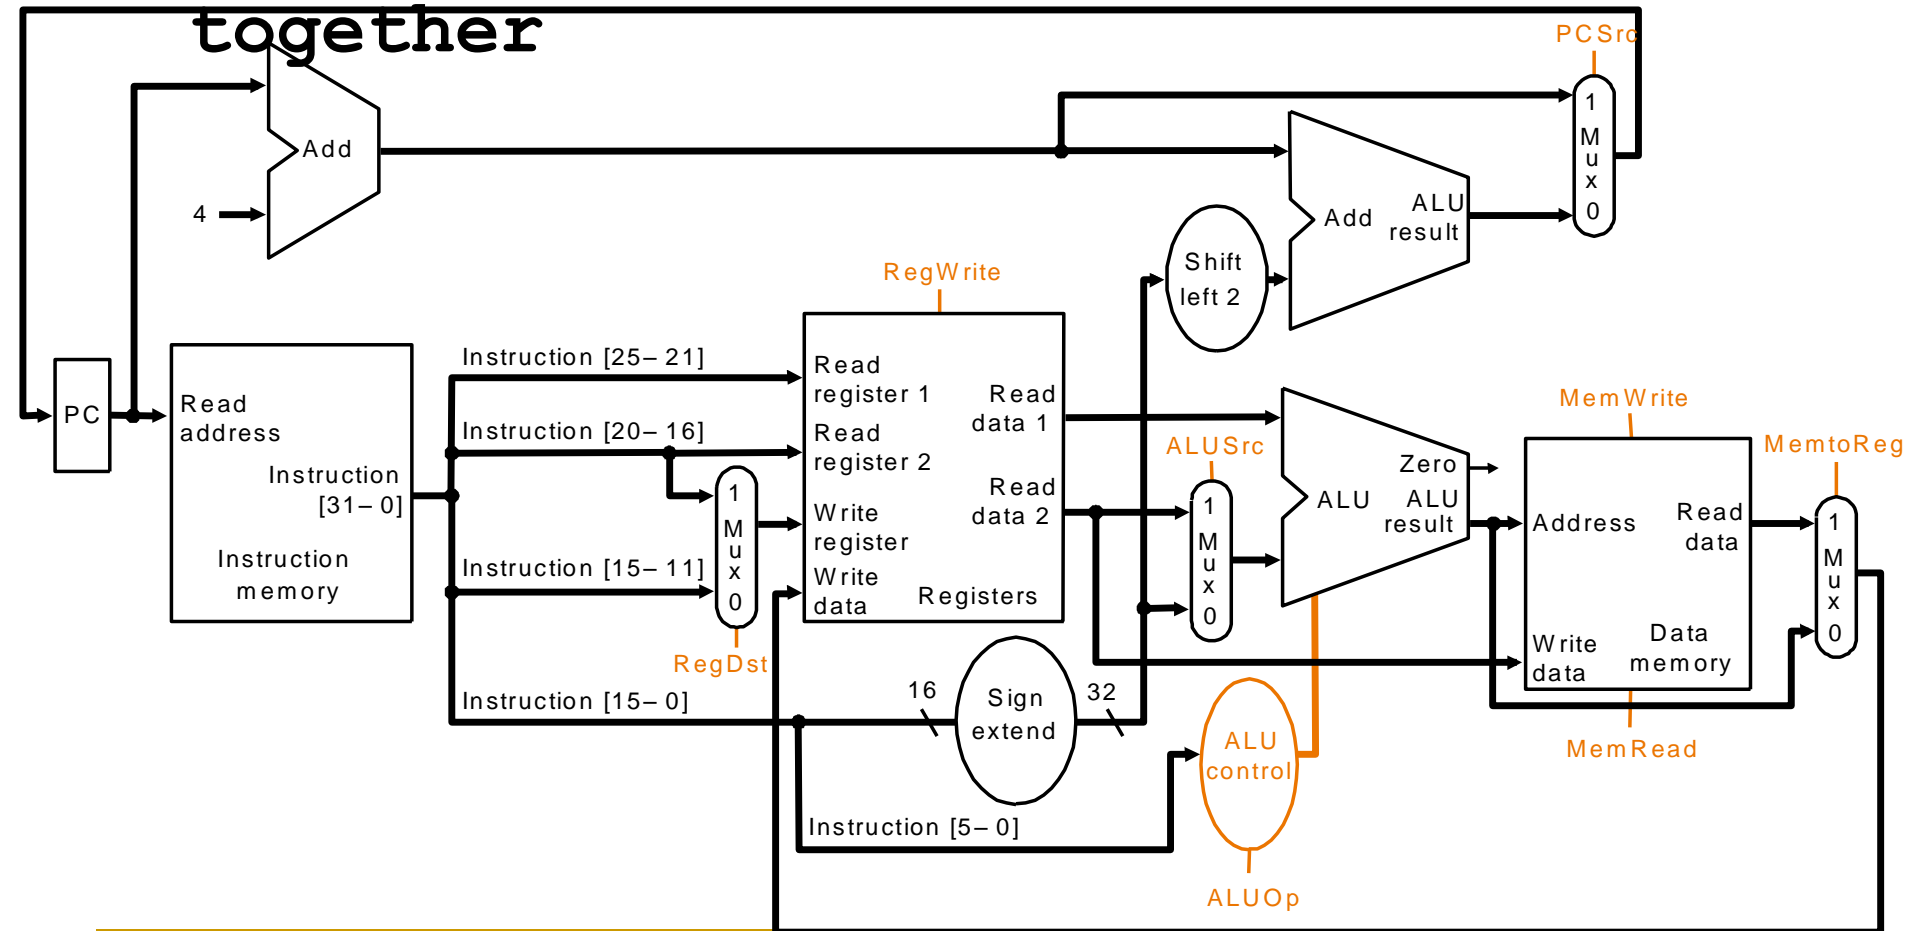


# Implementation of *beq*



# Building the Datapath


- Use multiplexors to stitch them together



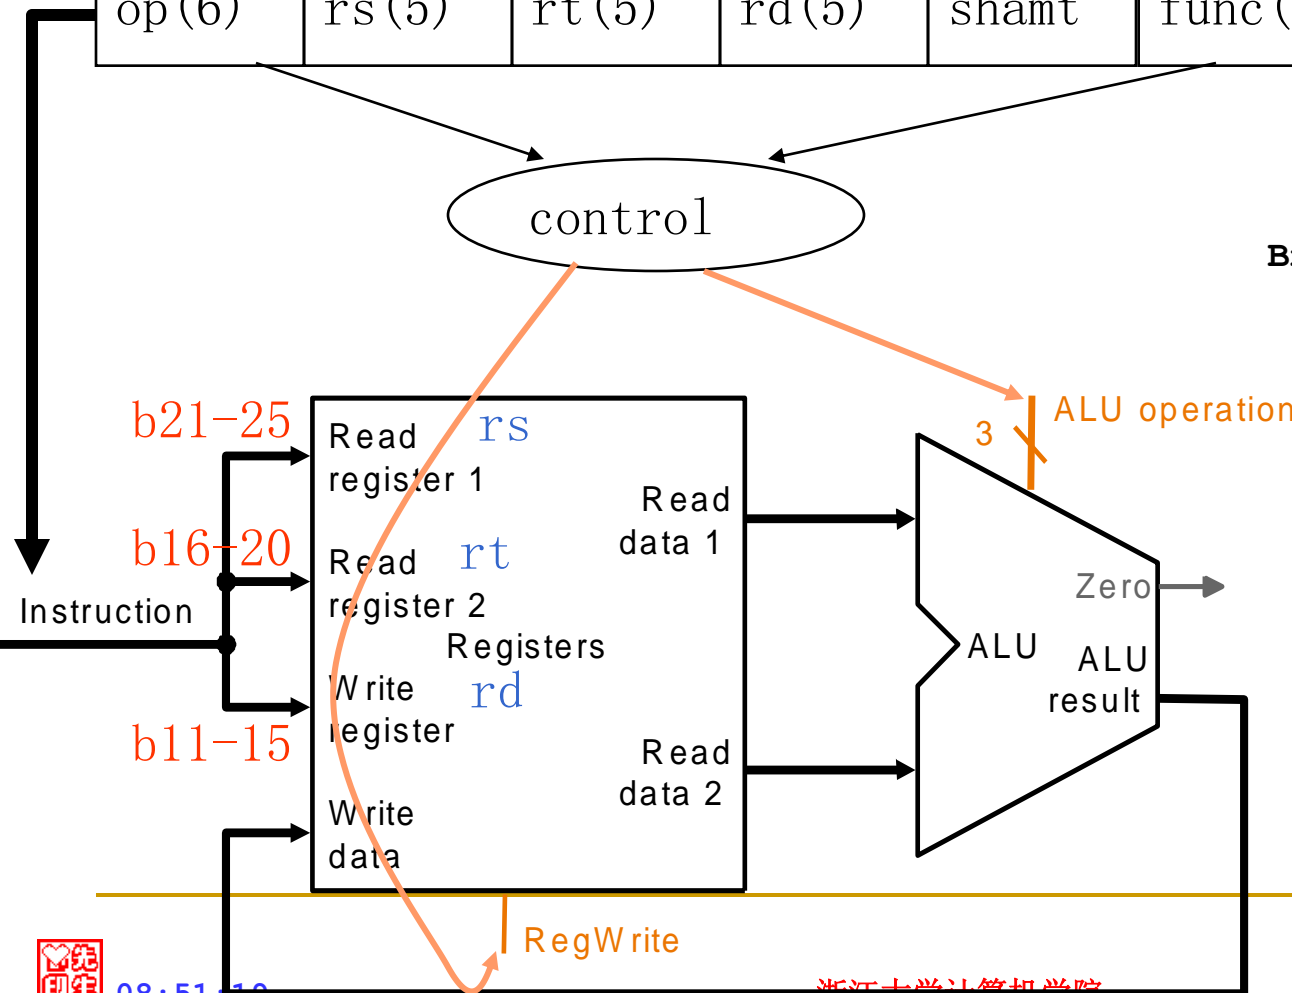


# The ALU control

instruction format:



Bnegate	op	uasge
0		00
and		
0	01	or
0	10	+
1	10	-
1	11	
slt		



# The ALU control

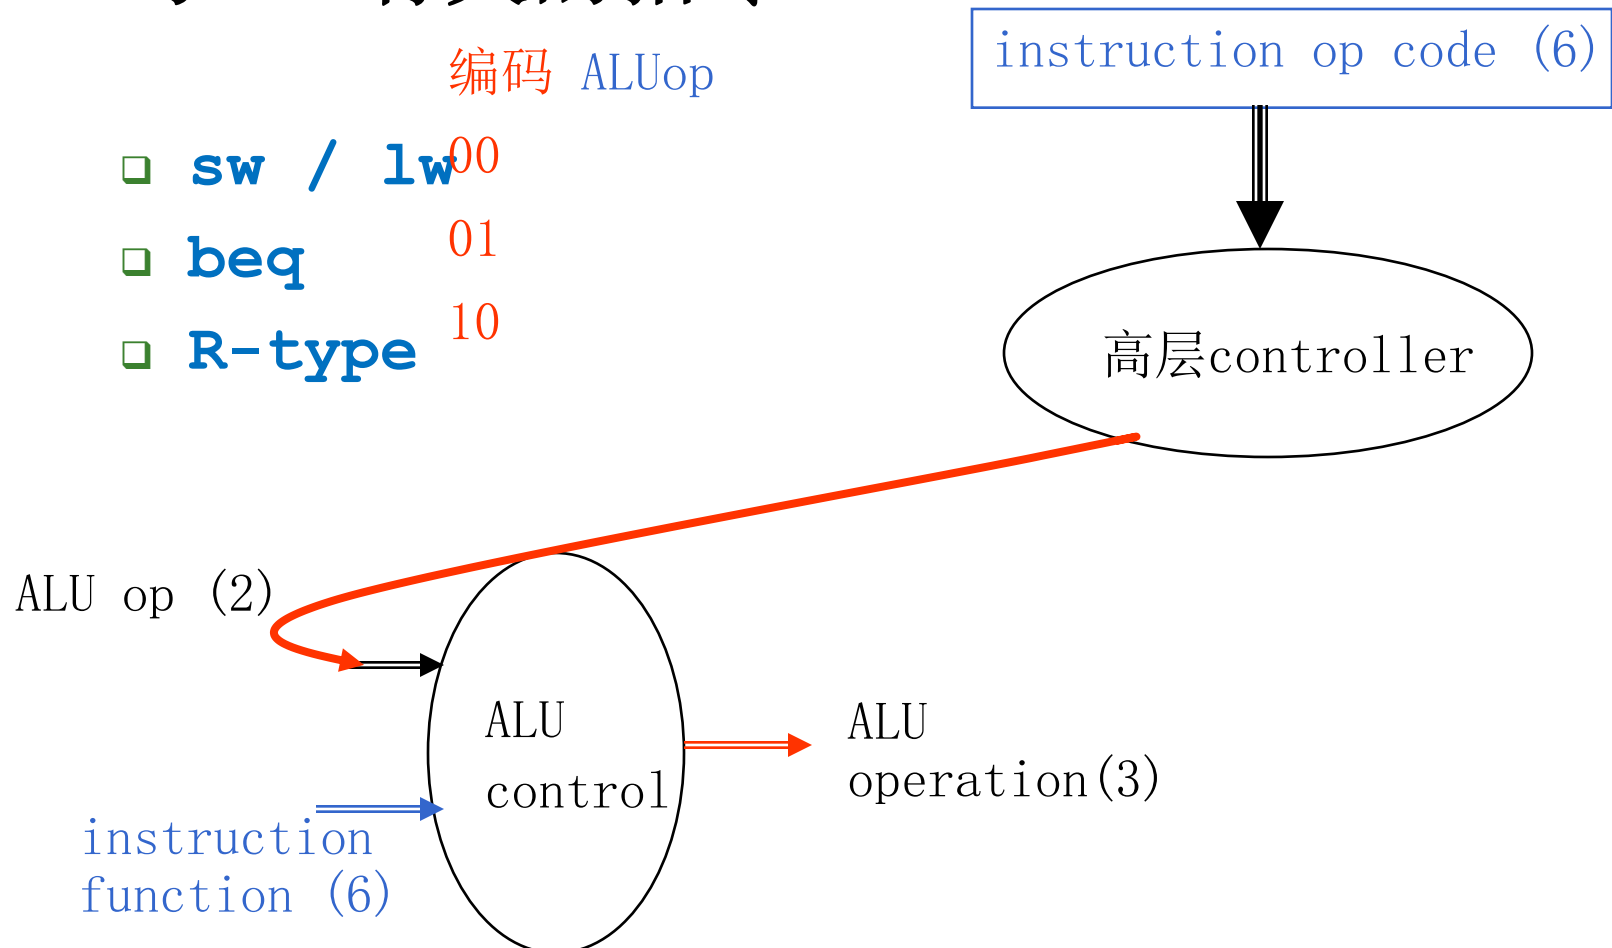
## ■ 与ALU有关的指令

编码 ALUop

□ **sw / lw** 00

□ **beq** 01

□ **R-type** 10



# The ALU control

- Must describe hardware to compute 3-bit ALU control input
  - given instruction type
    - 00 = lw, sw
    - 01 = beq,
    - 10 = R-type
  - function code for arithmetic

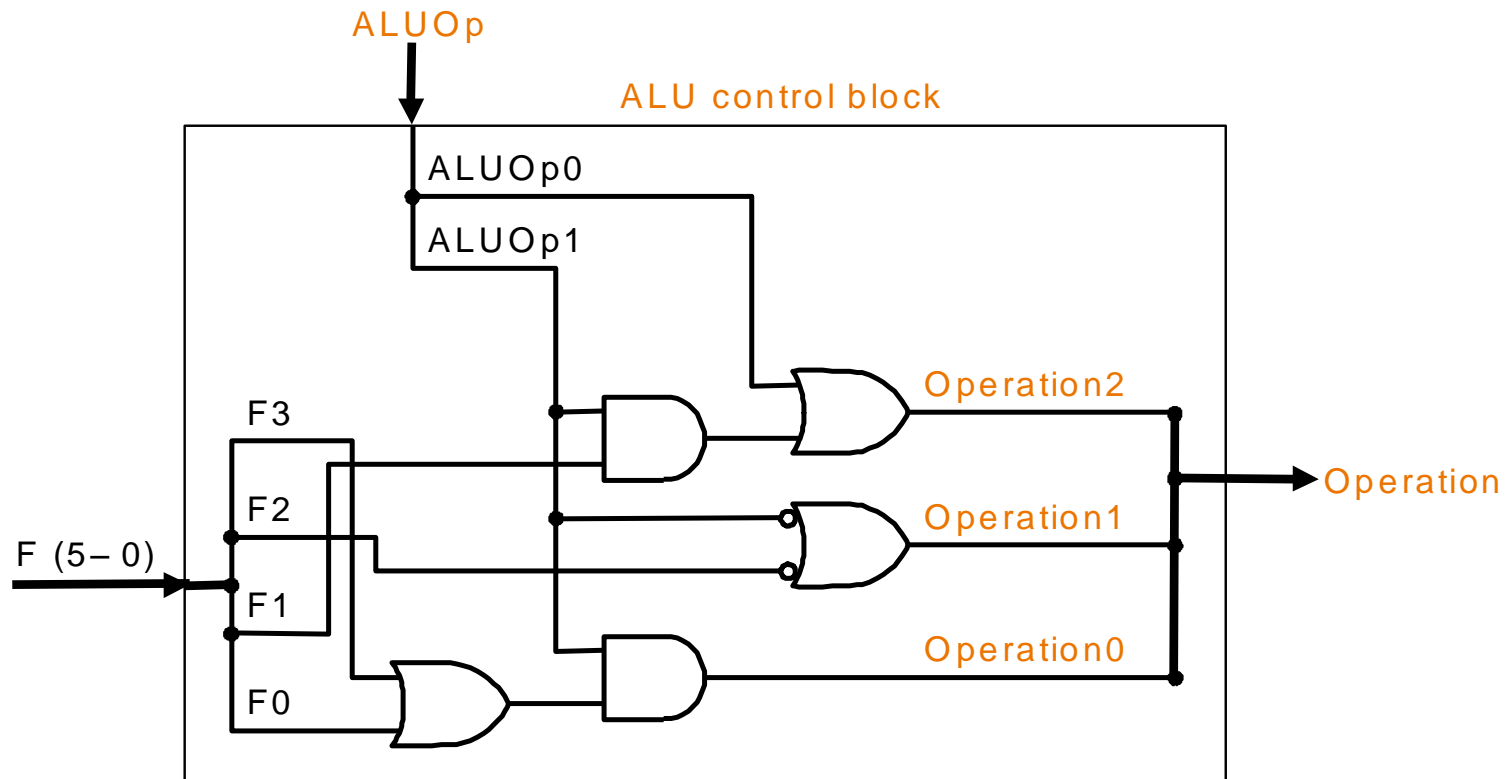
ALUOp computed from instruction type

- Describe into gates):

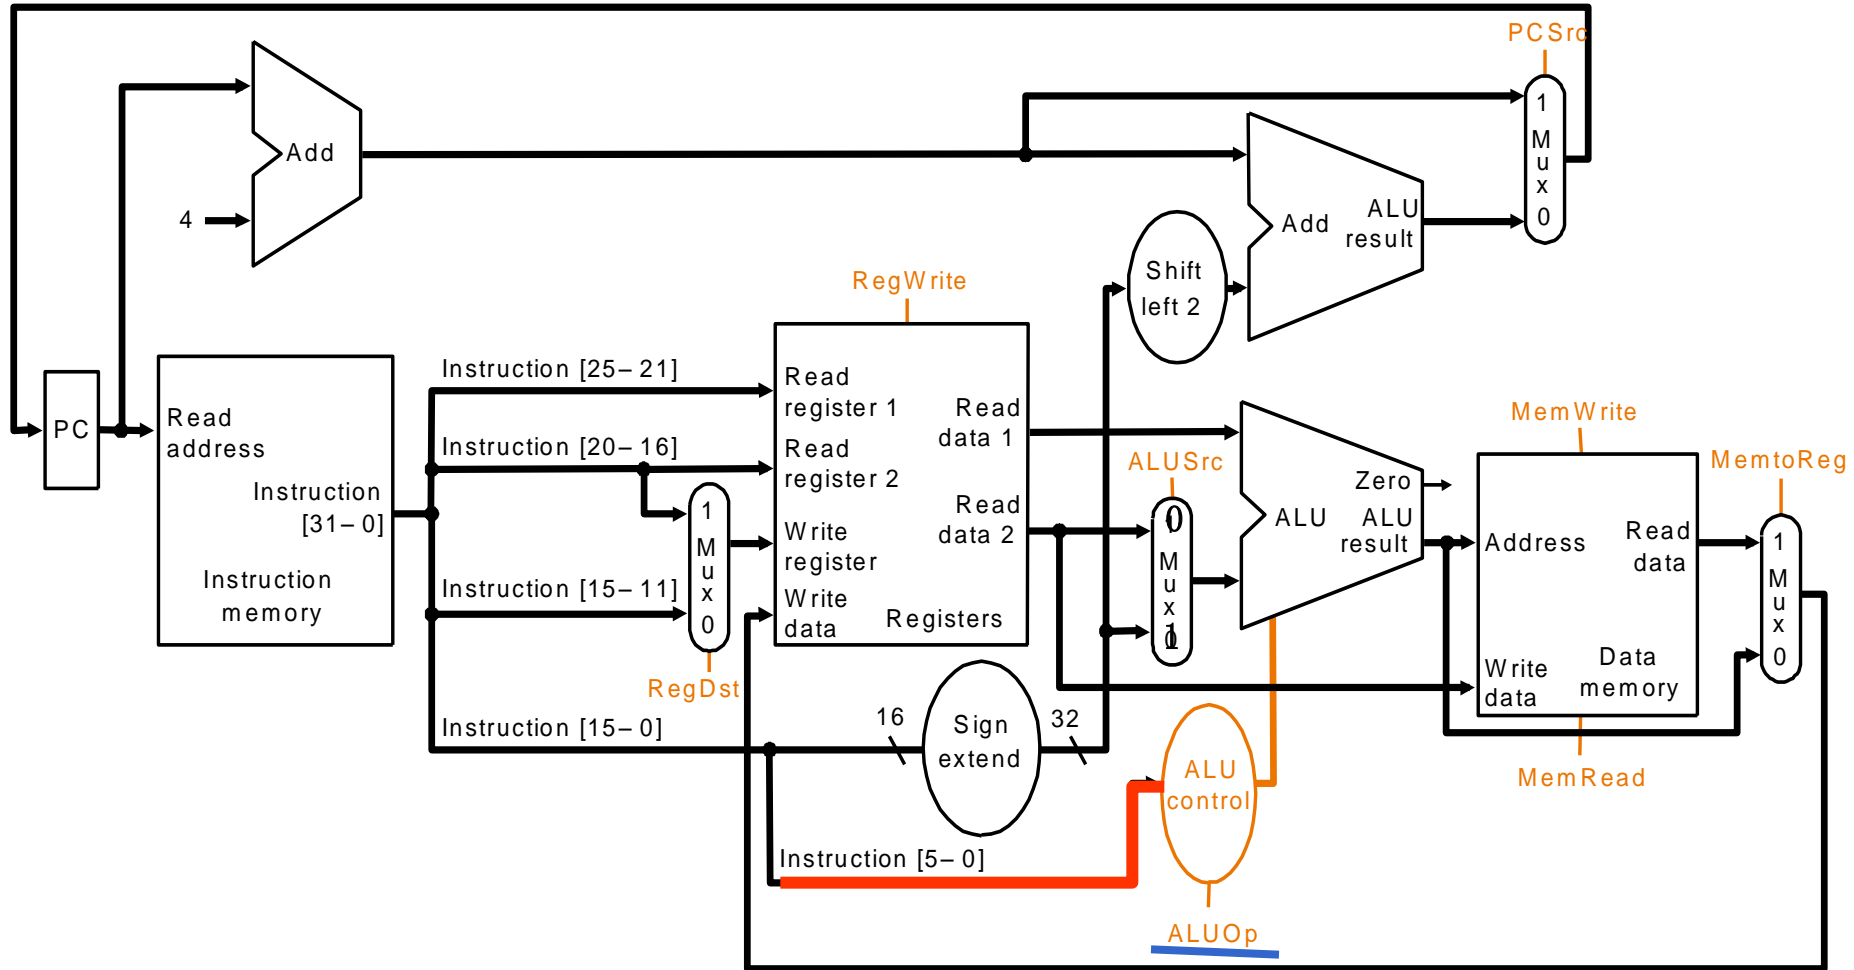
don't care

ALUOp		Funct field						Operation
ALUOp1	ALUOp0	F5	F4	F3	F2	F1	F0	
	0	X	X	X	X	X	X	010
X	1	X	X	X	X	X	X	110
1	X	X	X	0	0	0	0	010
1	X	X	X	0	0	1	0	110
1	X	X	X	0	1	0	0	000
1	X	X	X	0	1	0	1	001
1	X	X	X	1	0	1	0	111

# The ALU control





# The ALU control



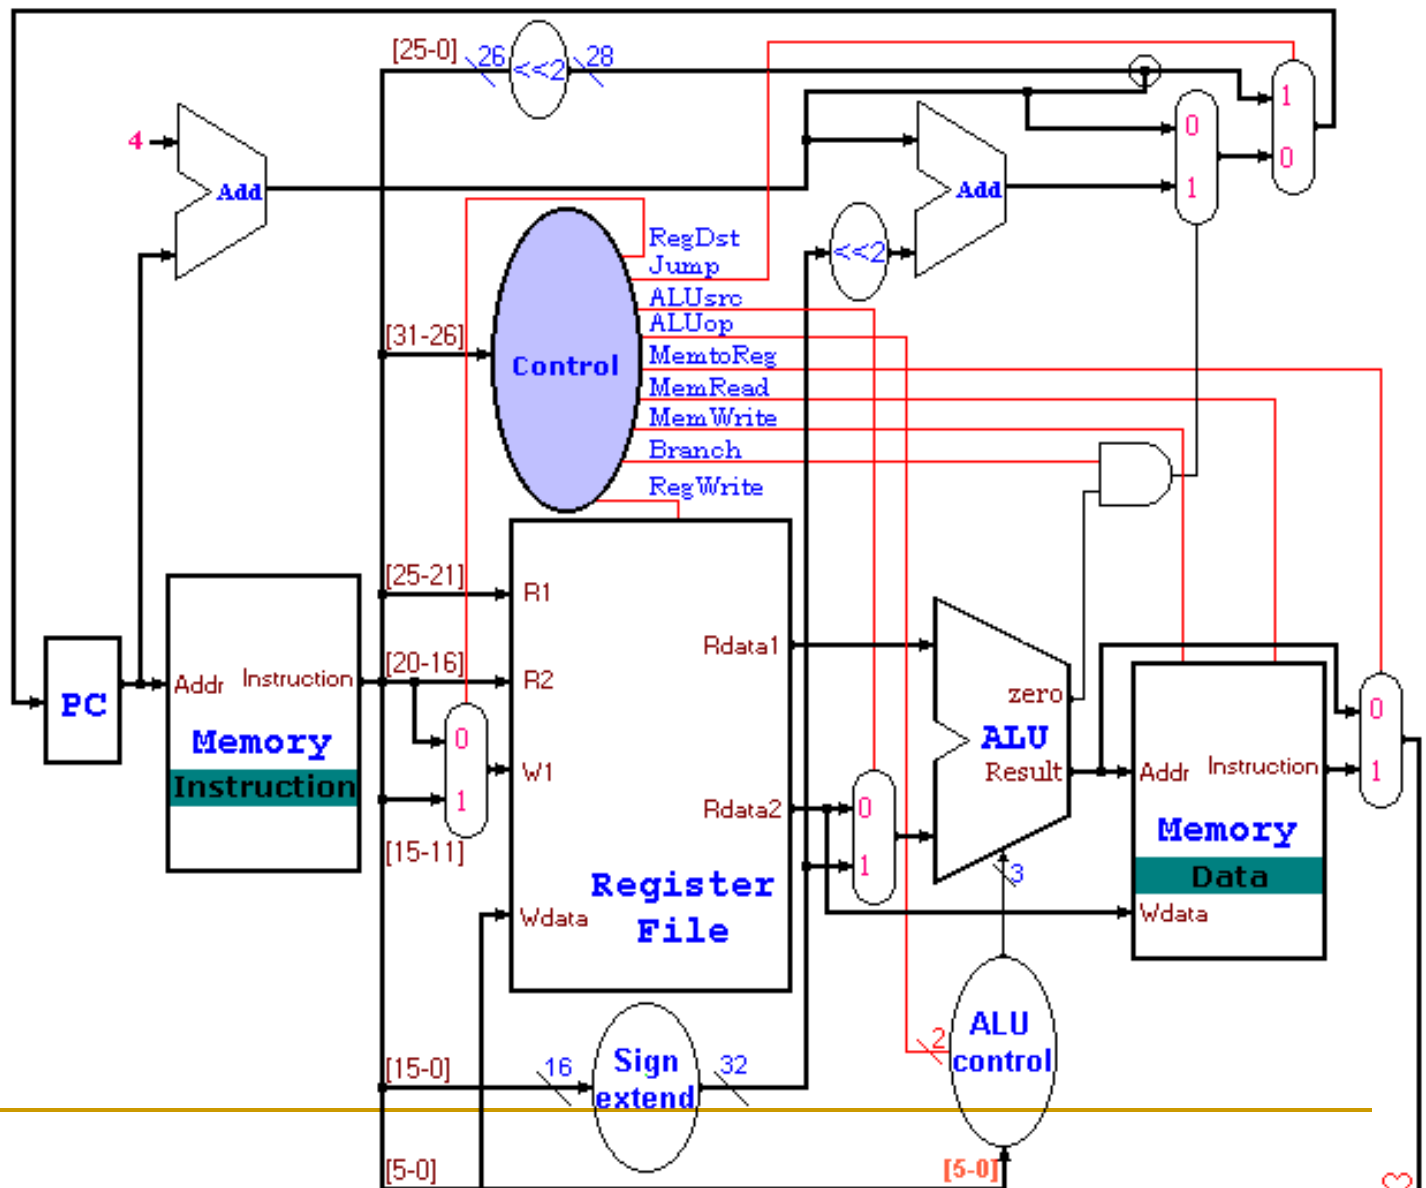
# Designing the Main Control Unit

- Main Control Unit function
  - ALU op (2)
  - other control signals (p. 359)
    - 4 Mux
    - 3 R/W





Instruction	RegDst	ALUSrc	Memto-Reg	Reg Write	Mem Read	Mem Write	Branch	ALUOp1	ALUp0
R-format	1	0	0	1	0	0	0	1	0
lw	0	1	1	1	1	0	0	0	0
sw	X	1	X	0	0	1	0	0	0
beq	X	0	X	0	0	0	1	0	1





# Control


- Simple combinational logic (truth tables)

R-type

lw 35

sw 43

beq 4



# j instruction

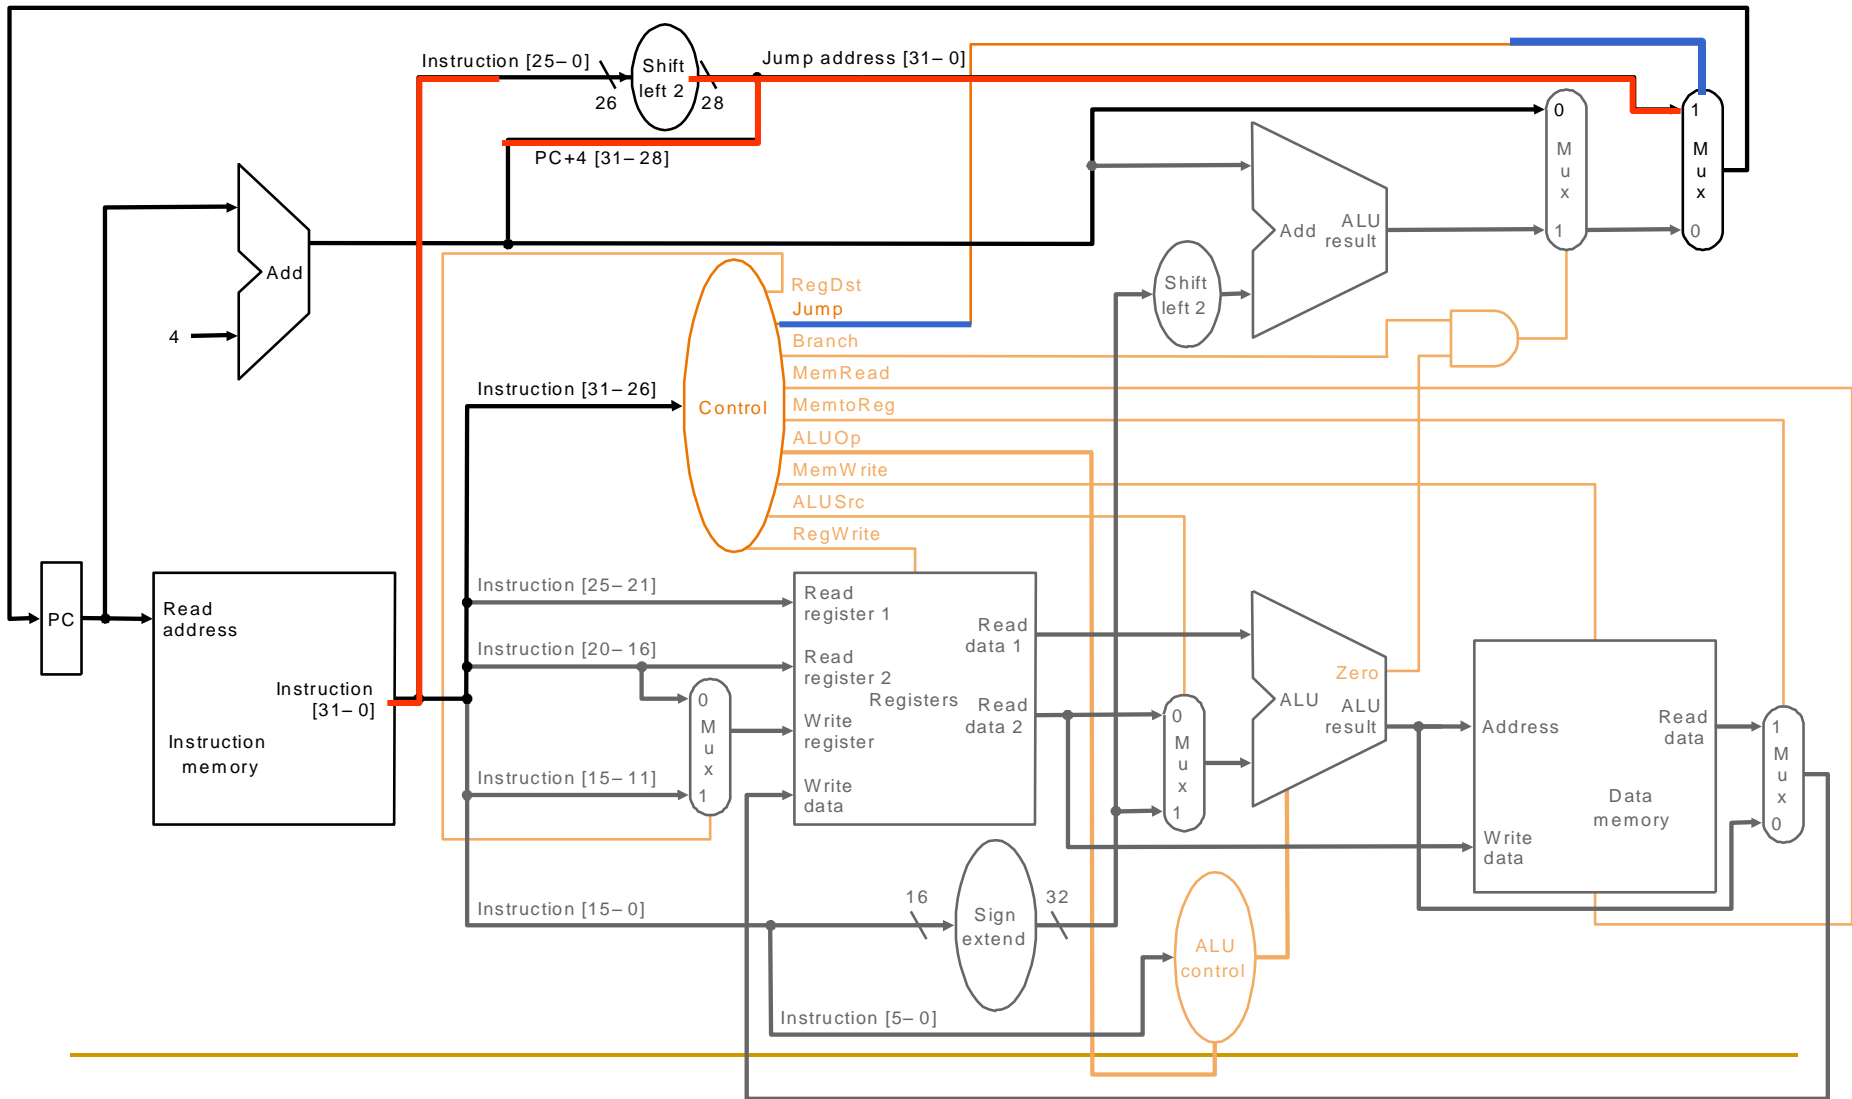
## ■ instruction format

### □ j Label



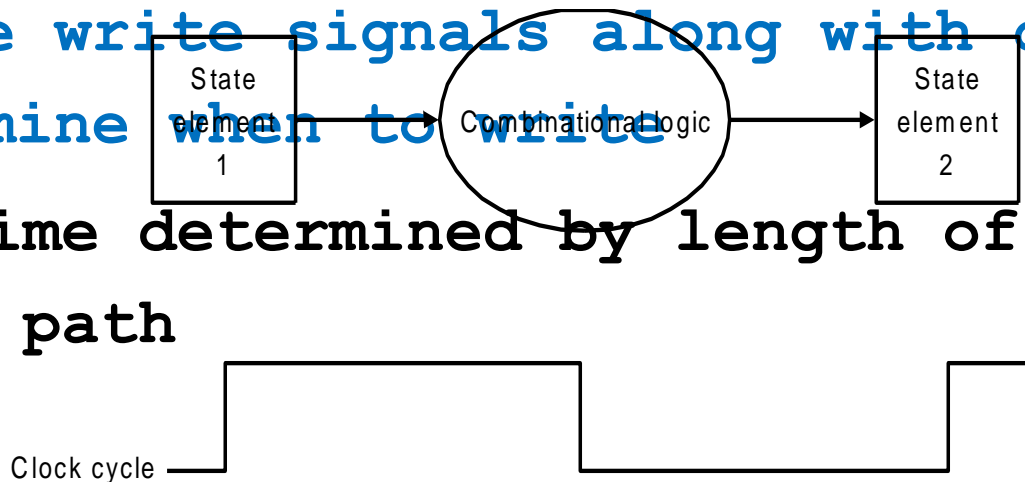
## ■ Implementation

$$\square pc = pc_{28 \sim 31} + \frac{26\text{bits-address}}{4} *$$



# Simple Control Structure

- All of the logic is combinational
- We wait for everything to settle down, and the right thing to be done
  - ALU might not produce right answer?right away
  - we use write signals along with clock to determine when to write
- Cycle time determined by length of the longest path

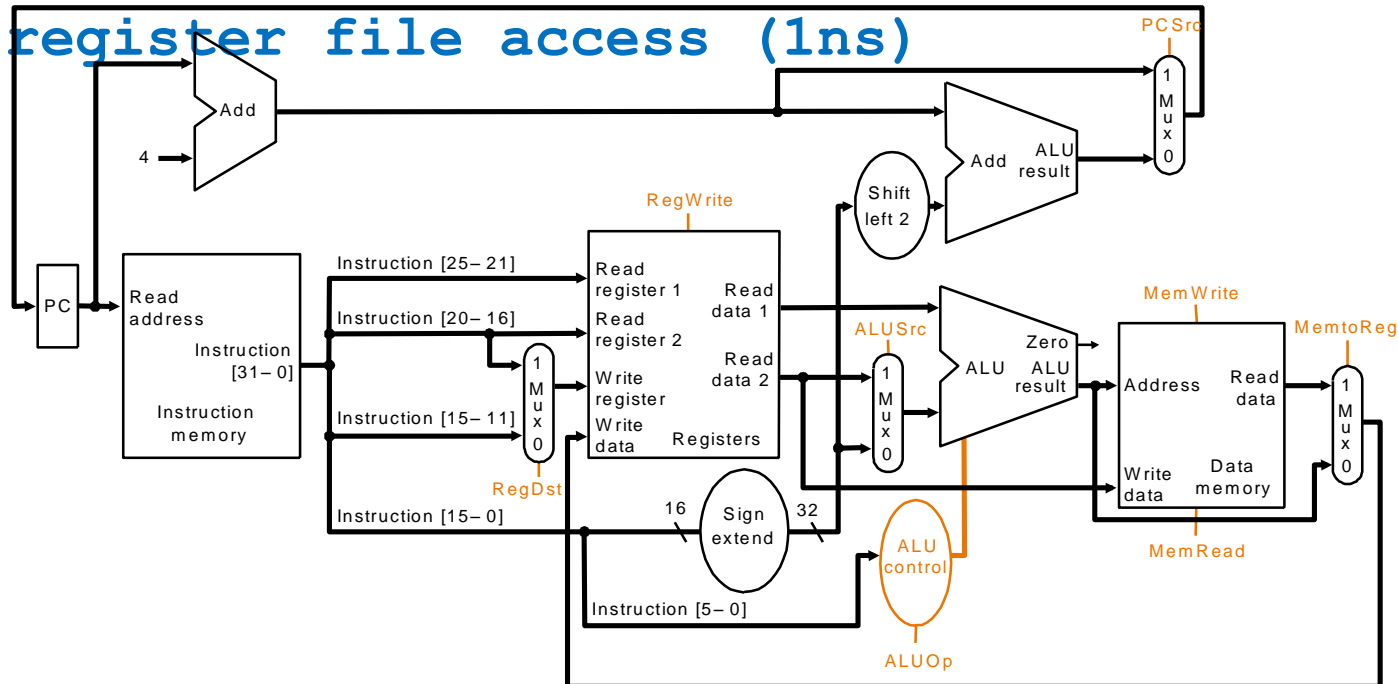


*We are ignoring some details like setup and hold times*

# Single Cycle

- Calculate cycle time assuming negligible delays except:

- memory (2ns), ALU and adders (2ns), register file access (1ns)



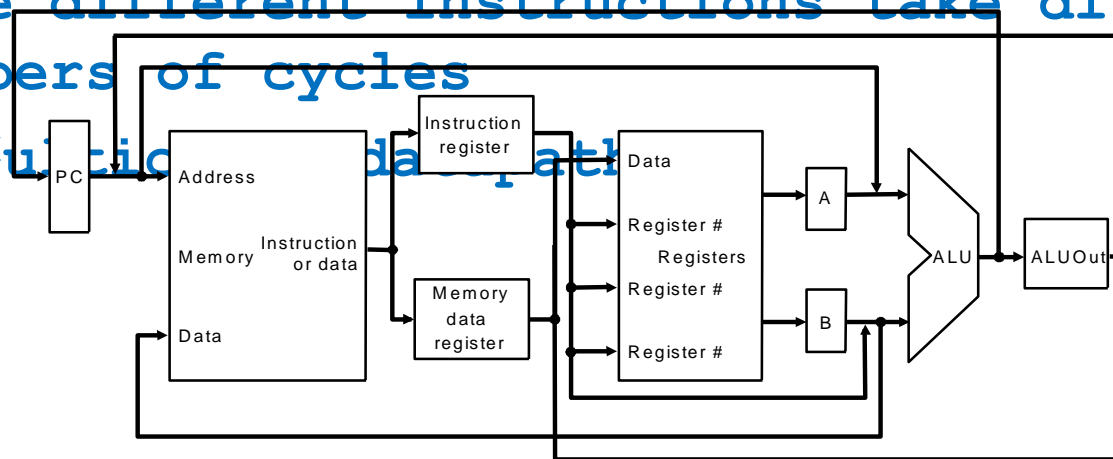
# Where we are headed

## ■ Single Cycle Problems:

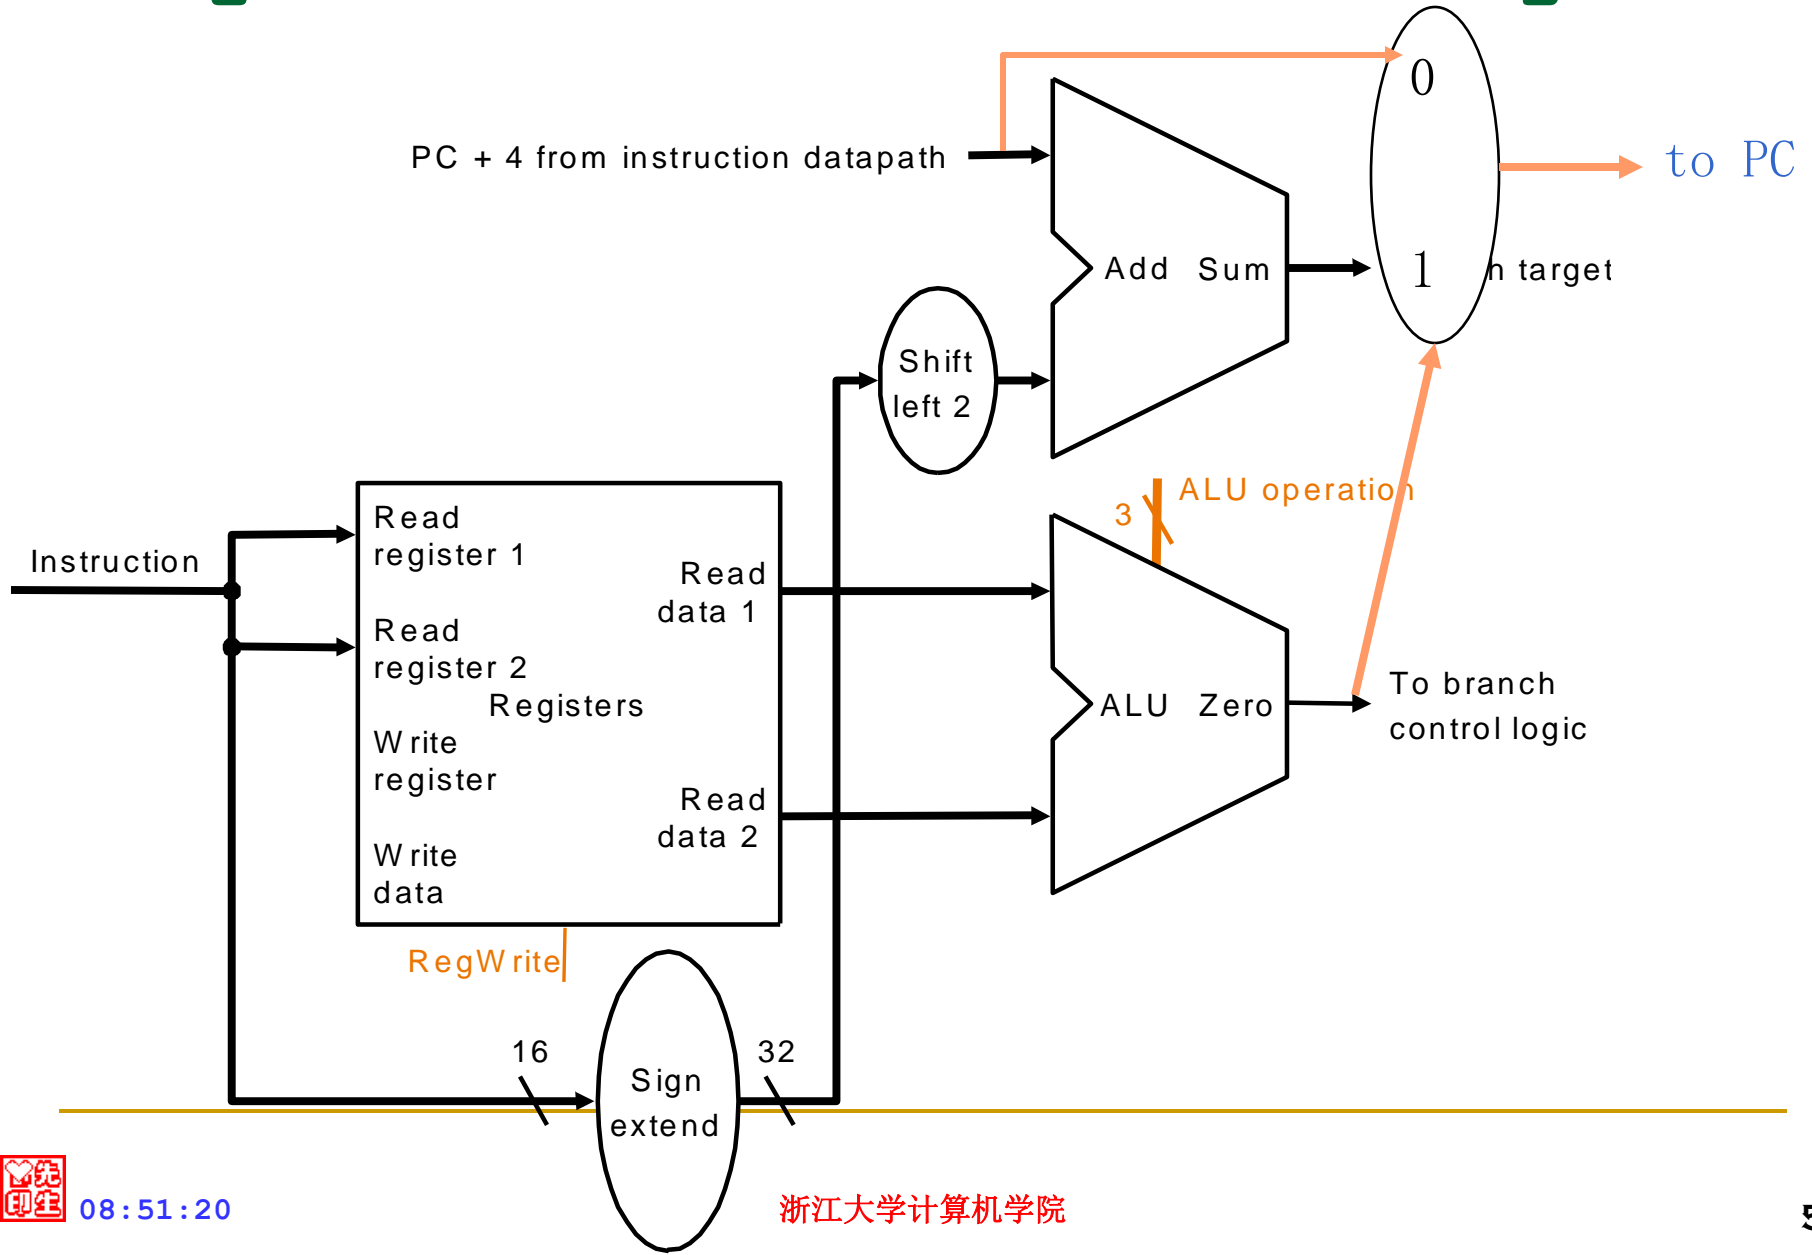
- ❑ what if we had a more complicated instruction like floating point?
- ❑ wasteful of area

## ■ One Solution:

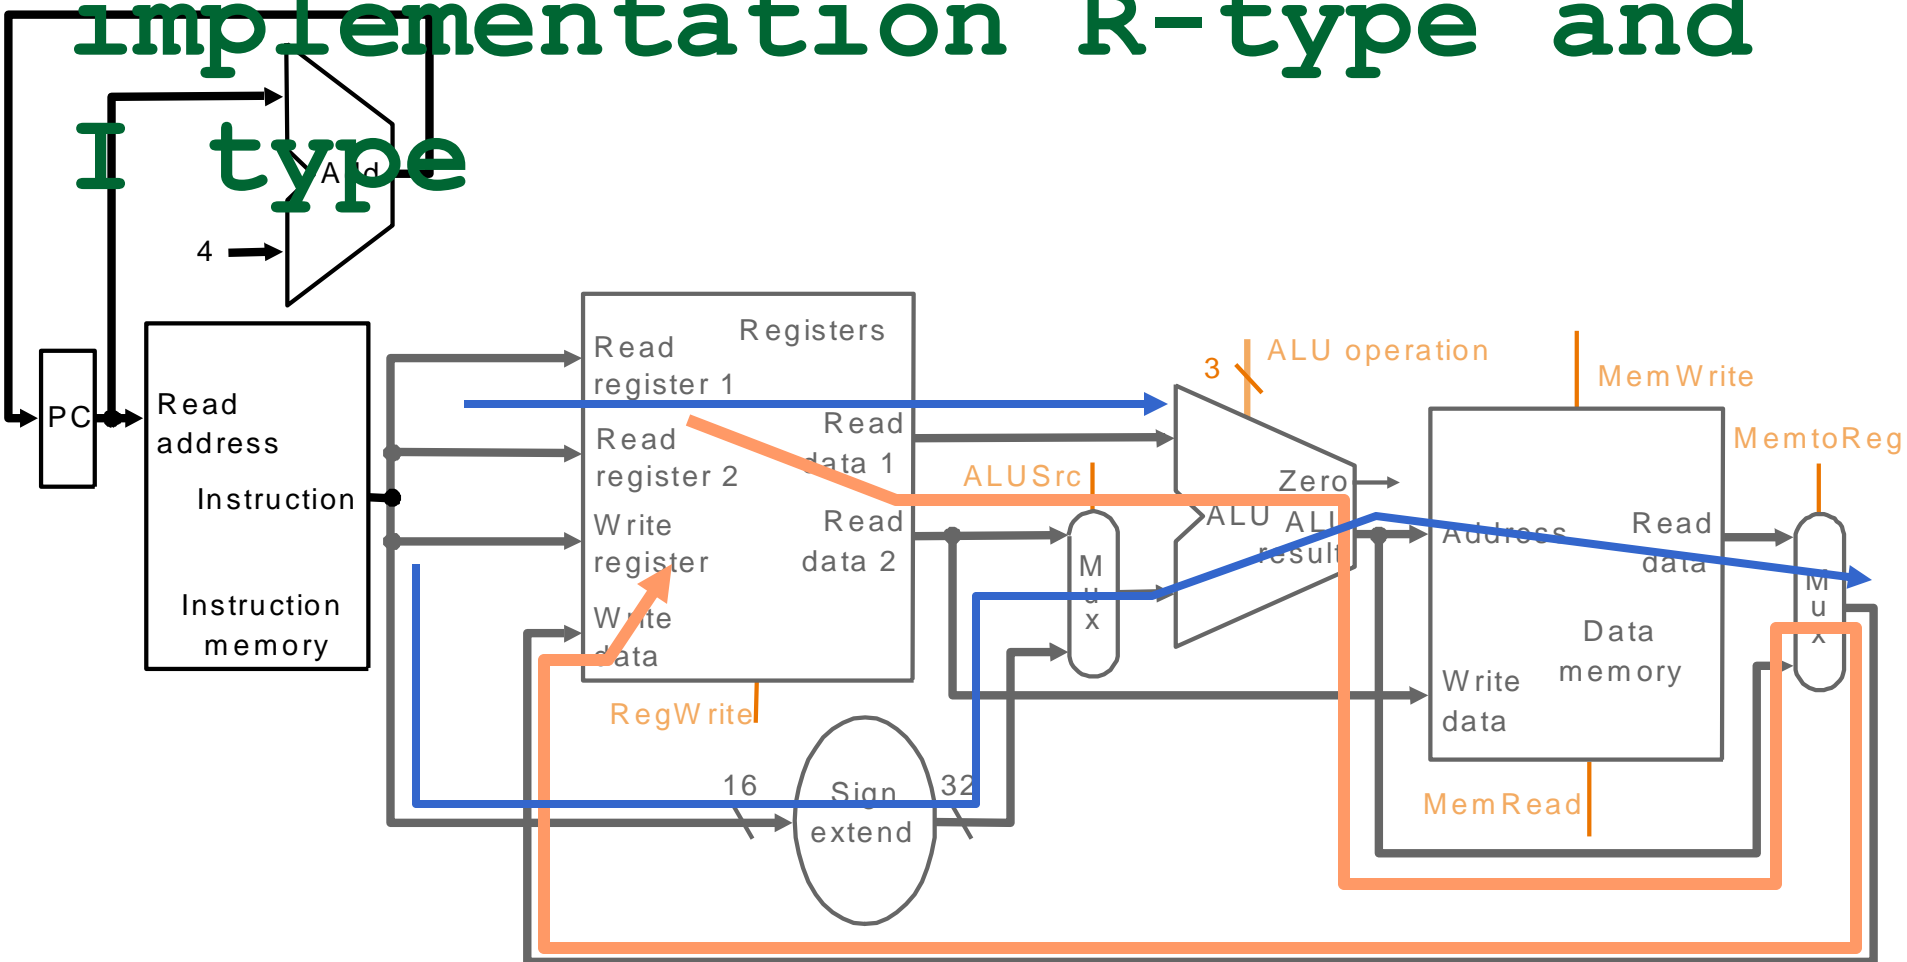
- ❑ use a **pipelined** cycle time
- ❑ have different instructions take different numbers of cycles
- ❑ a **dedicated** data path



# Implementation of *beq*




# combine the implementation R-type and I type





# Building the Datapath

- Use multiplexors to stitch them together





# Control

- Selecting the operations to perform (ALU, read/write, etc.)
- Controlling the flow of data (multiplexor inputs)
- Information comes from the 32 bits of the instruction




# Control

## ■ Example:

add \$8, \$17, \$18      Instruction Format:

000000      10001      10010      01000  
 00000      100000




## ■ ALU's operation based on instruction type and function code

# Control

- e.g., what should the ALU do with this instruction

- Example: `lw $1, 100($2)`



op	rs	rt	16 bit offset
ALU control input			

000 AND

001 OR

010 add

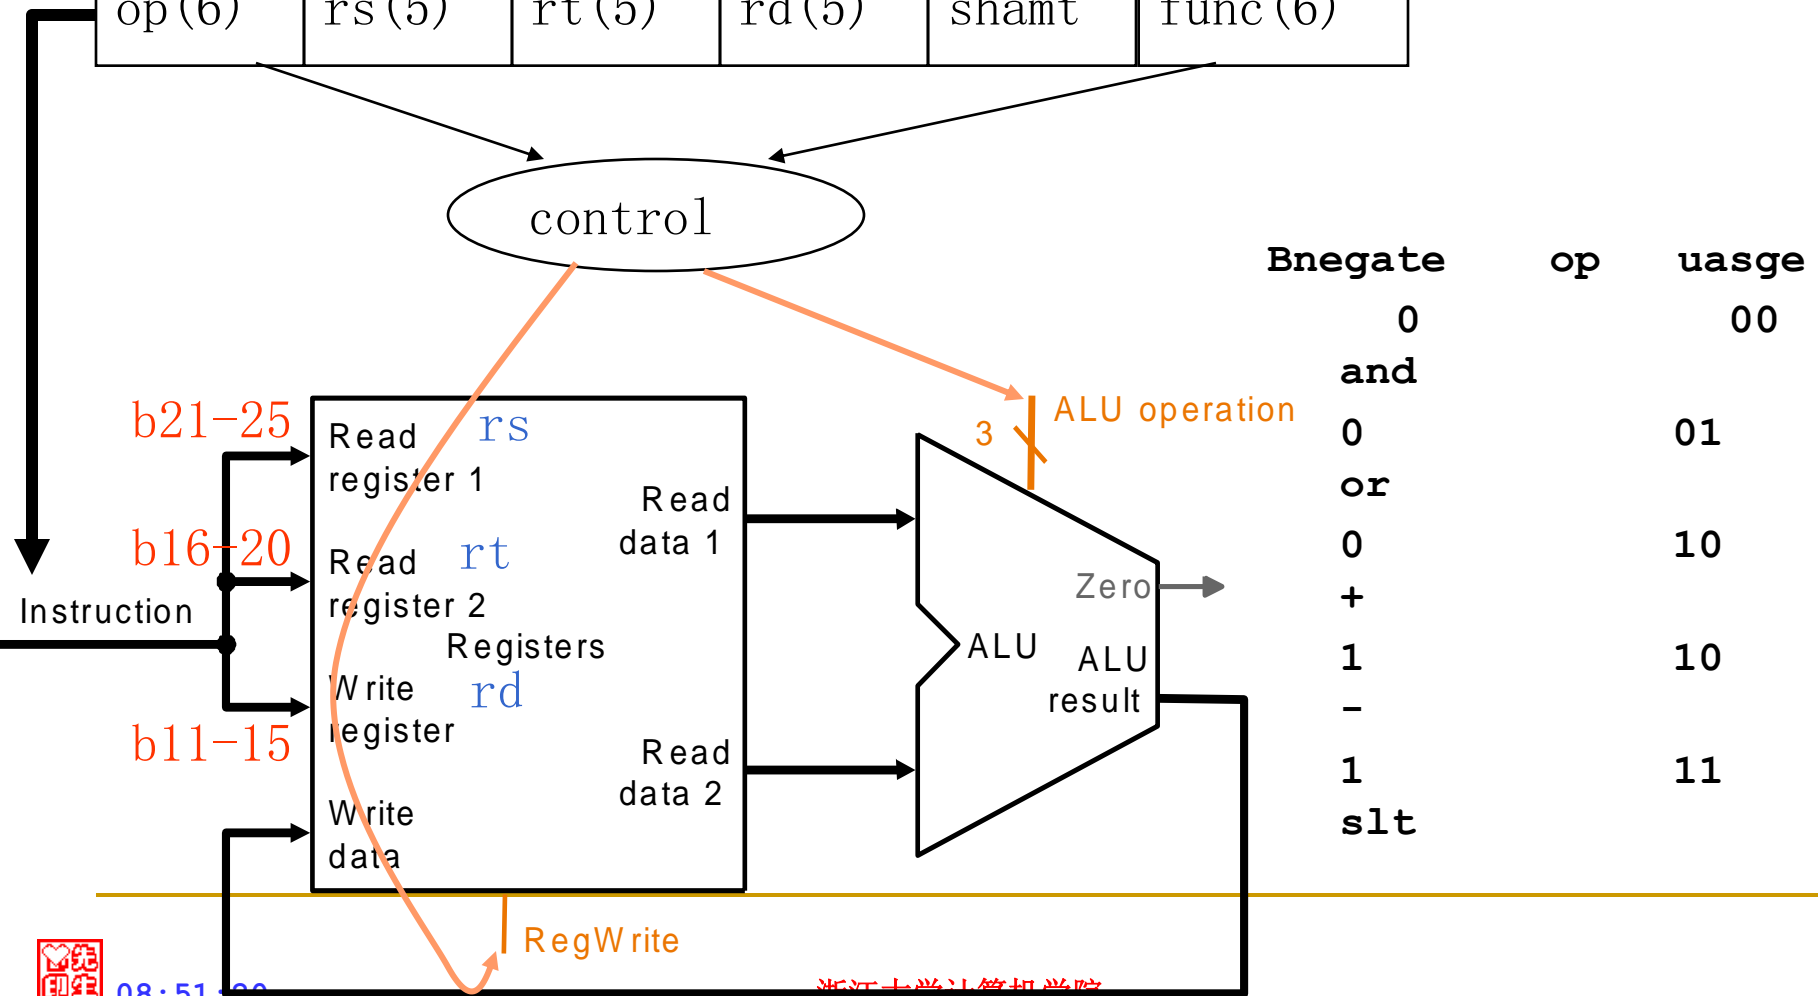

110 subtract

111 set-on-less-than

Why is the code for subtract 110 and not 011?

# The ALU control

instruction format:



# The ALU control


## ■ 与ALU有关的指令

编码 ALUop

□ **sw / lw** 00

□ **beq** 01

□ **R-type** 10



# The ALU control

- Must describe hardware to compute 3-bit ALU control input
  - given instruction type
    - 00 = lw, sw
    - 01 = beq,
    - 10 = R-type
  - function code for arithmetic

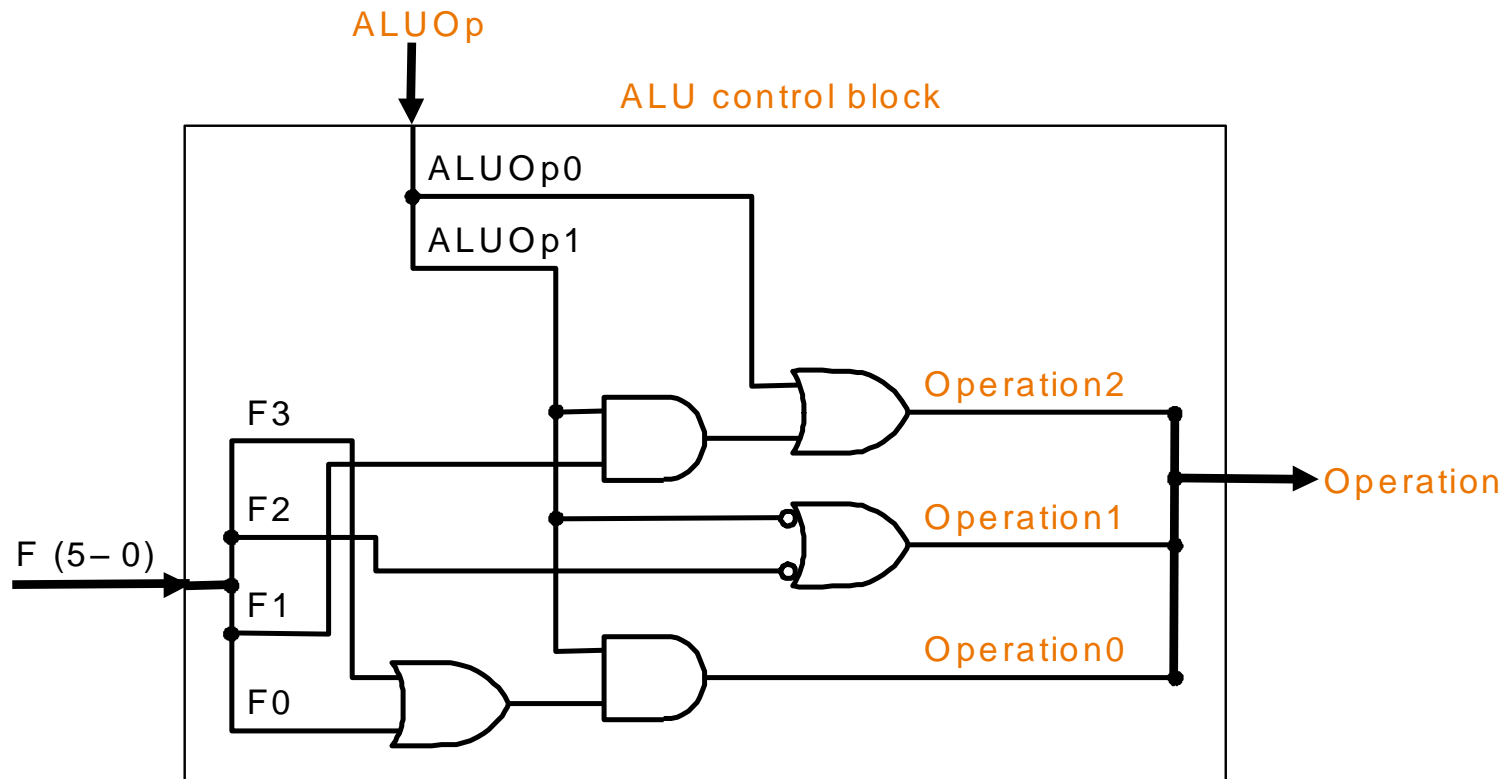
ALUOp computed from instruction type

- Describe into gates) : n turn

don't care

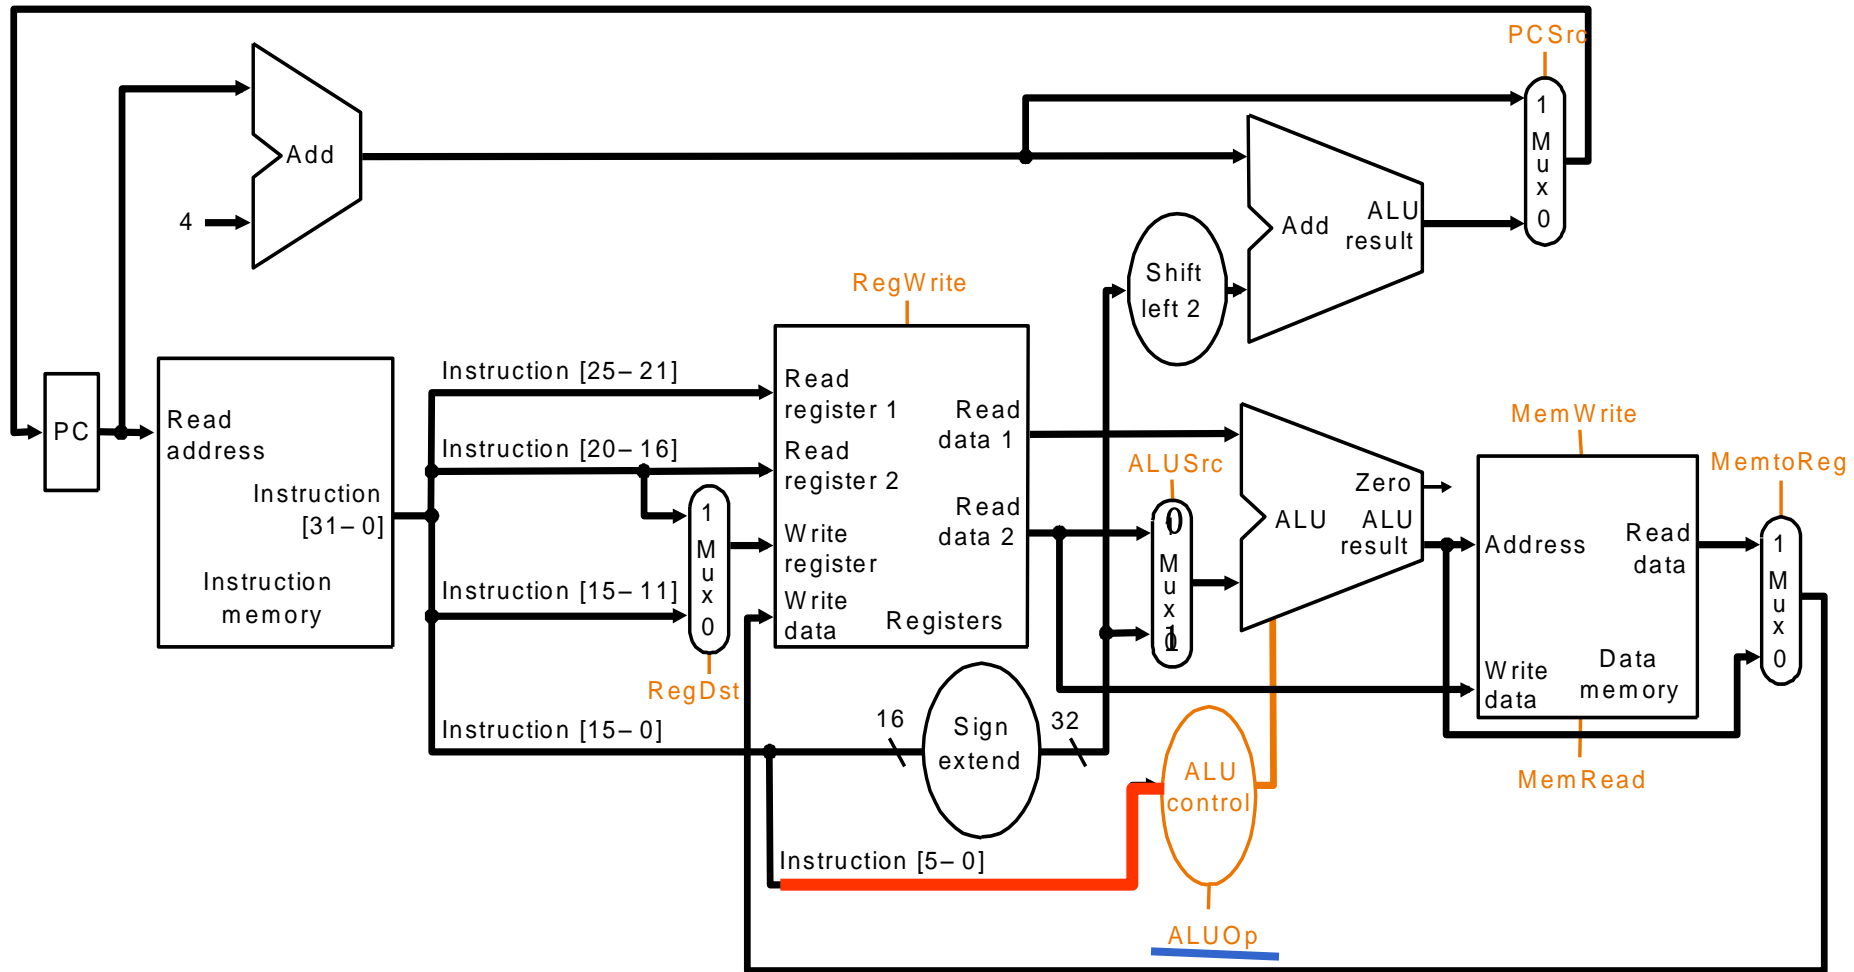
ALUOp		Funct field						Operation
ALUOp1	ALUOp0	F5	F4	F3	F2	F1	F0	
	0	X	X	X	X	X	X	010
X	1	X	X	X	X	X	X	110
1	X	X	X	0	0	0	0	010
	X	X	X	0	0	1	0	110
1	X	X	X	0	1	0	0	000
1	X	X	X	0	1	0	1	001
1	X	X	X	1	0	1	0	111

# The ALU control






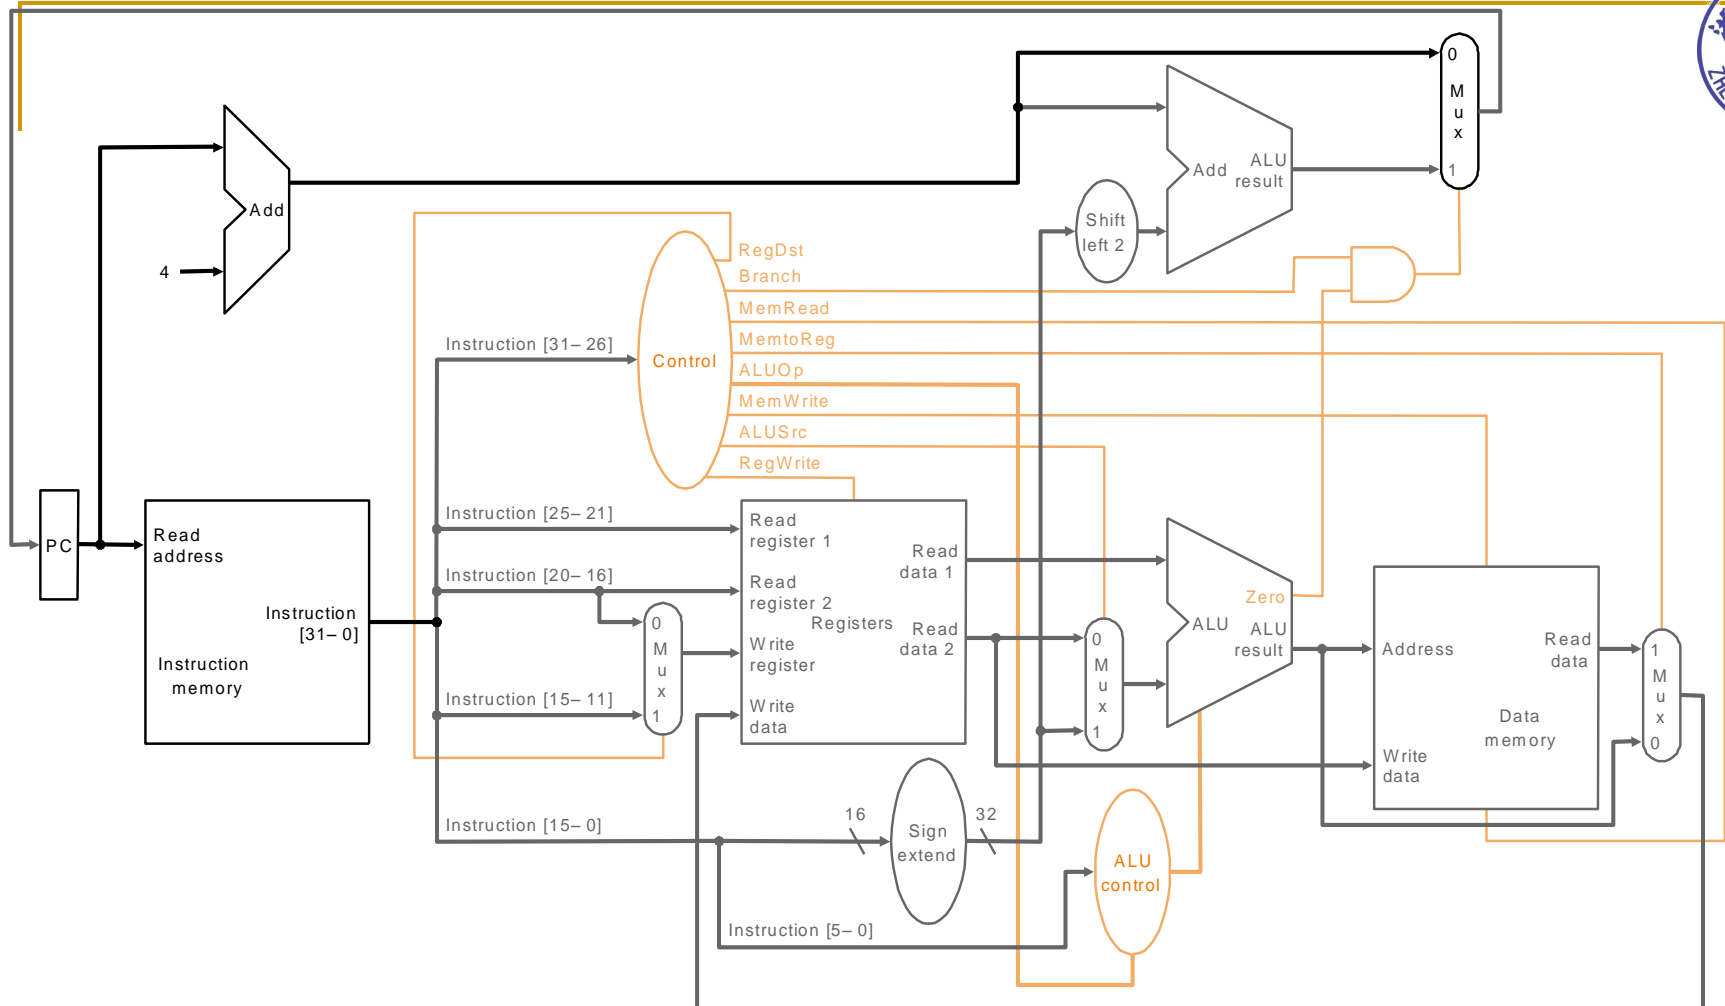
# The ALU control



# Designing the Main Control Unit

- Main Control Unit function
  - ALU op (2)
  - other control signals (p. 359)
    - 4 Mux
    - 3 R/W





Instruction	RegDst	ALUSrc	Memto-Reg	Reg Write	Mem Read	Mem Write	Branch	ALUOp1	ALUp0
R-format	1	0	0	1	0	0	0	1	0
lw	0	1	1	1	1	0	0	0	0
sw	X	1	X	0	0	1	0	0	0
beq	X	0	X	0	0	0	1	0	1

# Control

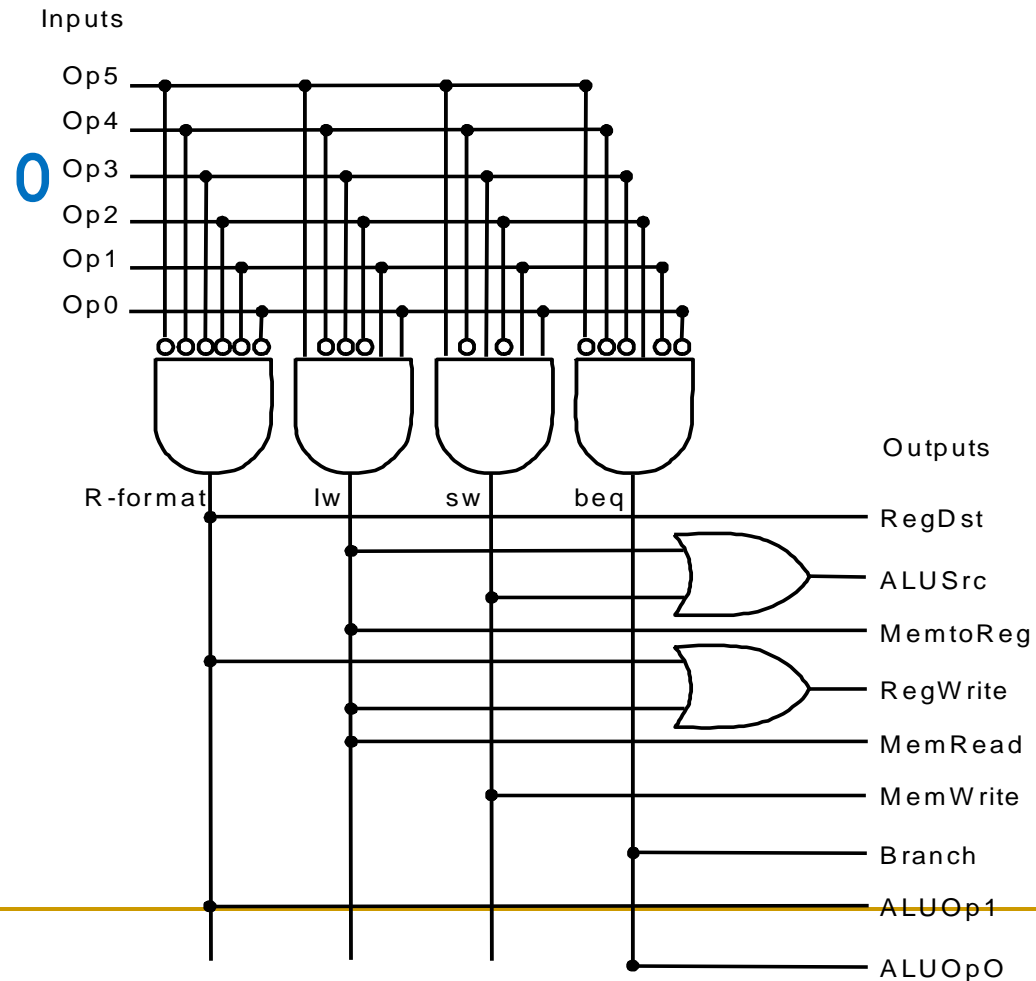
- Simple combinational logic (truth tables)

R-type

lw 35

sw 43

beq 4



# j instruction


## ■ instruction format

### □ j Label




## ■ Implementation

$$\square pc = pc_{28 \sim 31} + \frac{26\text{bits-address}}{4} *$$



# Our Simple Control

- **Structure** All of the logic is combinational
- We wait for everything to settle down, and the right thing to be done
  - ALU might not produce right answer? right away
  - we use write signals along with clock to determine when to write
- Cycle time determined by length of the longest path




*We are ignoring some details like setup and hold times*

# Single Cycle Implementation

■ Calculate cycle time assuming negligible delays except:

□ memory (2ns), ALU and adders (2ns), register file access (1ns)



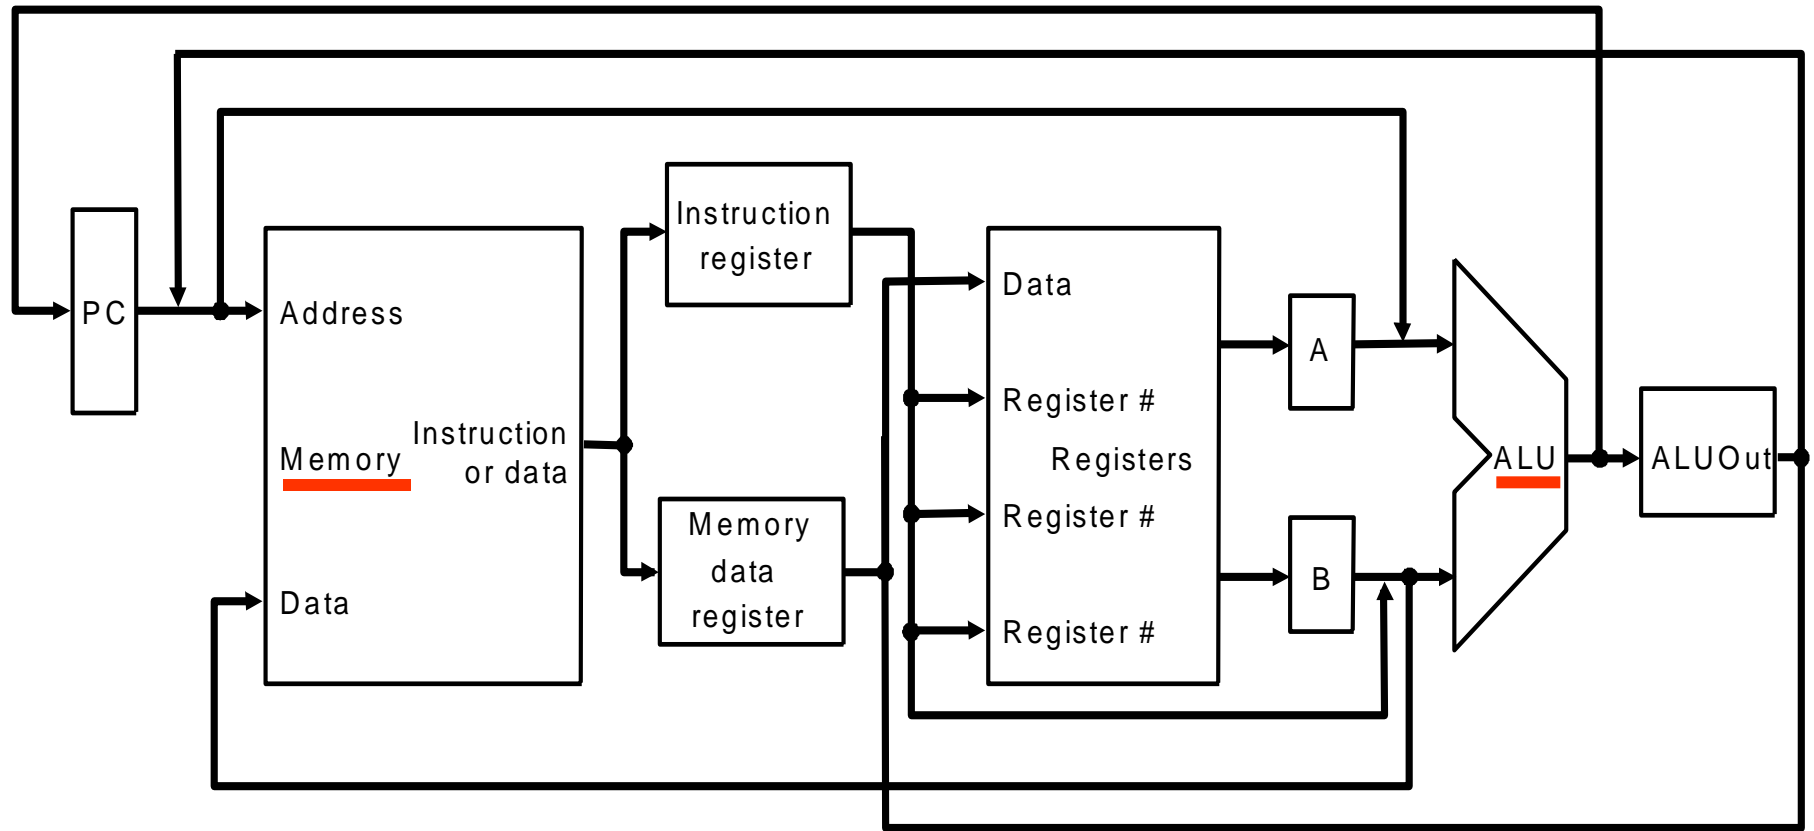


# Where we are headed

- Single Cycle Problems:
  - what if we had a more complicated instruction like floating point?
  - wasteful of area
- One Solution:
  - use a smaller cycle time
  - have different instructions take different numbers of cycles

# Multicycle Approach

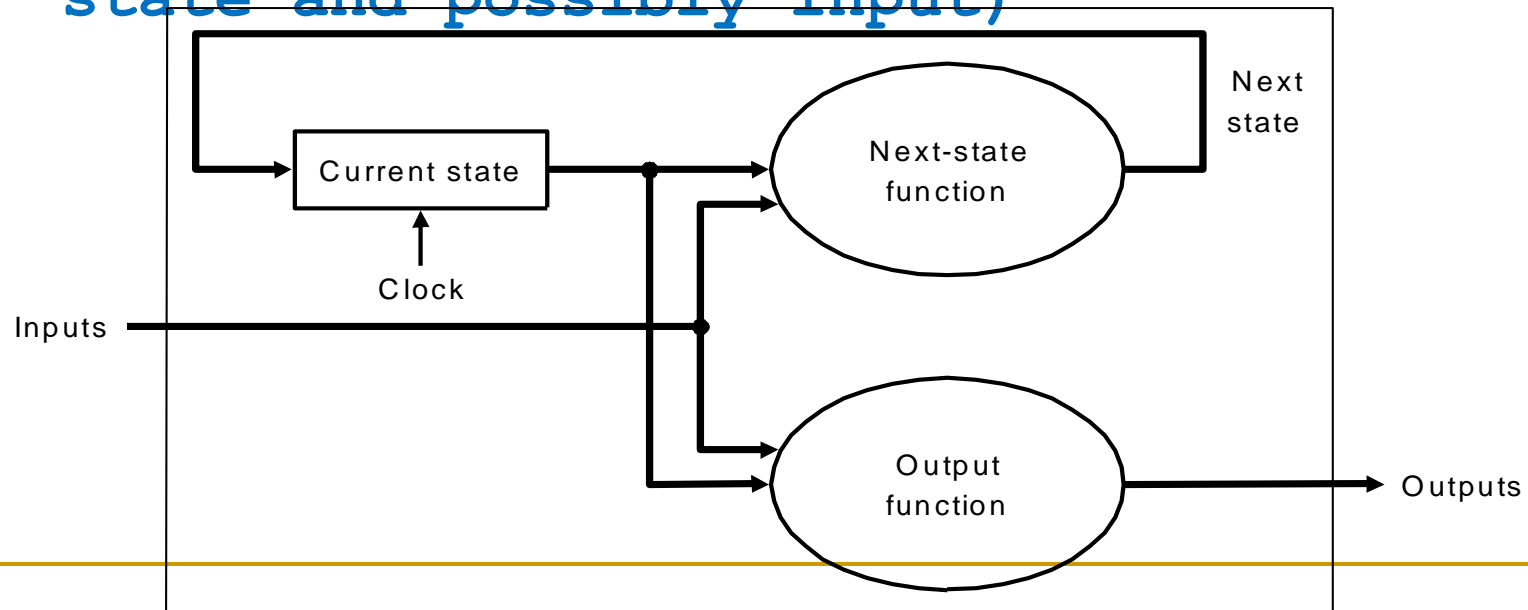
- We will be reusing functional units
  - ALU used to compute address and to increment PC
  - Memory used for instruction and data
- Our control signals will not be determined solely by instruction
  - e.g., what should the ALU do for a subtract instruction?
- We will use a finite state machine for control



# Review: finite state machines

## ■ Finite state machines:

- a set of states
- next state function (determined by current state and the input)
- output function (determined by current state and possibly input)

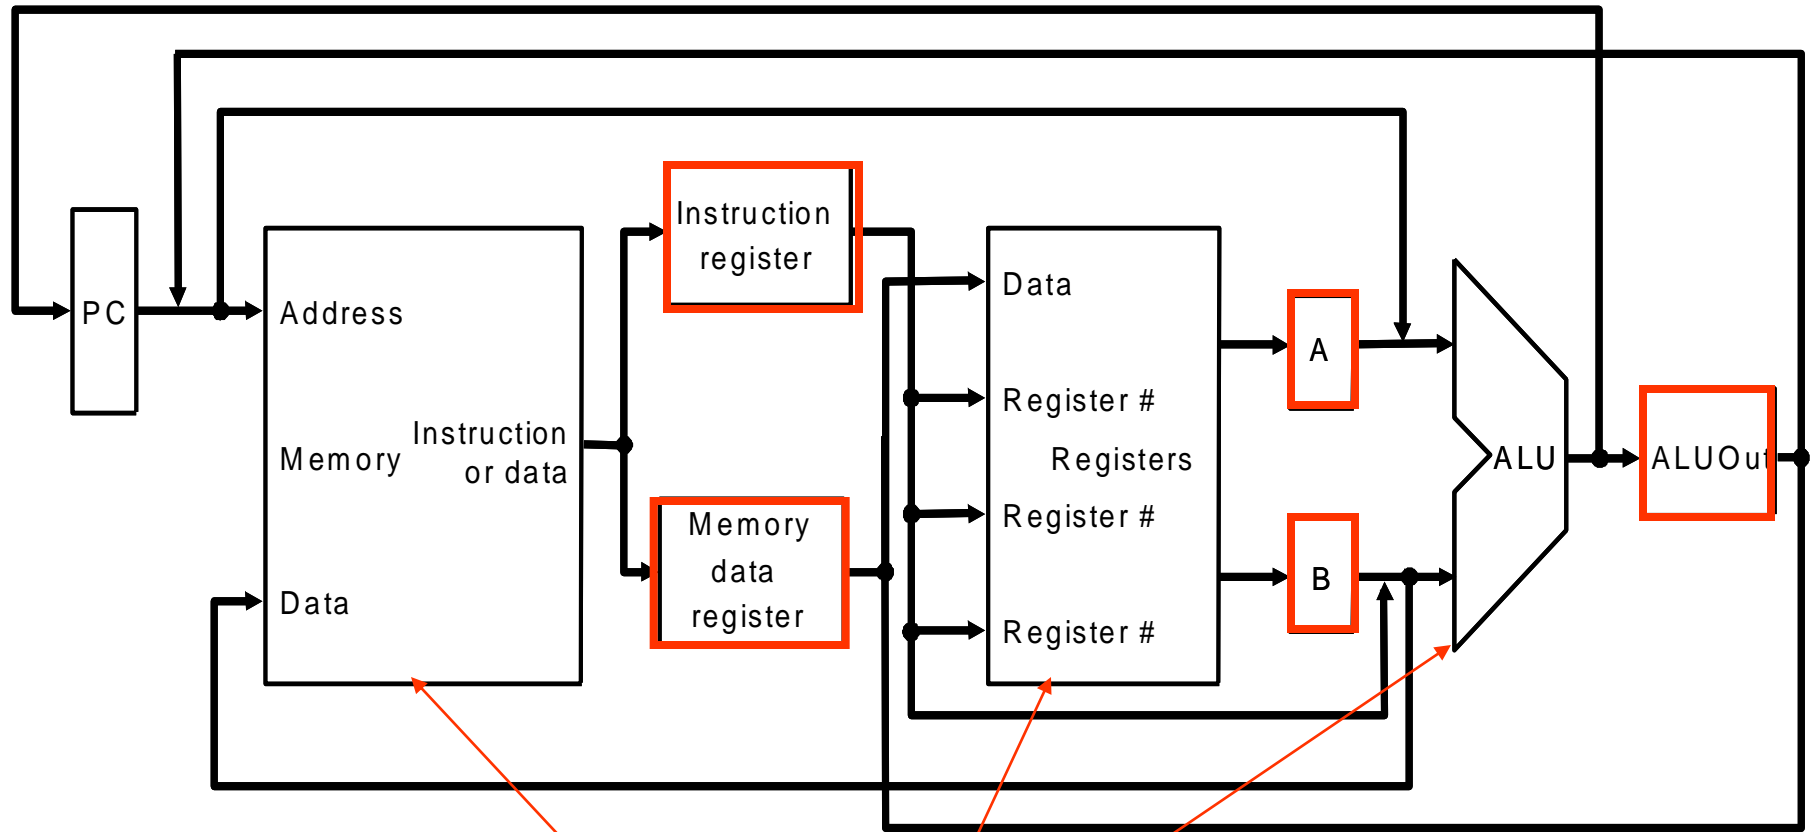


# Multicycle Approach

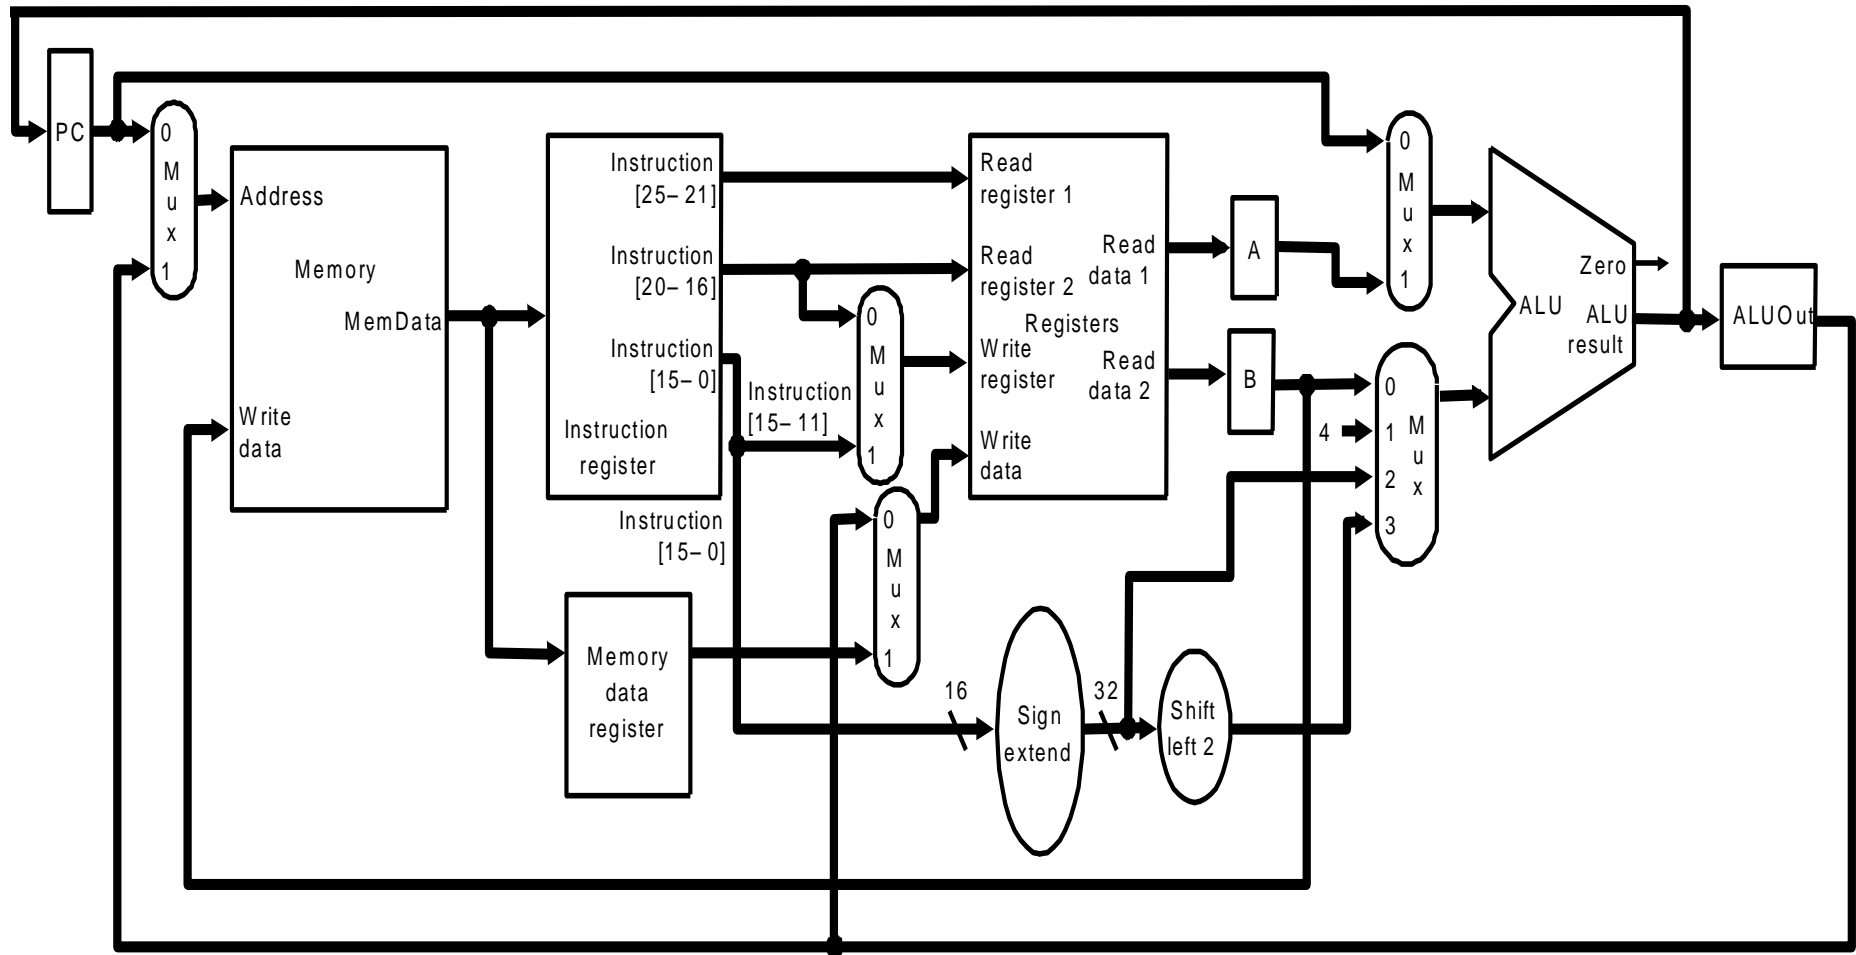
- Break up the instructions into steps, each step takes a cycle
  - balance the amount of work to be done
  - restrict each cycle to use only one major functional unit
- At the end of a cycle
  - store values for use in later cycles (easiest thing to do)
  - introduce additional internal registers

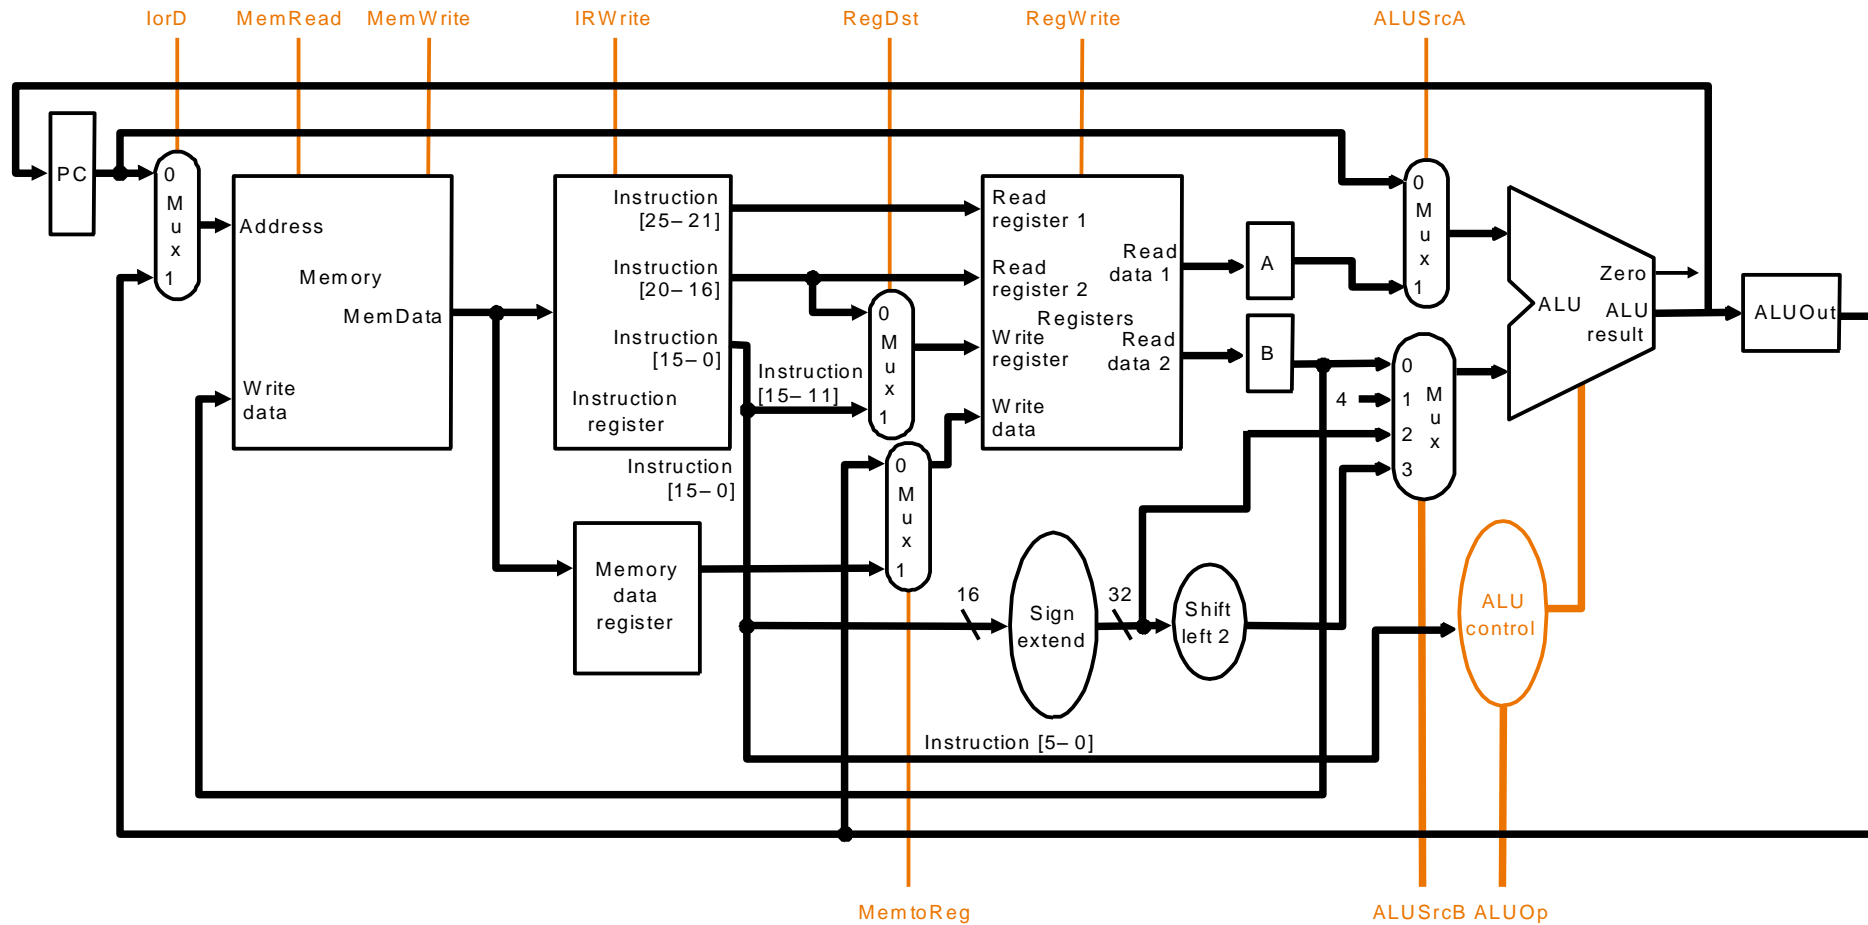
# Multicycle Approach

## ■ a Multicycle datapath



shared unit







## ■ PC 的改变方式

□ seq:  $pc = pc + 4$

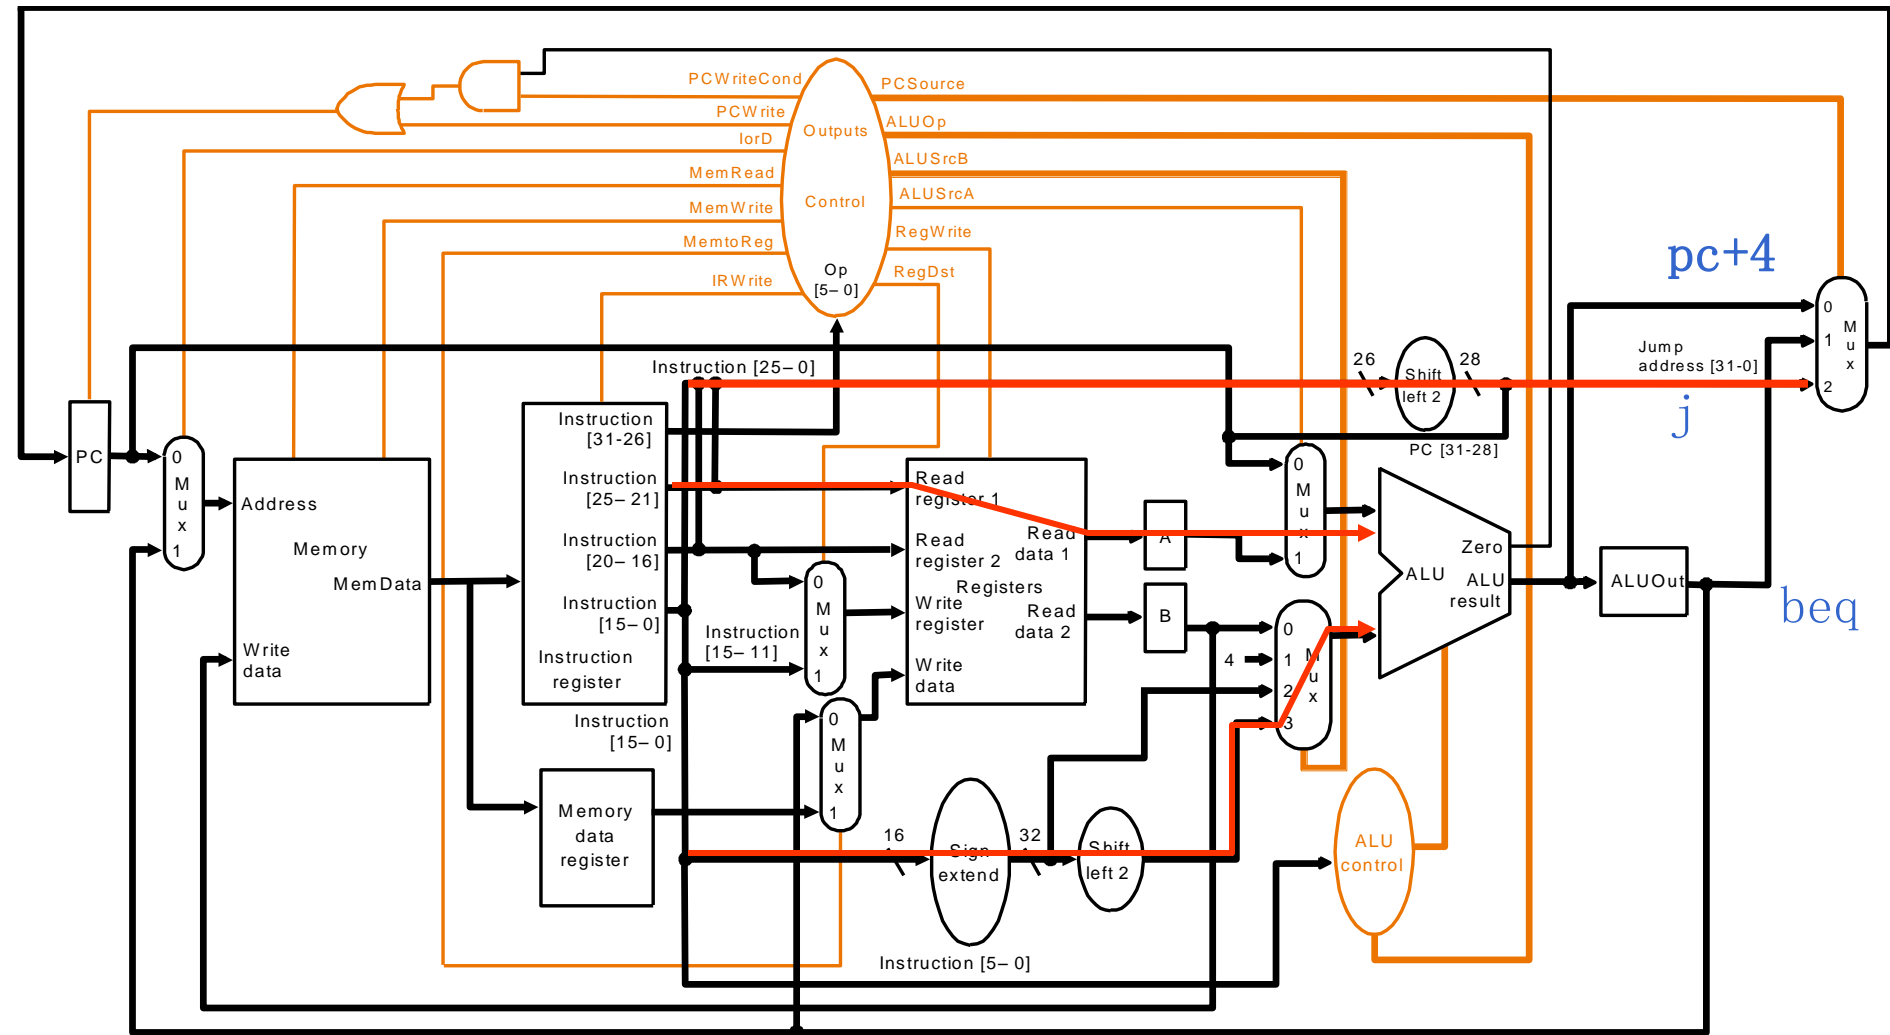
□ beq:  $pc = pc + offset * 4$

□ j :  $pc = pc_{31-28} + IR_{25-0} \ll 2$

seq:  $pc = pc + 4$

beq:  $pc = pc + offset * 4$

j :  $pc = pc_{31-28} + IR_{25-0} \ll 2$



# Five Execution Steps

- Instruction Fetch
- Instruction Decode and Register Fetch
- Execution, Memory Address Computation, or Branch Completion
- Memory Access or R-type instruction completion
- Write-back step

---

***INSTRUCTIONS TAKE FROM 3 - 5 CYCLES!***



# Step 1: Instruction

## Fetch

- Use PC to get instruction and put it in the Instruction Register.
- Increment the PC by 4 and put the result back in the PC.
- Can be described succinctly using RTL "Register-Transfer Language"

$IR = Memory[PC];$

$PC = PC + 4;$

Can we figure out the values of the control signals?

What is the advantage of updating the PC now?



# Step 2: Instruction

## Decode and Register Fetch

- Read registers `rs` and `rt` in case we need them
- Compute the branch address in case the instruction is a branch
- RTL:  
$$A = \text{Reg}[\text{IR}[25-21]];$$
$$B = \text{Reg}[\text{IR}[20-16]];$$
$$\text{ALUOut} = \text{PC} + (\text{sign-extend}(\text{IR}[15-0]) \ll 2);$$

We aren't setting any control lines based on the instruction type

(we are busy "decoding" it in our control logic)

# Step 3 (instruction dependent)

- ALU is performing one of three functions, based on instruction type
- Memory Reference ( lw / sw ):  
$$\text{ALUOut} = A + \text{sign-extend}(\text{IR}[15-0]);$$
- R-type:  
$$\text{ALUOut} = A \text{ op } B;$$
  
Branch:  
$$\text{if } (A == B) \text{ PC} = \text{ALUOut};$$
- jump:  
$$\text{pc} = \text{pc}_{31-28} + \text{IR}_{25-0} \ll 2$$



# Step 4 (R-type or memory access)

- Loads and stores access memory

$\text{MDR} = \text{Memory}[\text{ALUOut}] ;$

or

$\text{Memory}[\text{ALUOut}] = \text{B} ;$

- R-type instructions finish

$\text{Reg}[\text{IR}[15-11]] = \text{ALUOut} ;$

*The write actually takes place at the end of the cycle on the edge*



# Write-back step (step 5)



- `lw`

- `Reg[IR[20-16]] = MDR;`

*What about all the other instructions?*



# Summary :

Step name	Action for R-type instructions	Action for memory-reference instructions	Action for branches	Action for jumps
Instruction fetch	$IR = Memory[PC]$ $PC = PC + 4$			
Instruction decode/register fetch	$A = Reg [IR[25-21]]$ $B = Reg [IR[20-16]]$ $ALUOut = PC + (sign-extend (IR[15-0]) \ll 2)$			
Execution, address computation, branch/ jump completion	$ALUOut = A \text{ op } B$	$ALUOut = A + sign-extend (IR[15-0])$	if $(A == B)$ then $PC = ALUOut$	$PC = PC [31-28] \parallel (IR[25-0] \ll 2)$
Memory access or R-type completion	$Reg [IR[15-11]] = ALUOut$	Load: $MDR = Memory[ALUOut]$ or Store: $Memory [ALUOut] = B$		
Memory read completion		Load: $Reg[IR[20-16]] = MDR$		

# Simple Questions

- How many cycles will it take to execute this code?

```
lw $t2, 0($t3)
lw $t3, 4($t3) ←
beq $t2, $t3, Label #assume not
add $t5, $t2, $t3
sw $t5, 8($t3)
```

Label: ...

What is going on during the 8th cycle of execution?

- In what cycle does the actual addition of \$t2 and \$t3 takes place?

# Implementing the Control



- Value of control signals is dependent upon:
  - what instruction is being executed
  - which step is being performed
- Use the information were accumulated to specify a finite state machine
  - specify the finite state machine graphically, or
  - use microprogramming
- Implementation can be derived from ~~specification~~





# Computer Organization & Design

## 第07章: Pipeline

**楼学庆**

<http://10.214.47.99/>

[Email:hzlou@163.com](mailto:hzlou@163.com)



玉泉校区曹光彪东楼507室




08:51:21

浙江大学计算机学院


# Enhancing Performance with Pipelining

# Pipelining

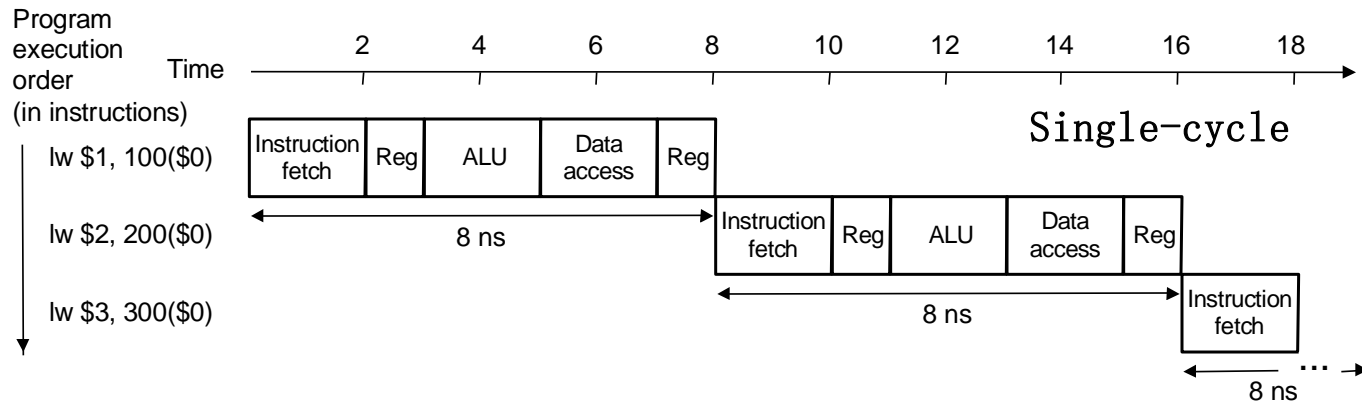
- Start work ASAP!! Do not waste time!



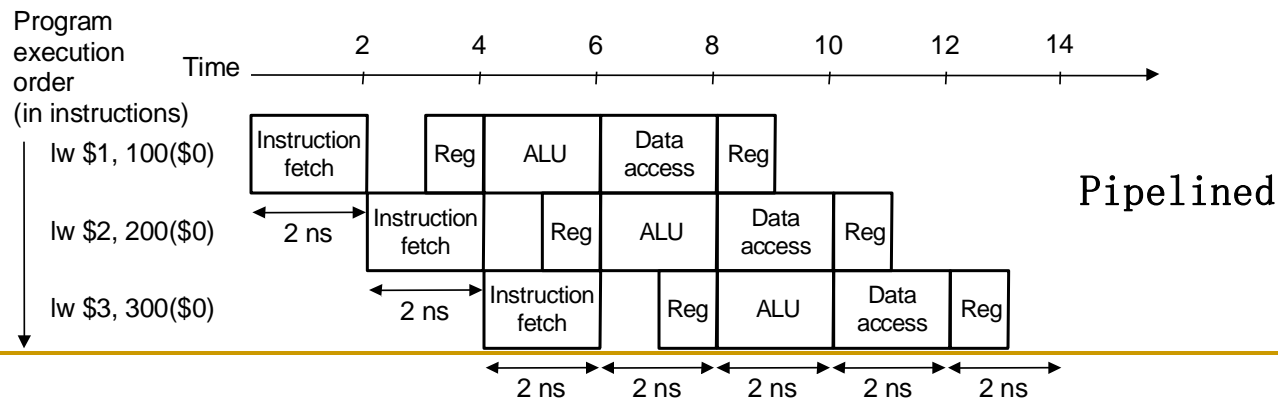
Assume 30 min. each task - wash, dry, fold, store - and that separate tasks use separate hardware and so can be overlapped



# Pipelined vs. Single-Cycle Instruction



Assume 2 ns for memory access, ALU operation; 1 ns for register access: therefore, single cycle clock 8 ns; pipelined clock cycle 2 ns.



# Pipelining: Keep in Mind

- Pipelining *does not* reduce latency of a single task, it *increases* throughput of entire workload
- Pipeline rate *limited* by longest stage
  - *potential* speedup = number pipe stages
  - *unbalanced* lengths of pipe stages reduces speedup
- Time to *fill* pipeline and time to *drain* it – when there is *slack* in the pipeline – reduces speedup



# Example Problem

- *Problem: for the laundry fill in the following table when*
  1. *the stage lengths are 30, 30, 30 30 min., resp.*
  2. *the stage lengths are 20, 20, 60, 20 min., resp.*

Person	Unpipelined finish time	Pipeline 1 finish time	Ratio unpipelined to pipeline 1
1			
2			
3			
4			
n			

- *Come up with a formula for pipeline speed-up!*

# Pipelining MIPS

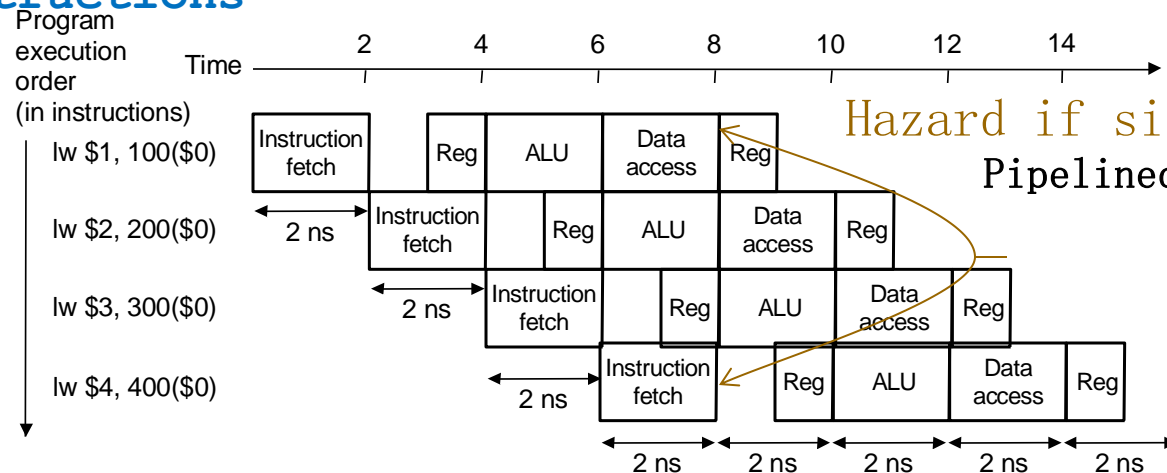
- What makes it easy with MIPS?
  - *all instructions are same length*
    - so fetch and decode stages are similar for all instructions
  - *just a few instruction formats*
    - simplifies instruction decode and makes it possible in one stage
  - *memory operands appear only in load/stores*
    - so memory access can be deferred to exactly one later stage
  - *operands are aligned in memory*
    - one data transfer instruction requires one memory access stage

# Pipelining MIPS

- What makes it hard?
  - *structural hazards*: different instructions, at different stages, in the pipeline want to use the same hardware resource
  - *control hazards*: succeeding instruction, to put into pipeline, depends on the outcome of a previous branch instruction, already in pipeline
  - *data hazards*: an instruction in the pipeline requires data to be computed by a previous instruction still in the pipeline
- Before actually building the pipelined datapath and control we first briefly examine these potential hazards individually...

# Structural Hazards

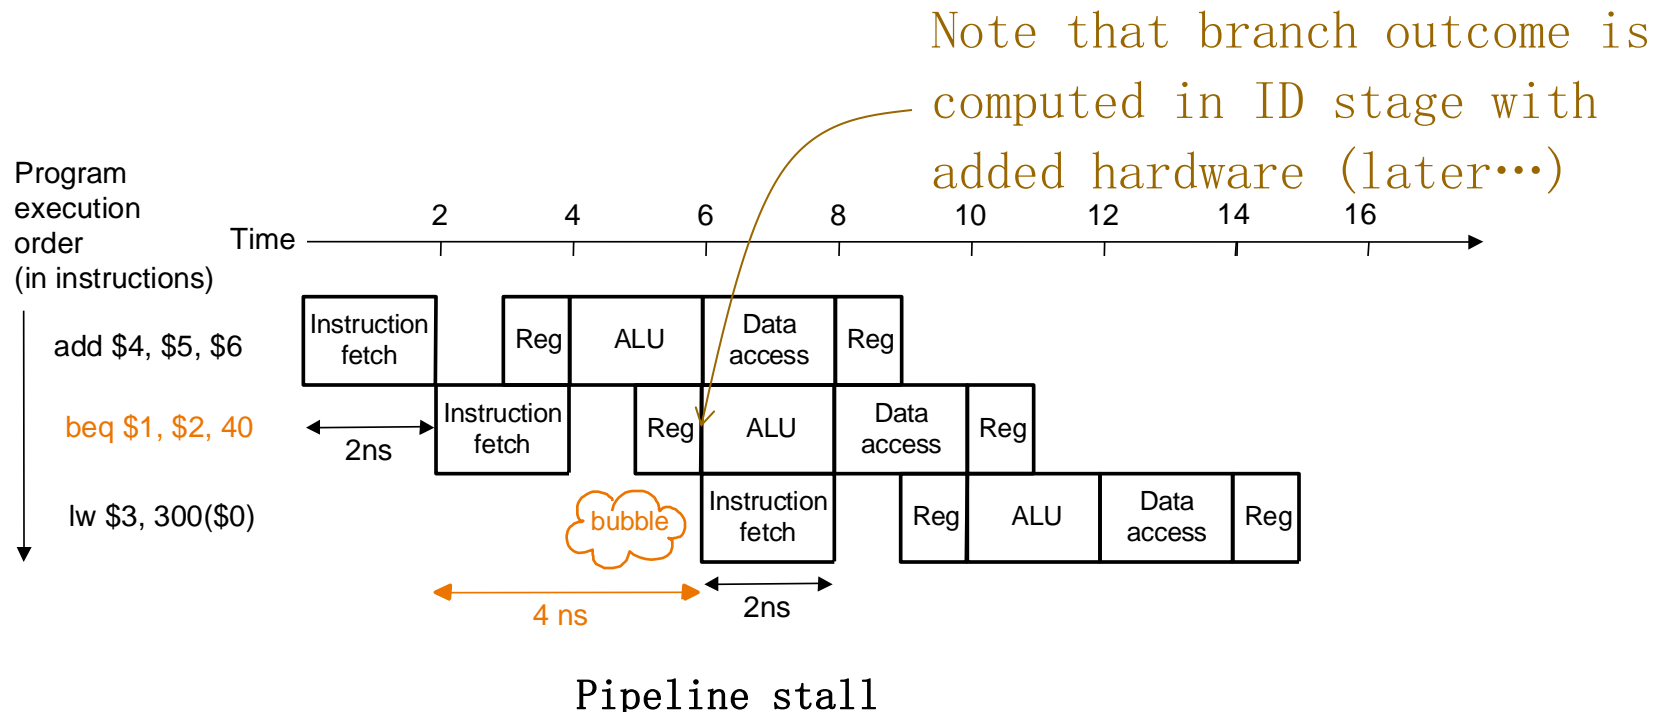
- **Structural hazard:** inadequate hardware to simultaneously support all instructions in the pipeline in the same clock cycle
- E.g., suppose *single - not separate -* instruction and data memory in pipeline below with *one read port*
  - then a structural hazard between first and fourth *lw* instructions



- **MIPS was designed to be pipelined: structural hazards are easy to avoid!**

# Control Hazards

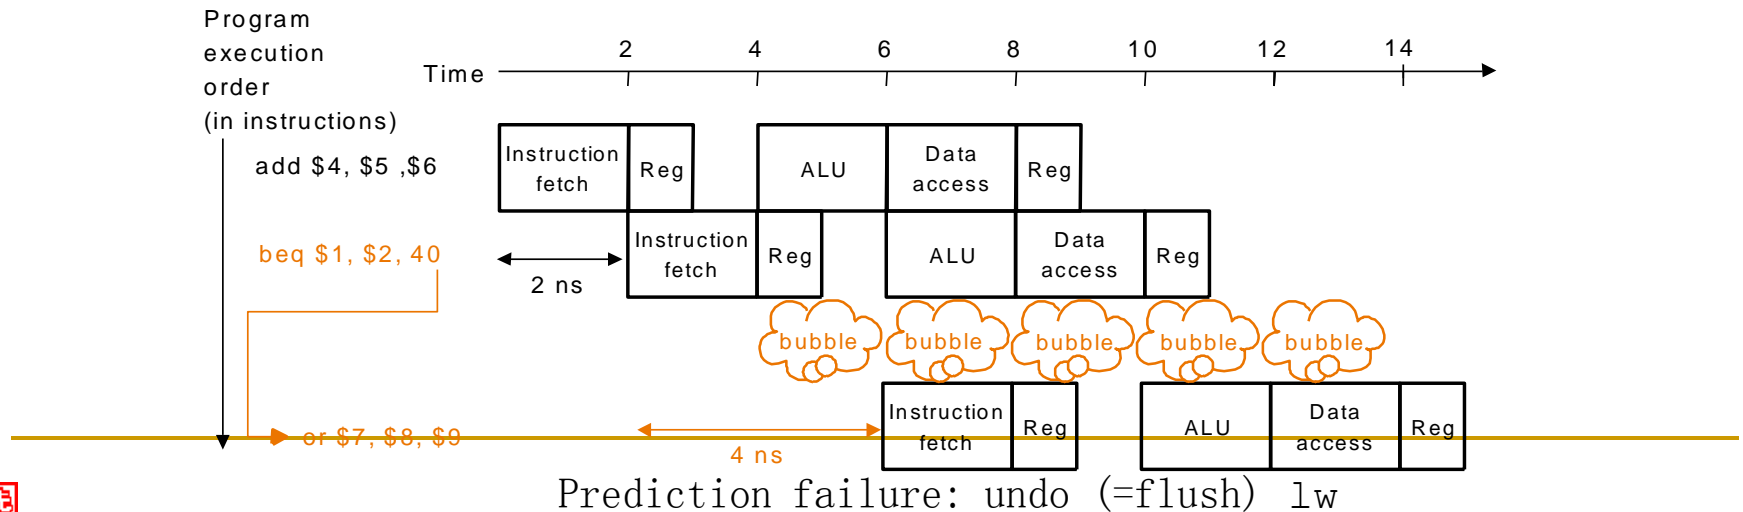
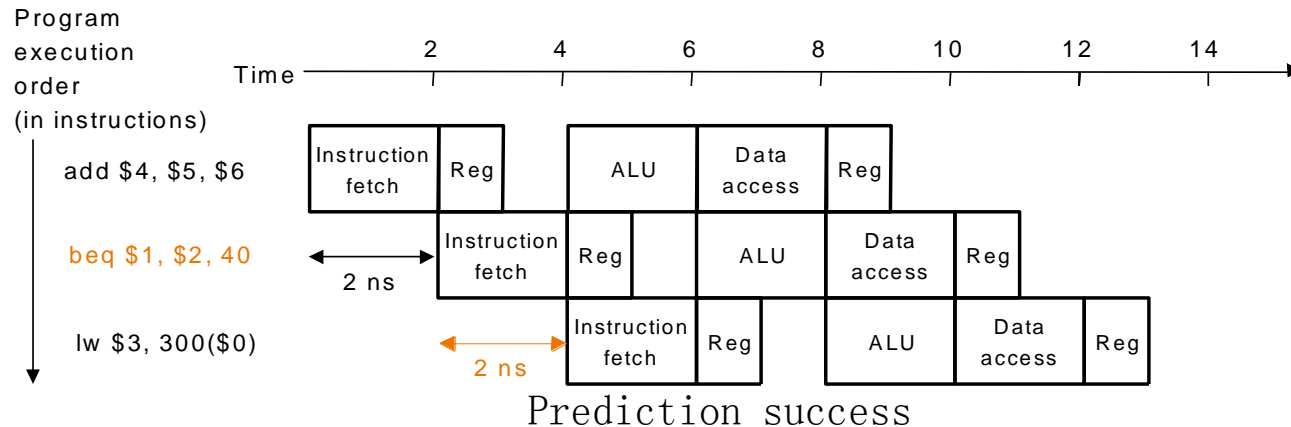
- *Control hazard*: need to make a decision based on the result of a previous instruction still executing in pipeline
- Solution 1 *Stall* the pipeline



# Control Hazards

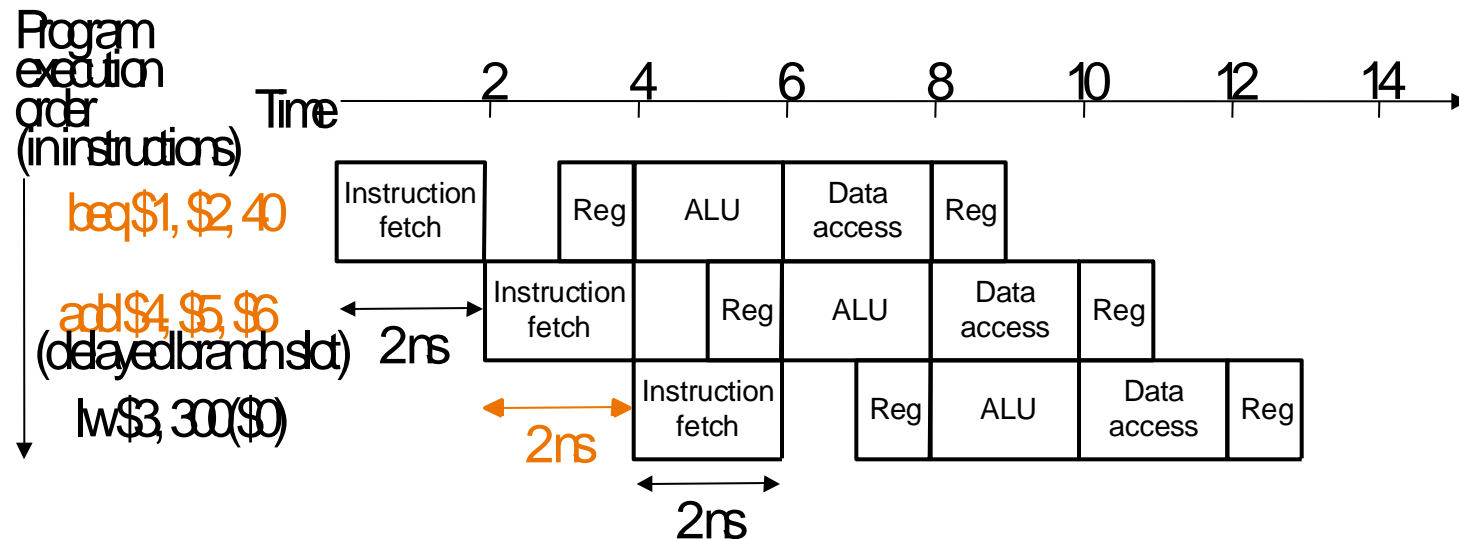
## ■ Solution 2 Predict branch outcome

□ e.g., predict *branch-not-taken* :



# Control Hazards

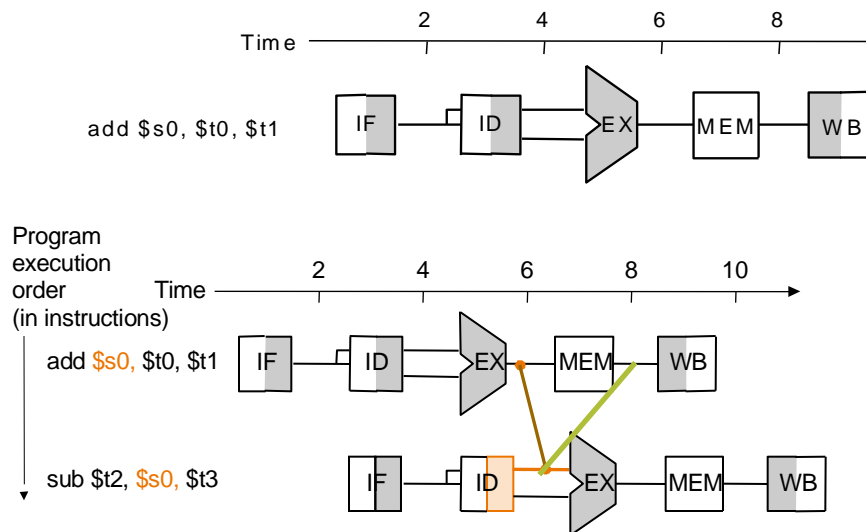
- Solution 3 Delayed branch: always execute the sequentially next statement with the branch executing after one instruction delay - compiler's job to find a statement that can be put in the slot that is independent of branch outcome
  - MIPS does this - but it is an option in SPIM (Simulator -> Settings)



Delayed branch `beq` is followed by `add` that is independent of branch outcome

# Data Hazards

- **Data hazard:** instruction needs data from the result of a previous instruction still executing in pipeline
- **Solution Forward data if possible...**



Instruction pipeline diagram:  
shade indicates use -  
left=write, right=read

Without forwarding - blue line -  
data has to go back in time;  
with forwarding - red line -

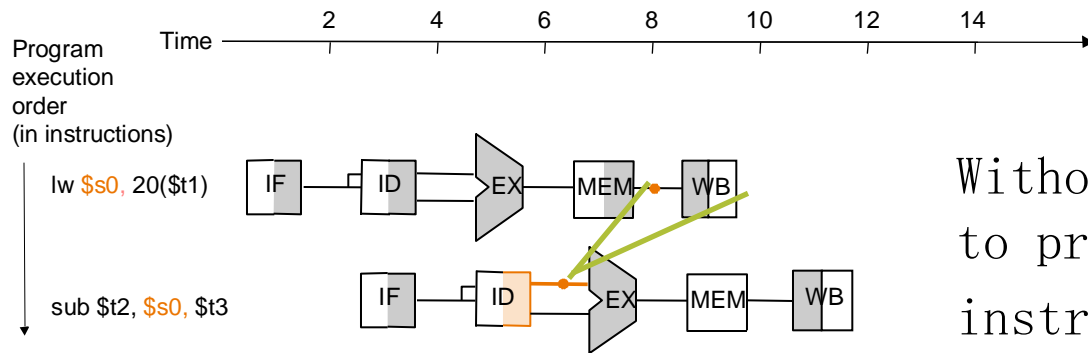
data is available in time



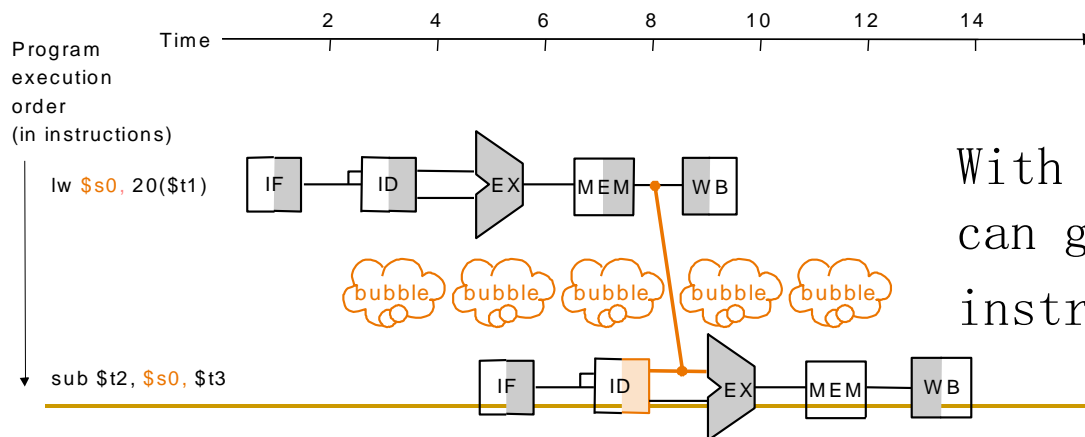
# Data Hazards

## ■ Forwarding may not be enough

- e.g., if an R-type instruction following a load uses the result of the load – called *load-use data hazard*



Without a stall it is impossible to provide input to the sub instruction in time



With a one-stage stall, forwarding can get the data to the sub instruction in time



# Reordering Code to Avoid Pipeline Stall (Software Solution)

## ■ Example:

```
lw $t0, 0($t1)
lw $t2, 4($t1)
sw $t2, 0($t1)
sw $t0, 4($t1)
```

← Data hazard

## ■ Reordered code:

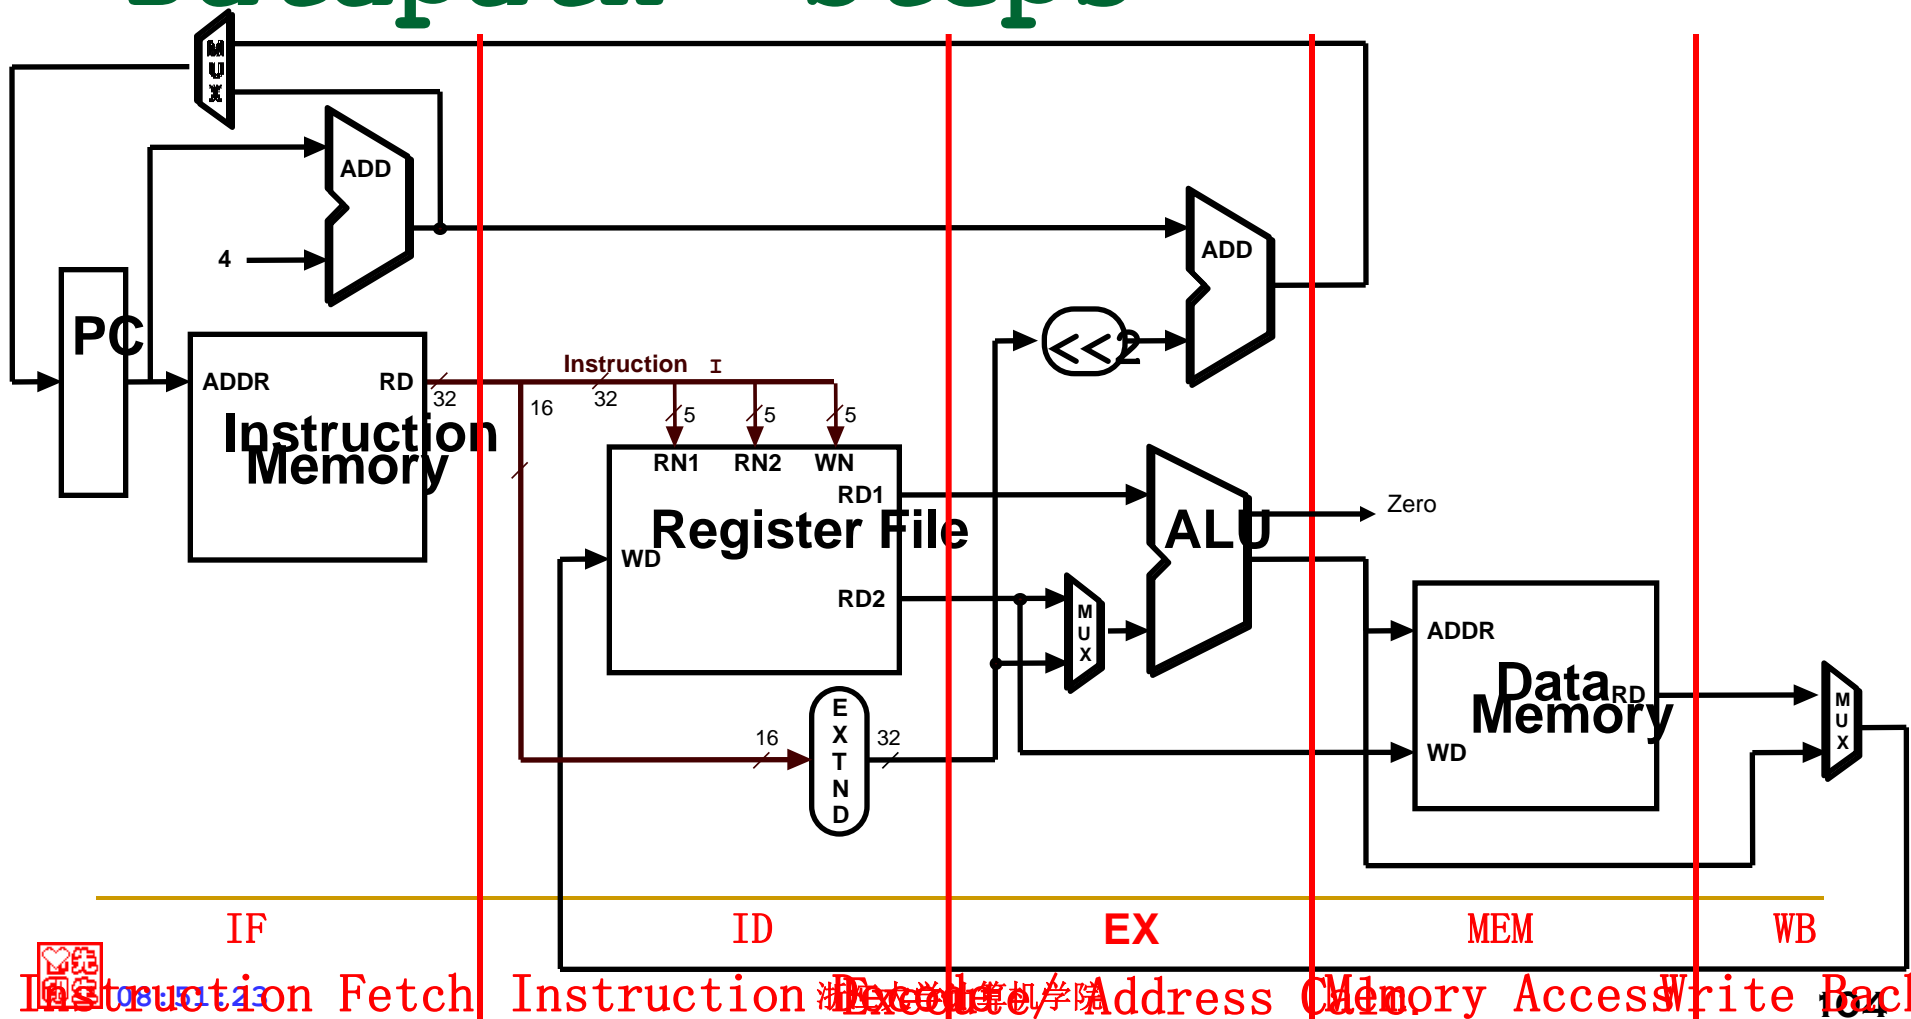
```
lw $t0, 0($t1)
lw $t2, 4($t1)
sw $t0, 4($t1)
sw $t2, 0($t1)
```

← Interchanged

# Pipelined Datapath

- We now move to actually building a pipelined datapath
- First recall the 5 steps in instruction execution
  1. Instruction Fetch & PC Increment (IF)
  2. Instruction Decode and Register Read (ID)
  3. Execution or calculate address (EX)
  4. Memory access (MEM)
  5. Write result into register (WB)
- Review: single-cycle processor
  - all 5 steps done in a single clock cycle
  - dedicated hardware required for each step
- *What happens if we break the execution into multiple cycles, but keep the extra hardware?*

# Review - Single-Cycle Datapath "Steps"

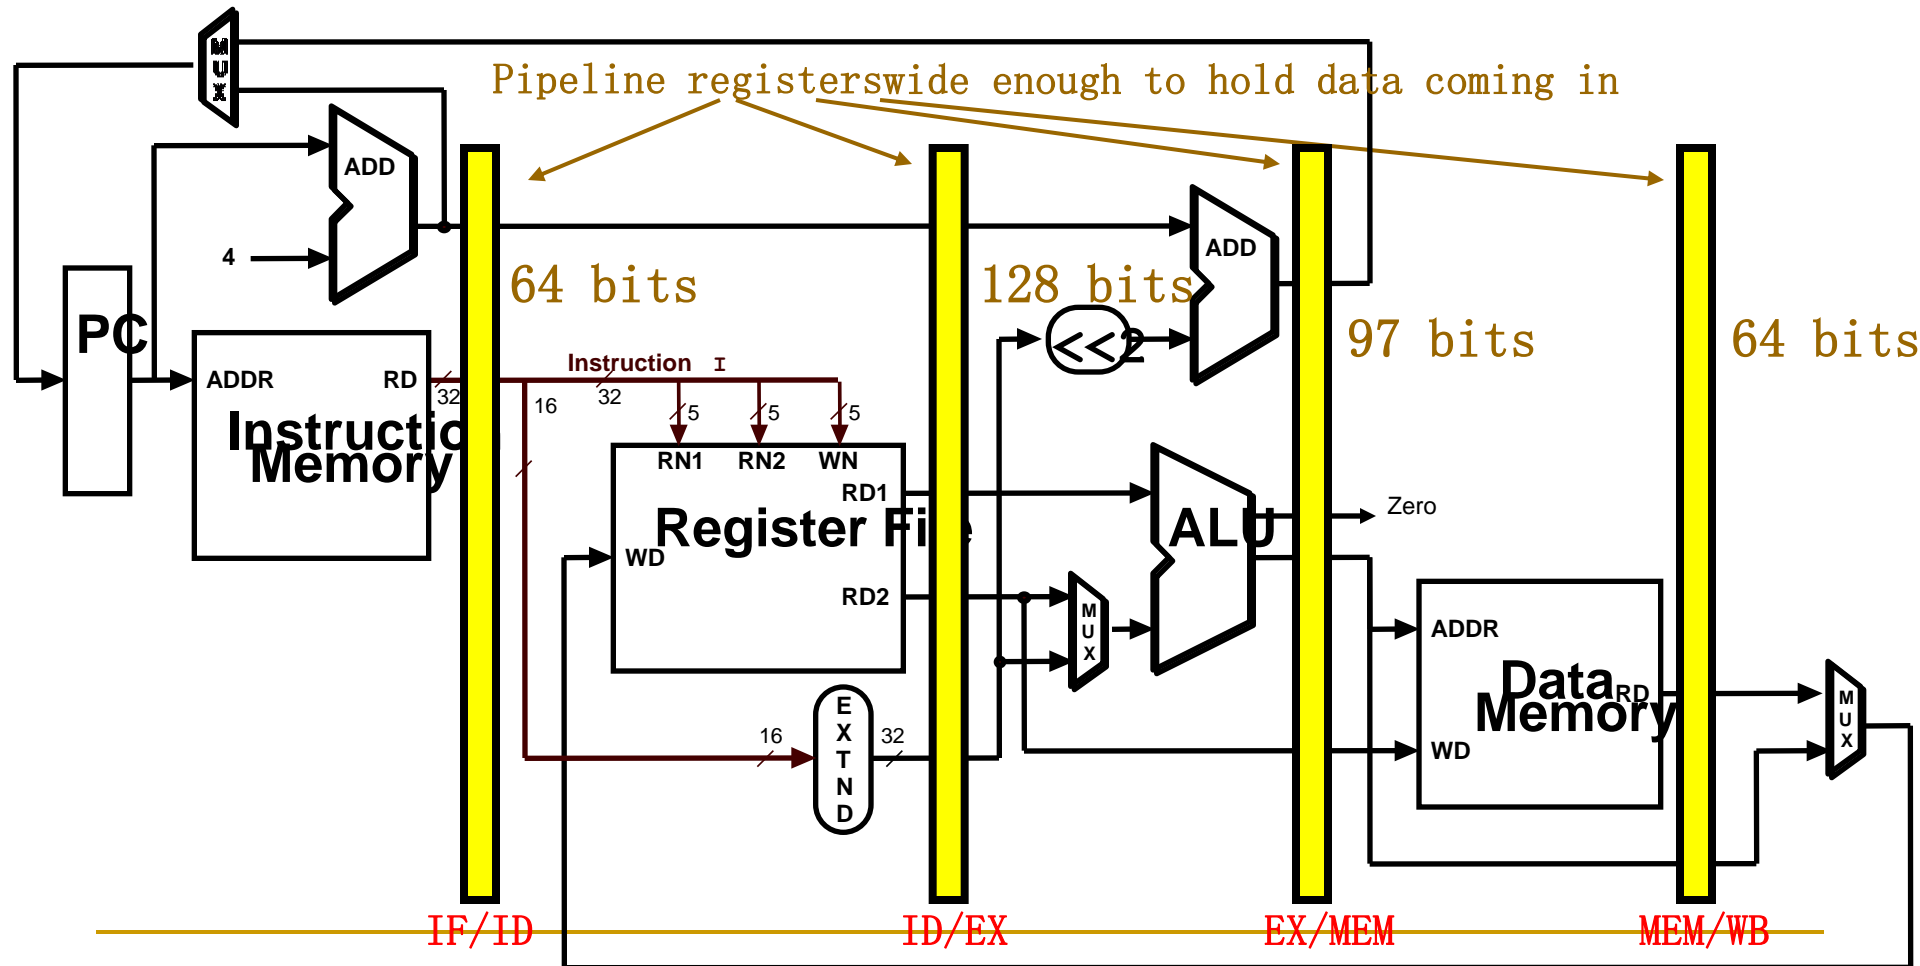




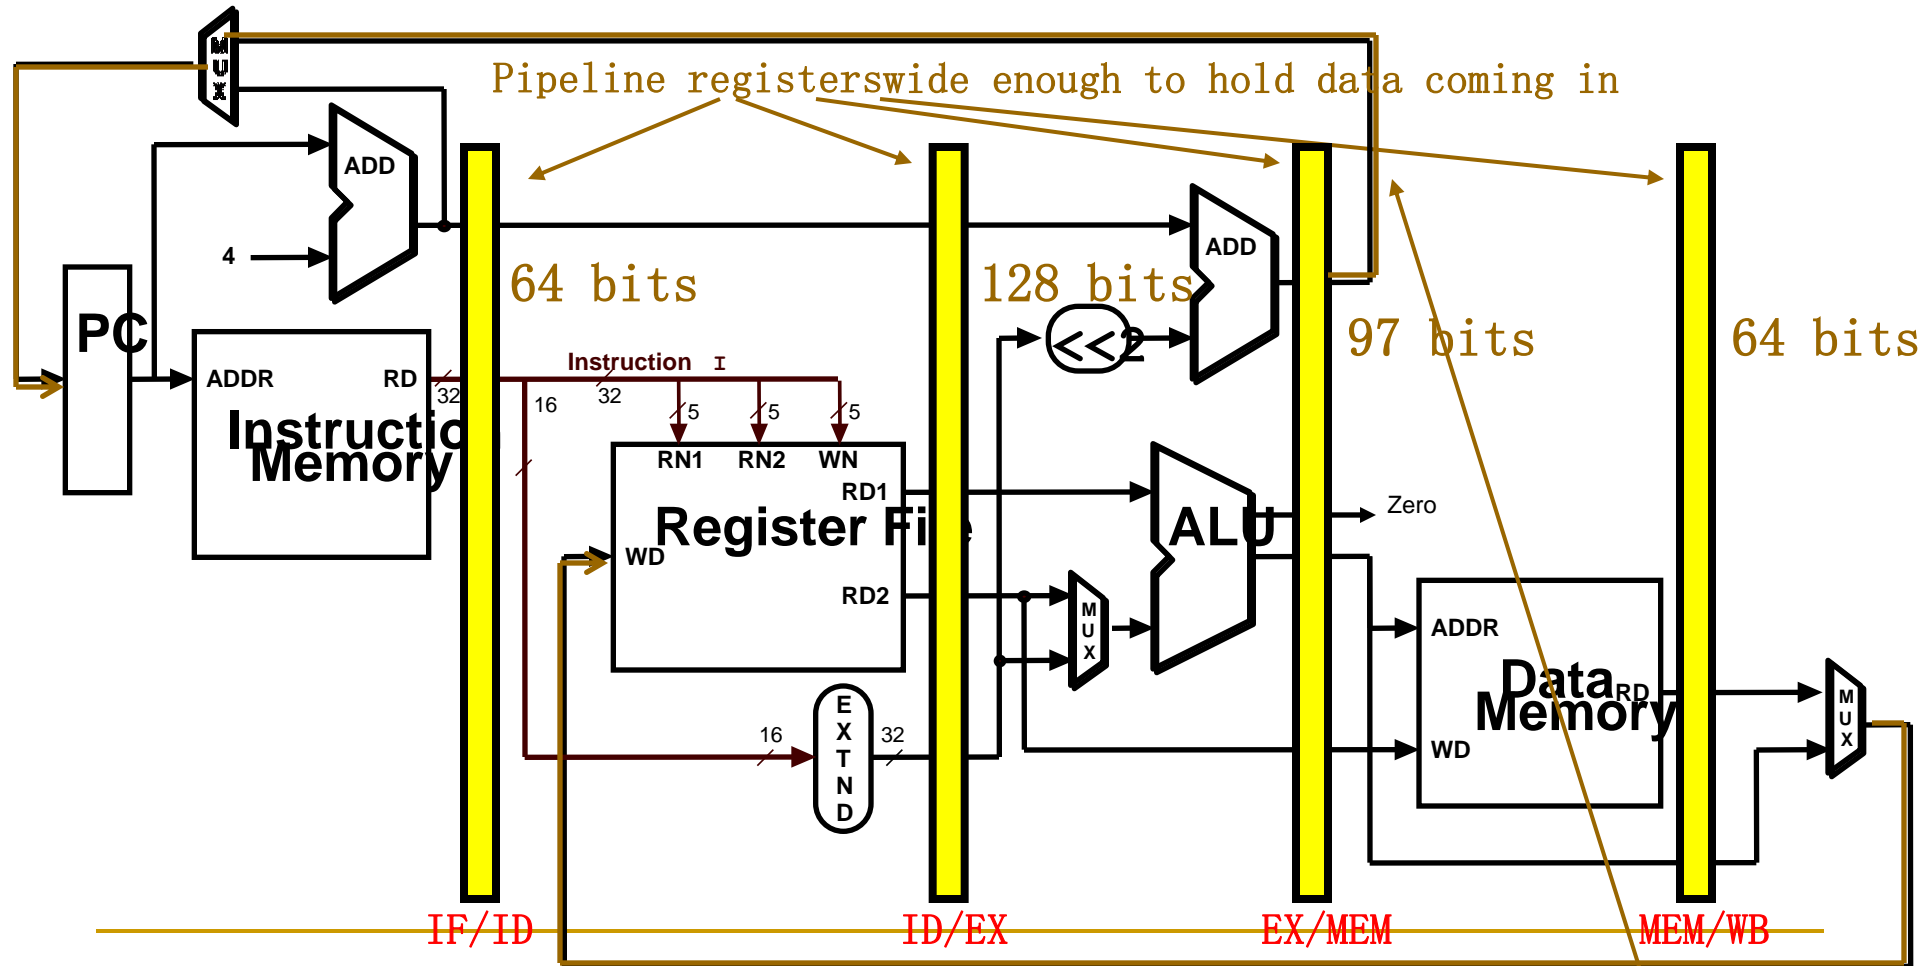
# Pipelined Datapath – Key Idea

- *What happens if we break the execution into multiple cycles, but keep the extra hardware?*
  - *Answer: We may be able to start executing a new instruction at each clock cycle - pipelining*
- ...but we shall need extra registers to hold data between cycles – *pipeline registers*

# Pipelined Datapath

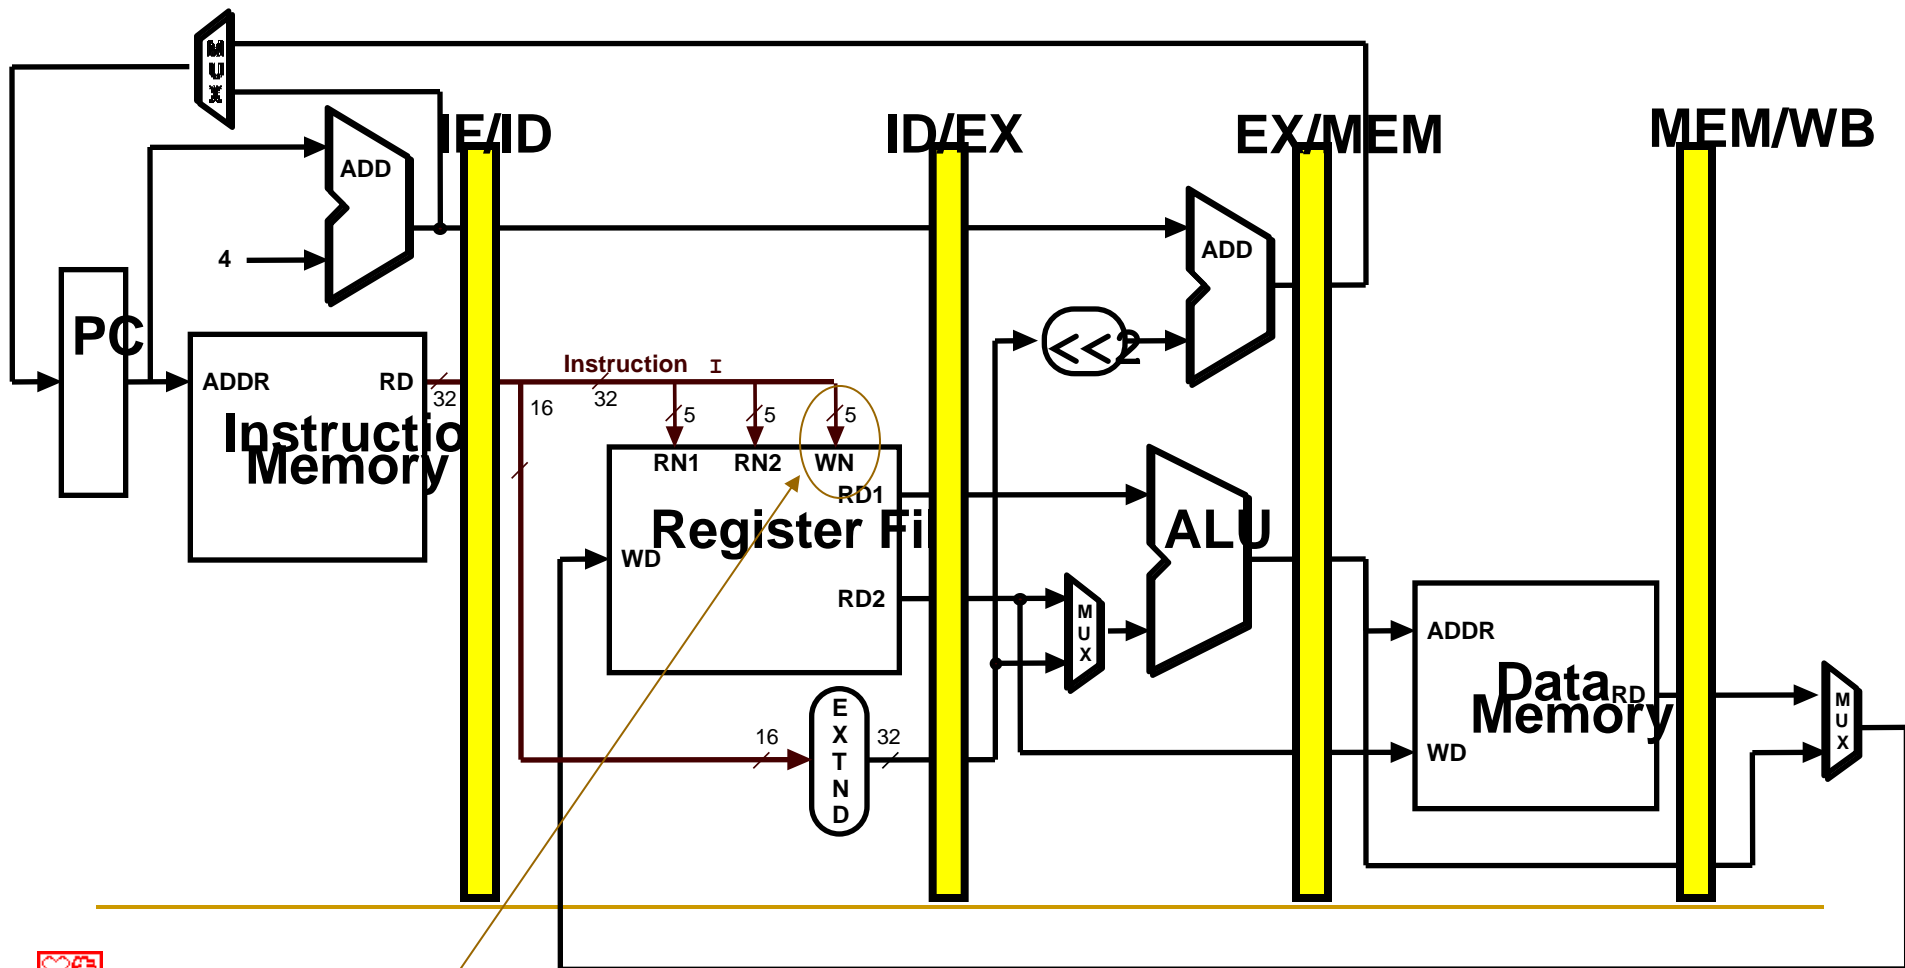


# Pipelined Datapath



Only data flowing right to left may cause hazard..., why?

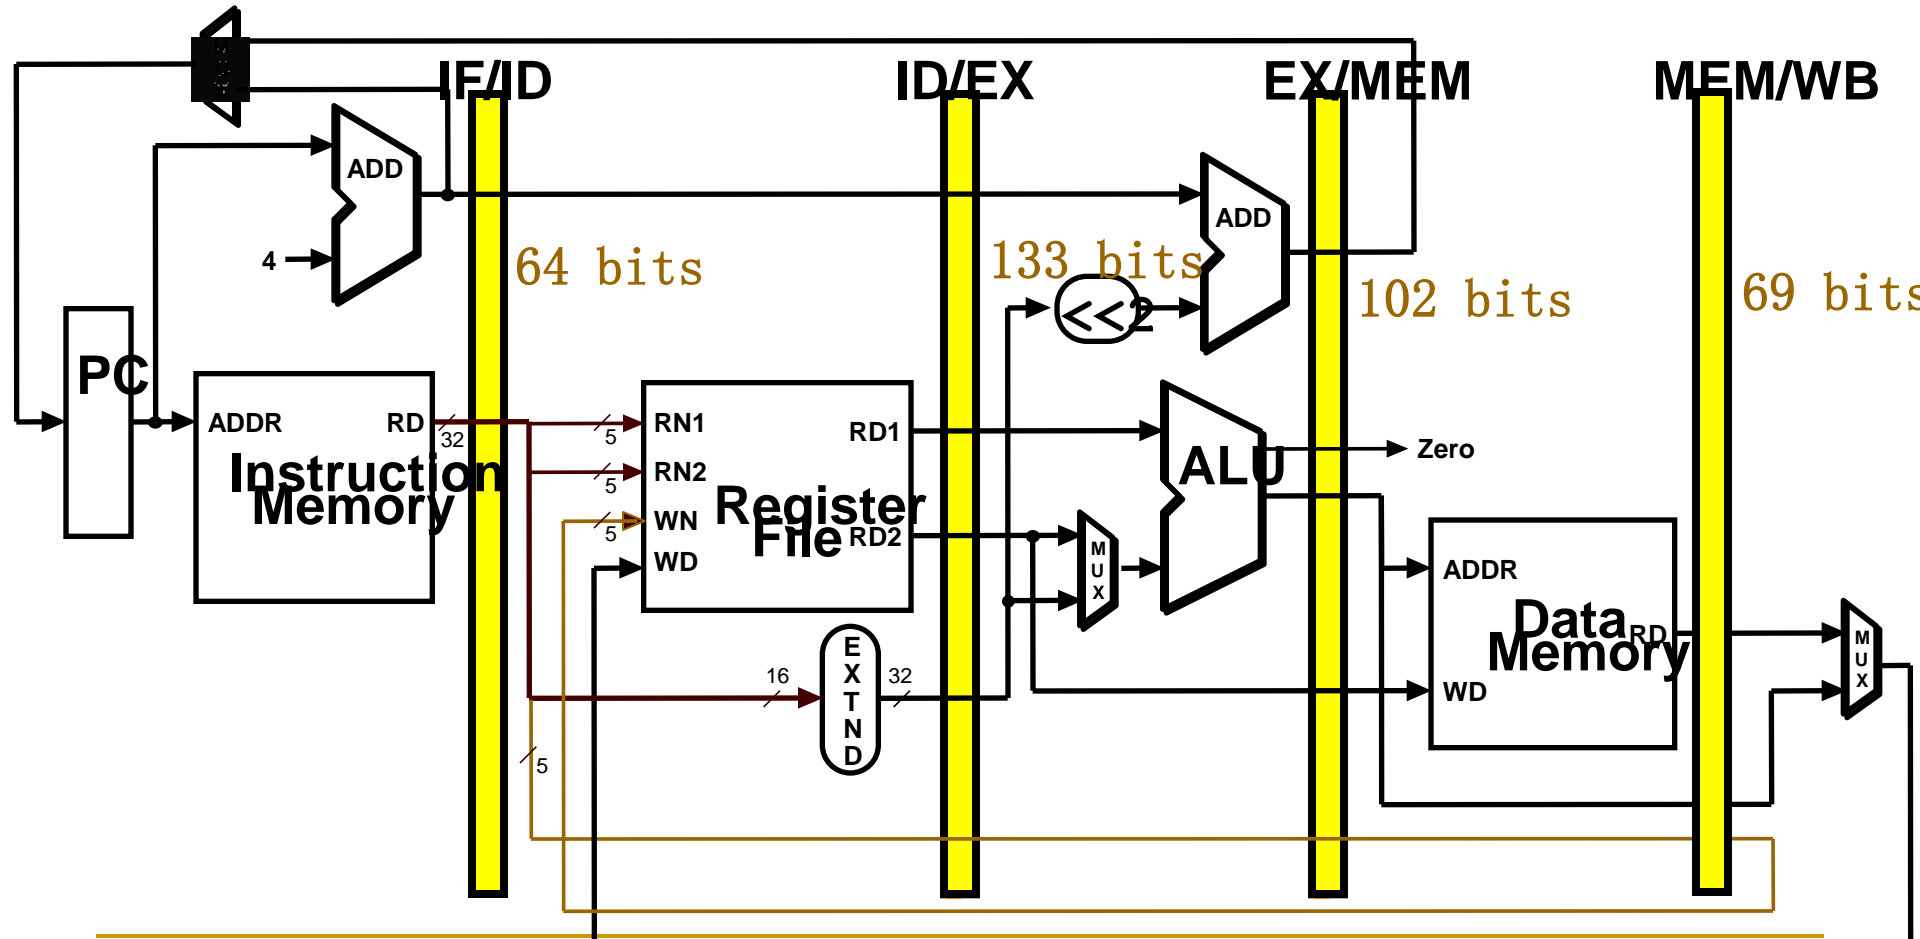
# Bug in the Datapath



Write register number comes from *later* instruction!



# Corrected Datapath



Destination register number is also passed through ID/EX, EX/MEM and MEM/WB registers, which are now wider by 5 bits

# Pipelined Example

- Consider the following instruction sequence:

```
lw  $t0, 10($t1)
```

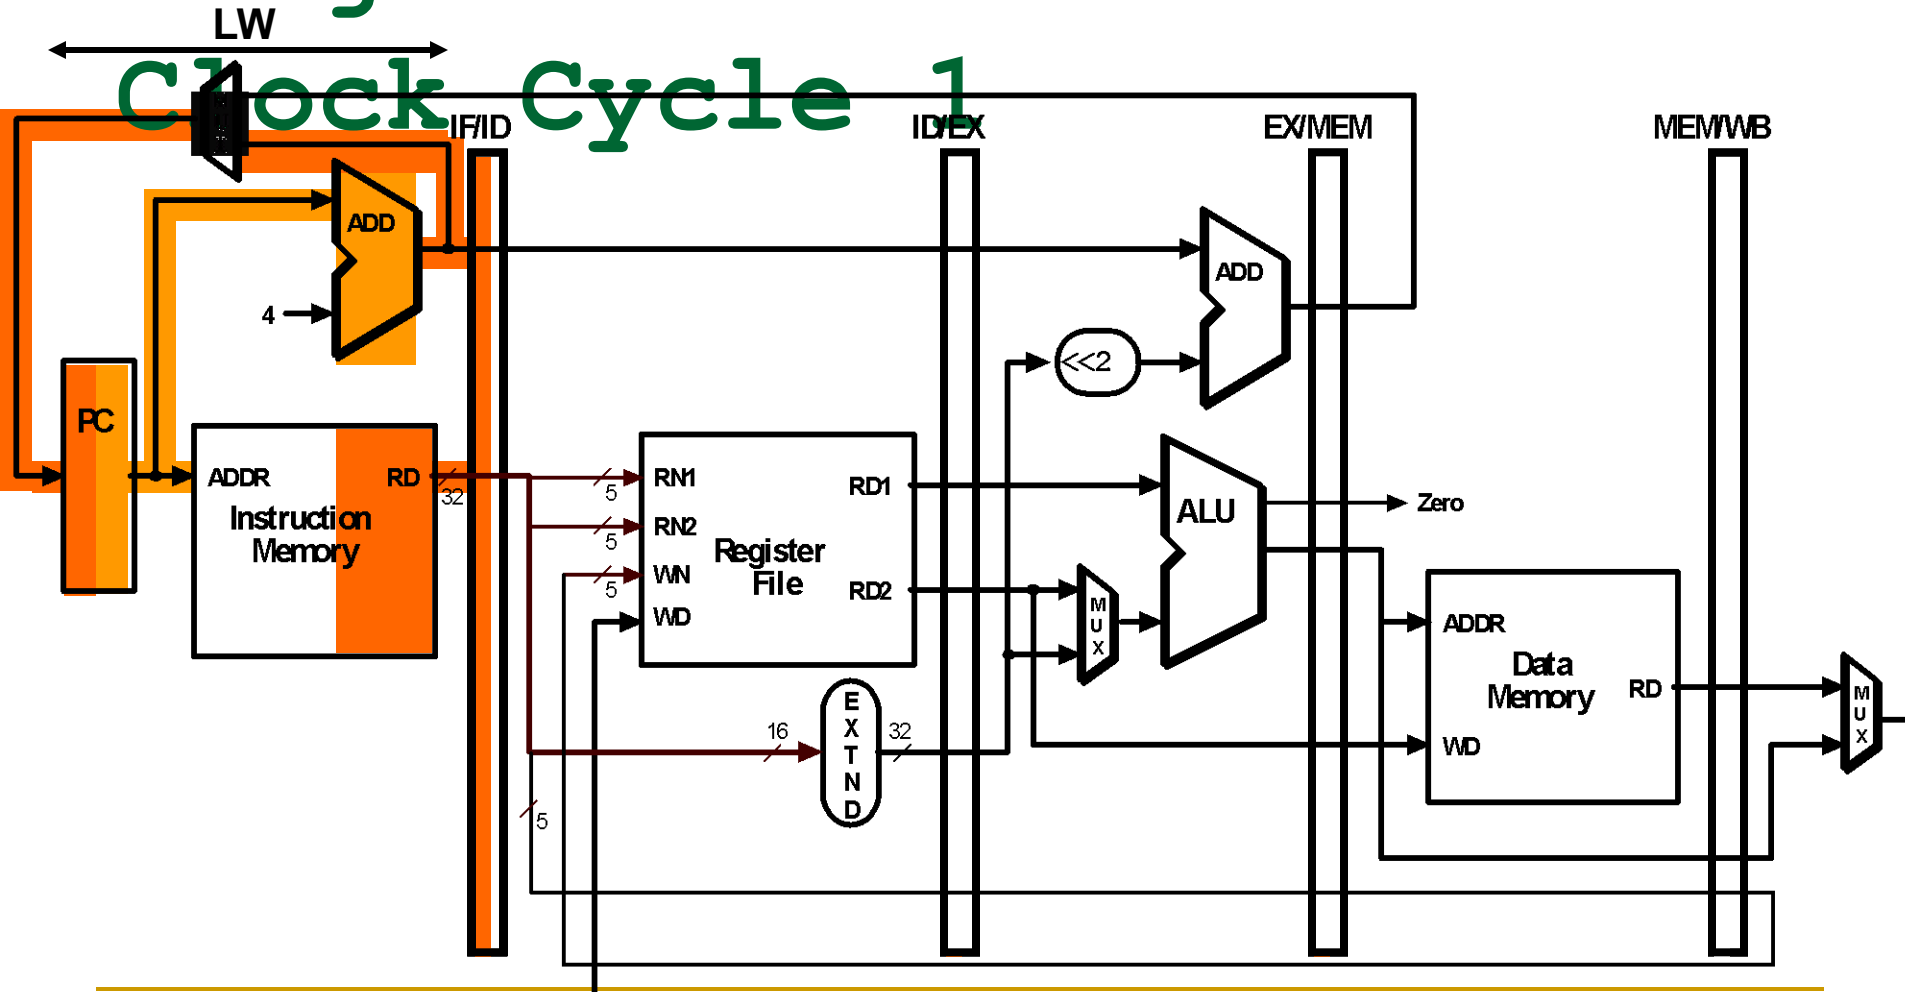
```
sw  $t3, 20($t4)
```

```
add $t5, $t6, $t7
```

```
sub $t8, $t9, $t10
```

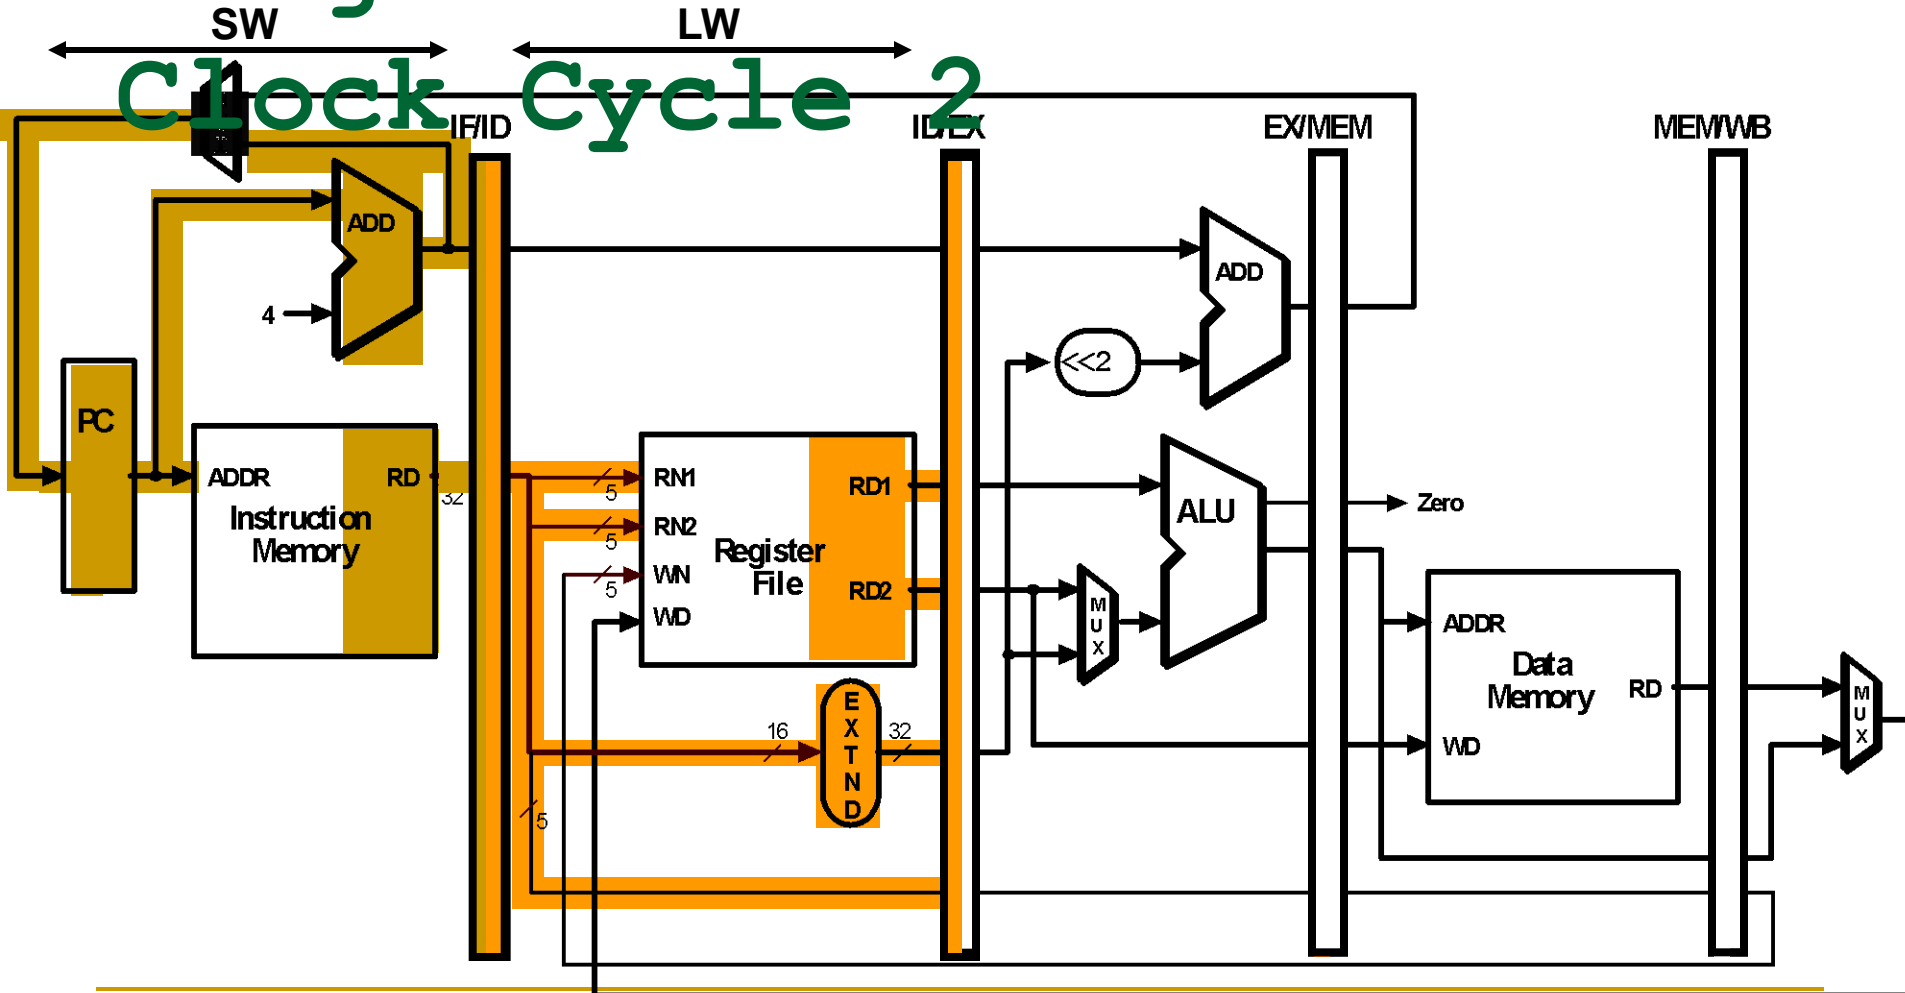
# Single-Clock-Cycle

## Diagram:



# Single-Clock-Cycle

## Diagram:

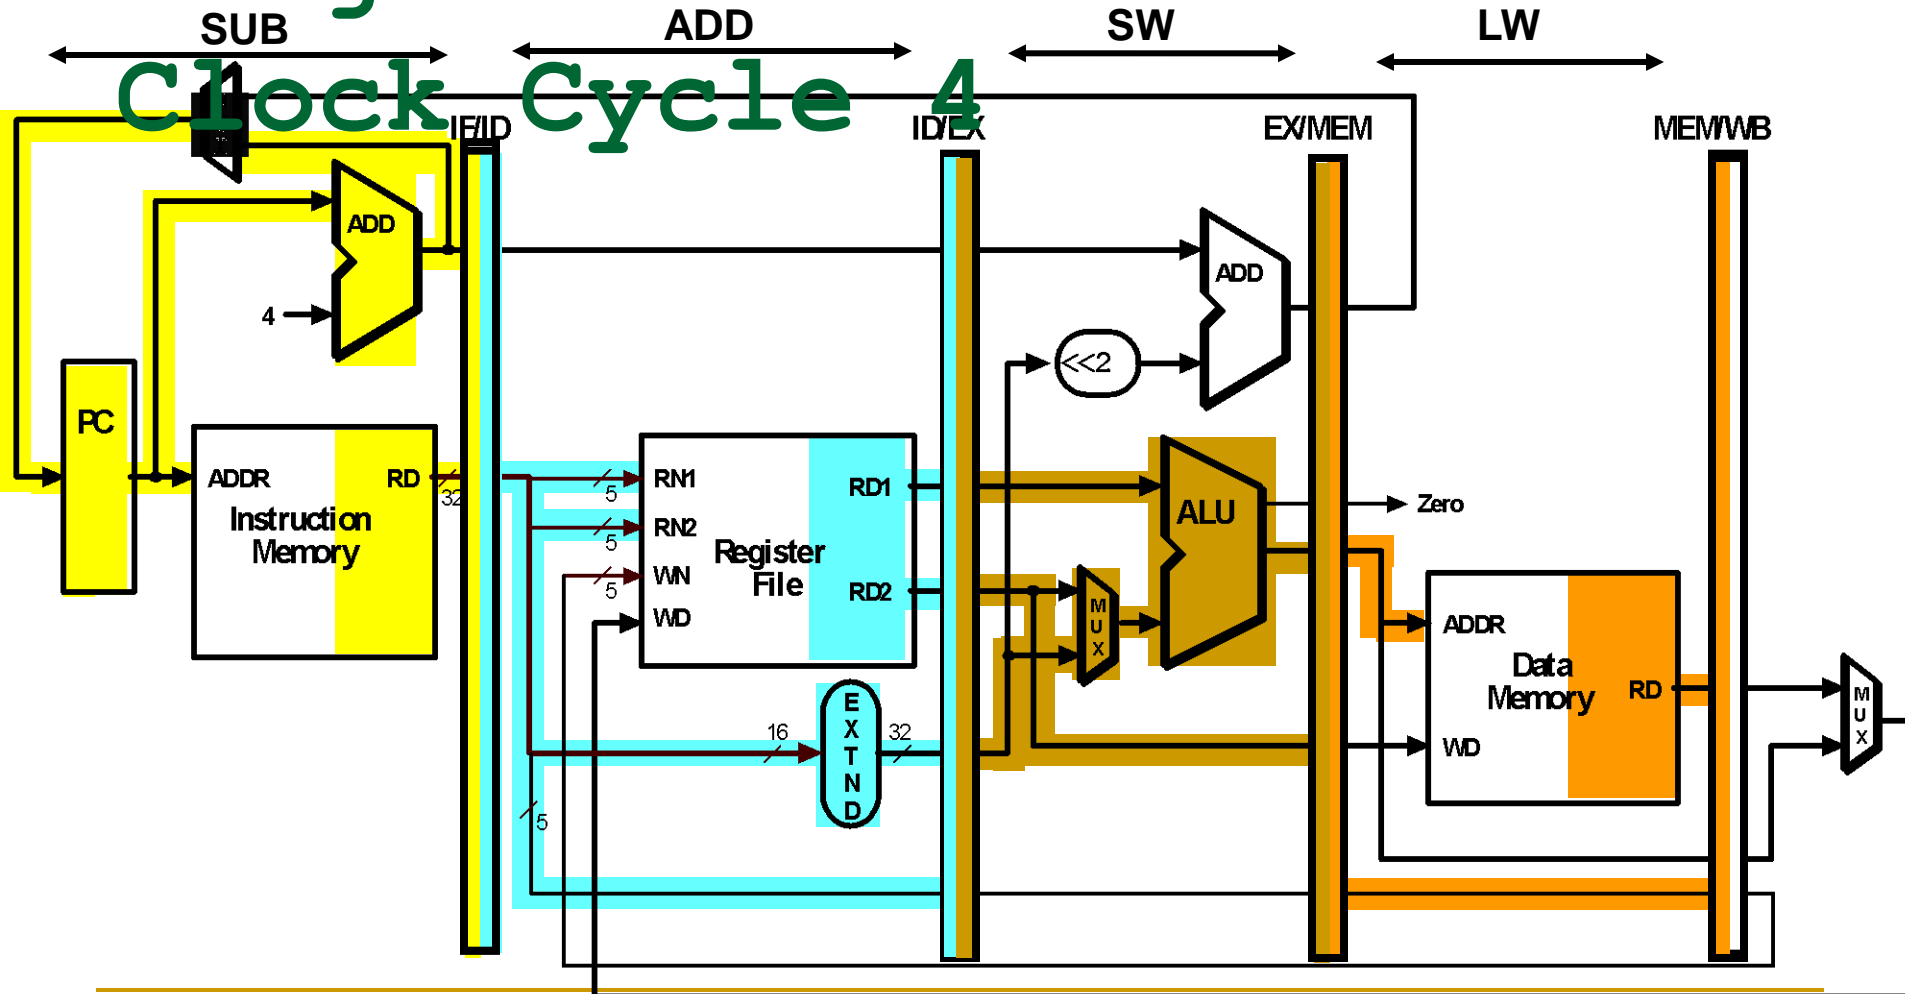


## Diagram:

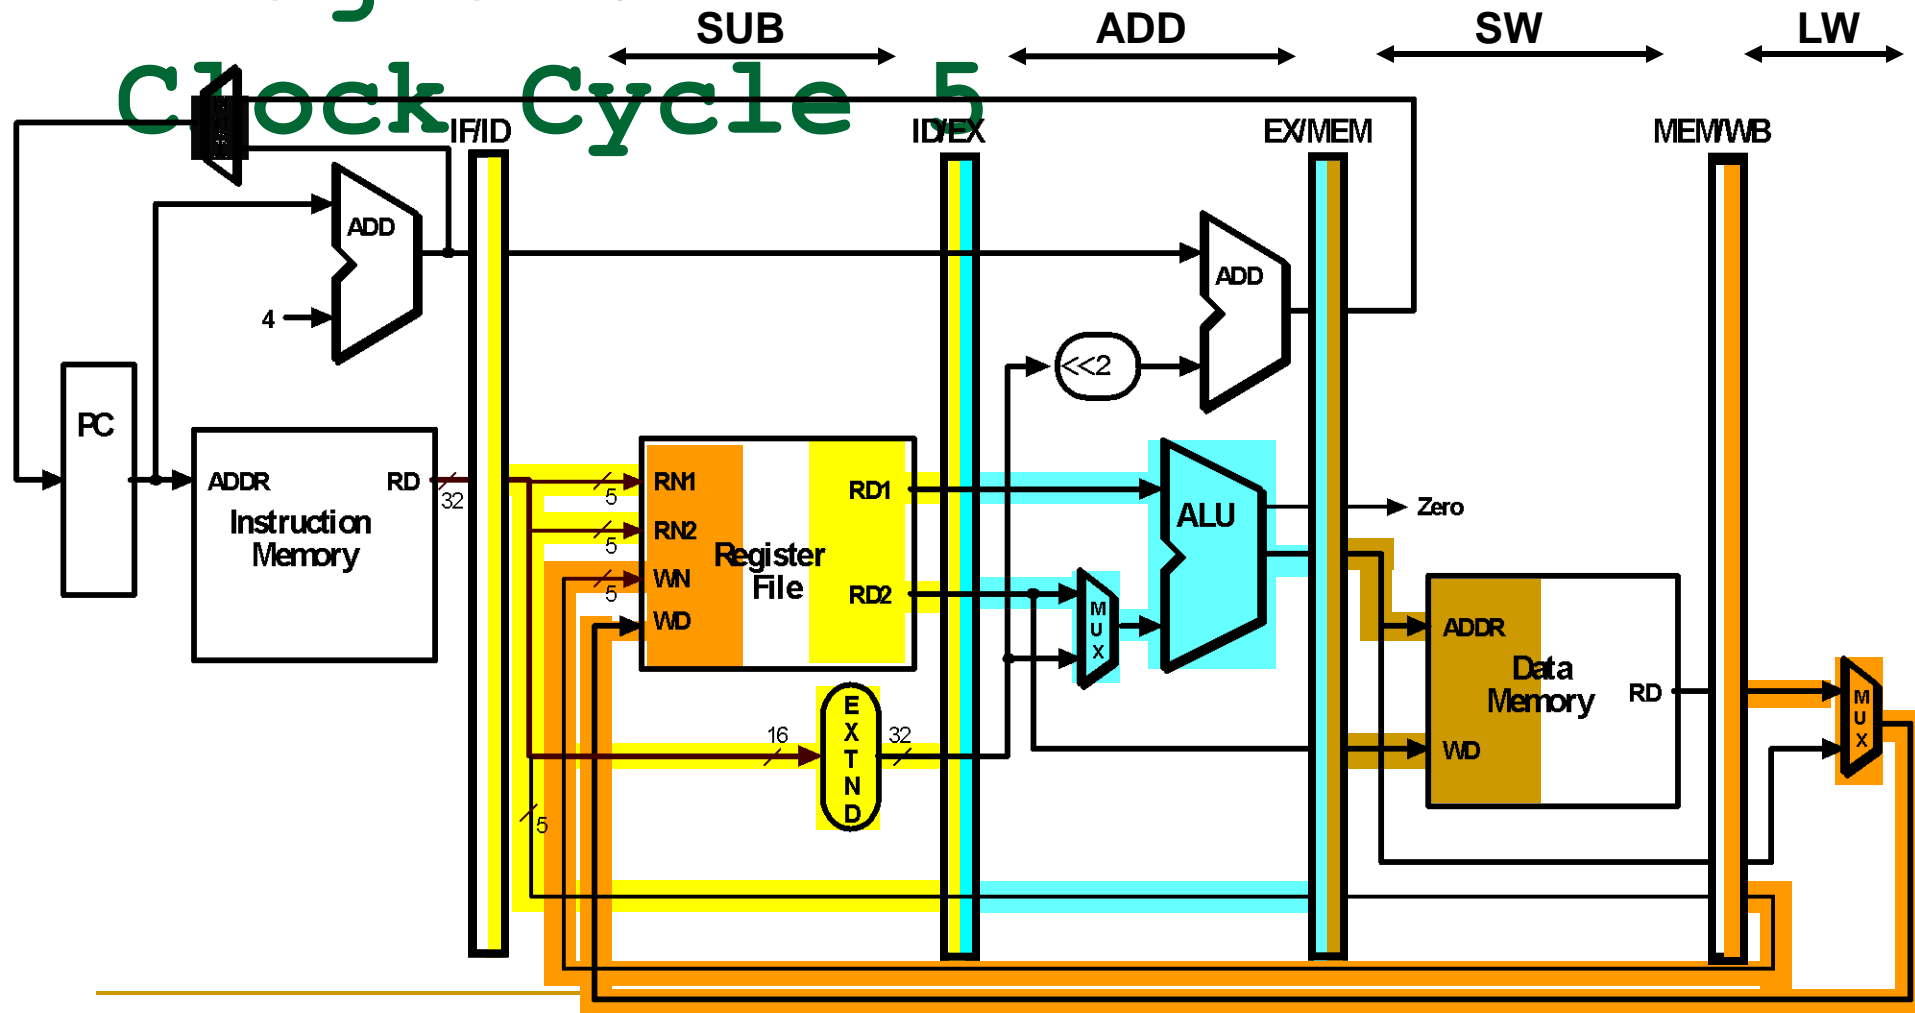
The diagram illustrates the internal state of a MIPS processor during Clock Cycle 3. The stages are labeled ADD, SW, and LW. The Register File contains values for RN1, RN2, VN, and WD. The ALU output is Zero. The Data Memory output is 0x00000000. The final MUX output is 0x00000000.

# Single-Clock-Cycle

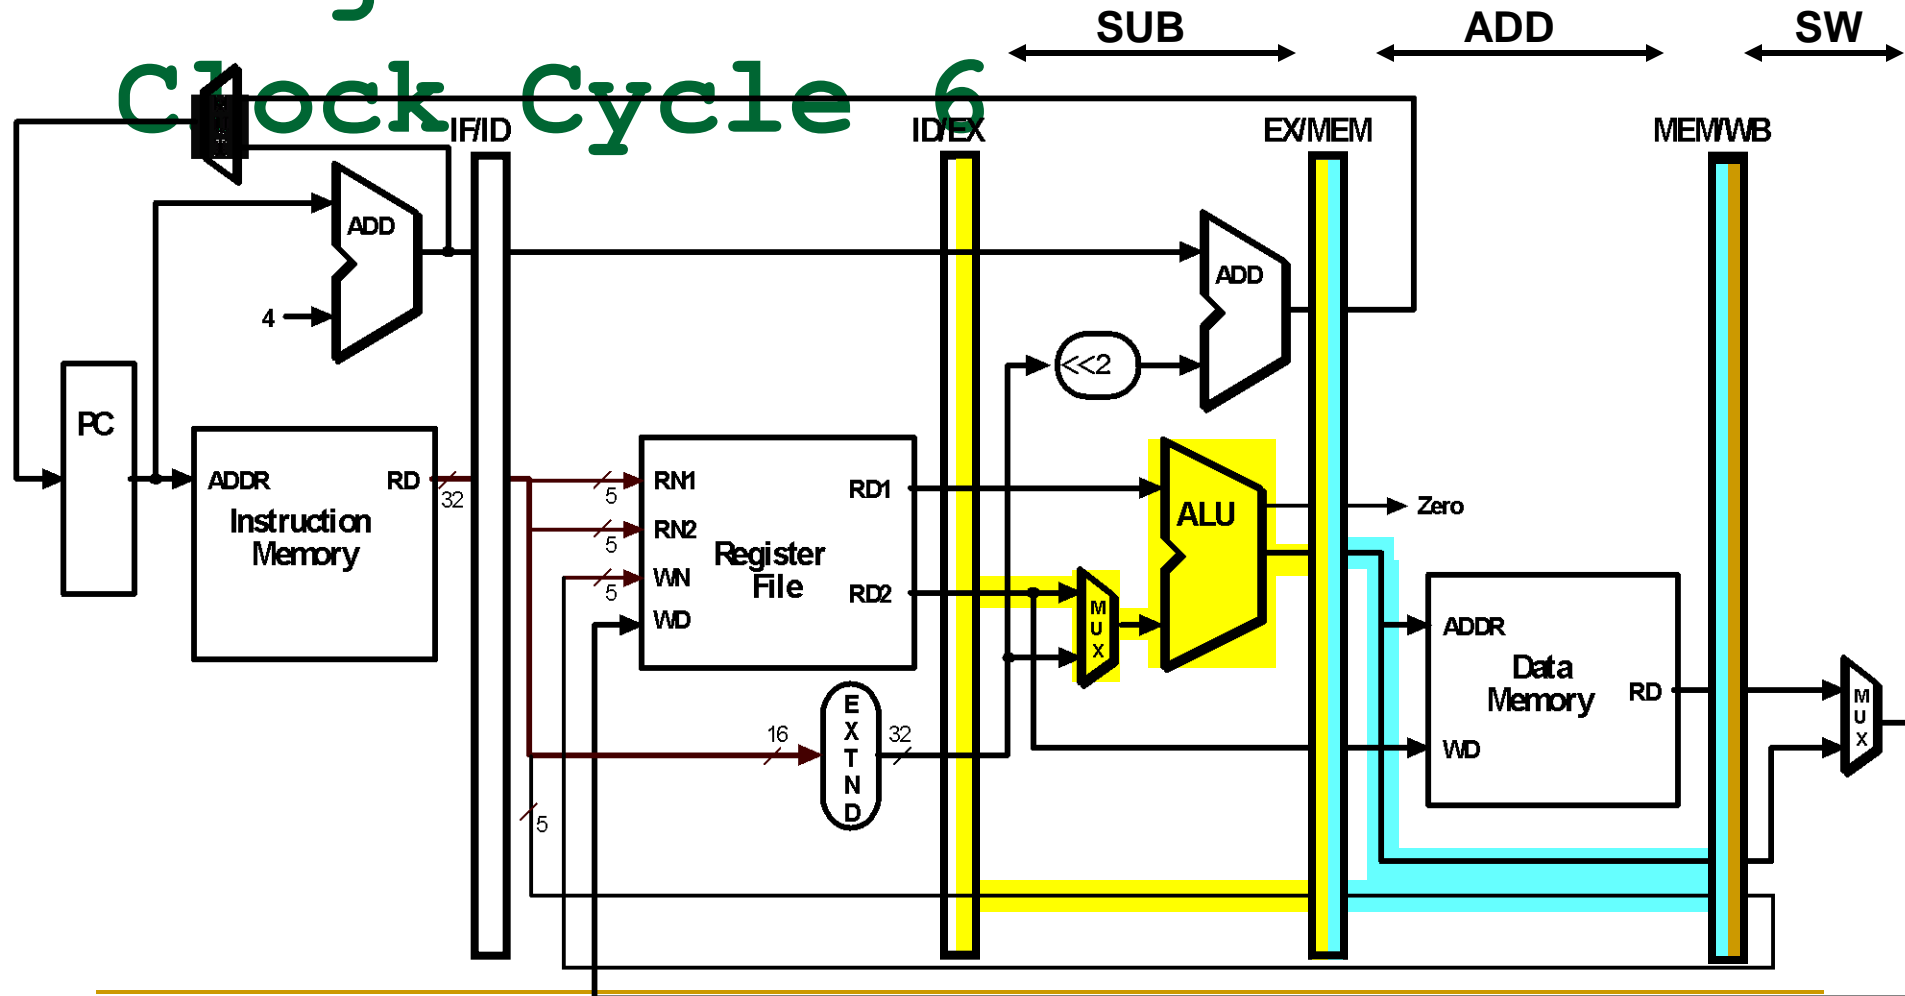
## Diagram:



# Single-Clock-Cycle Diagram: Clock Cycle 5



# Single-Clock-Cycle Diagram: Clock Cycle 6

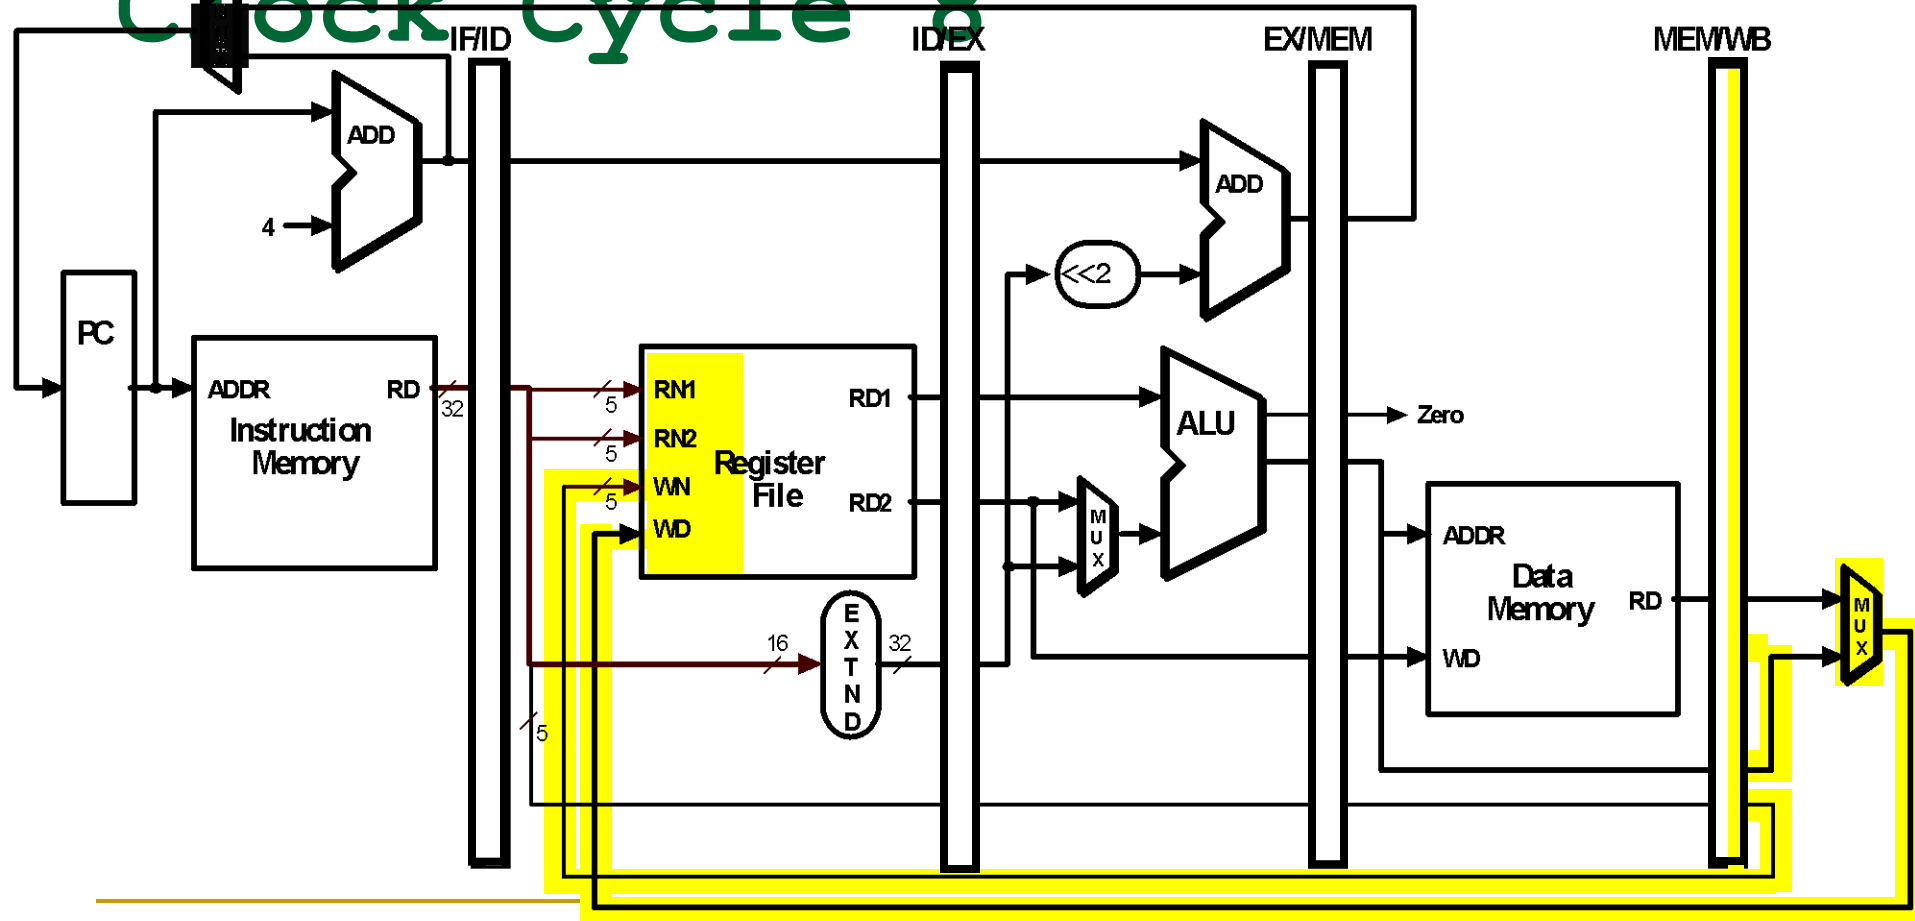




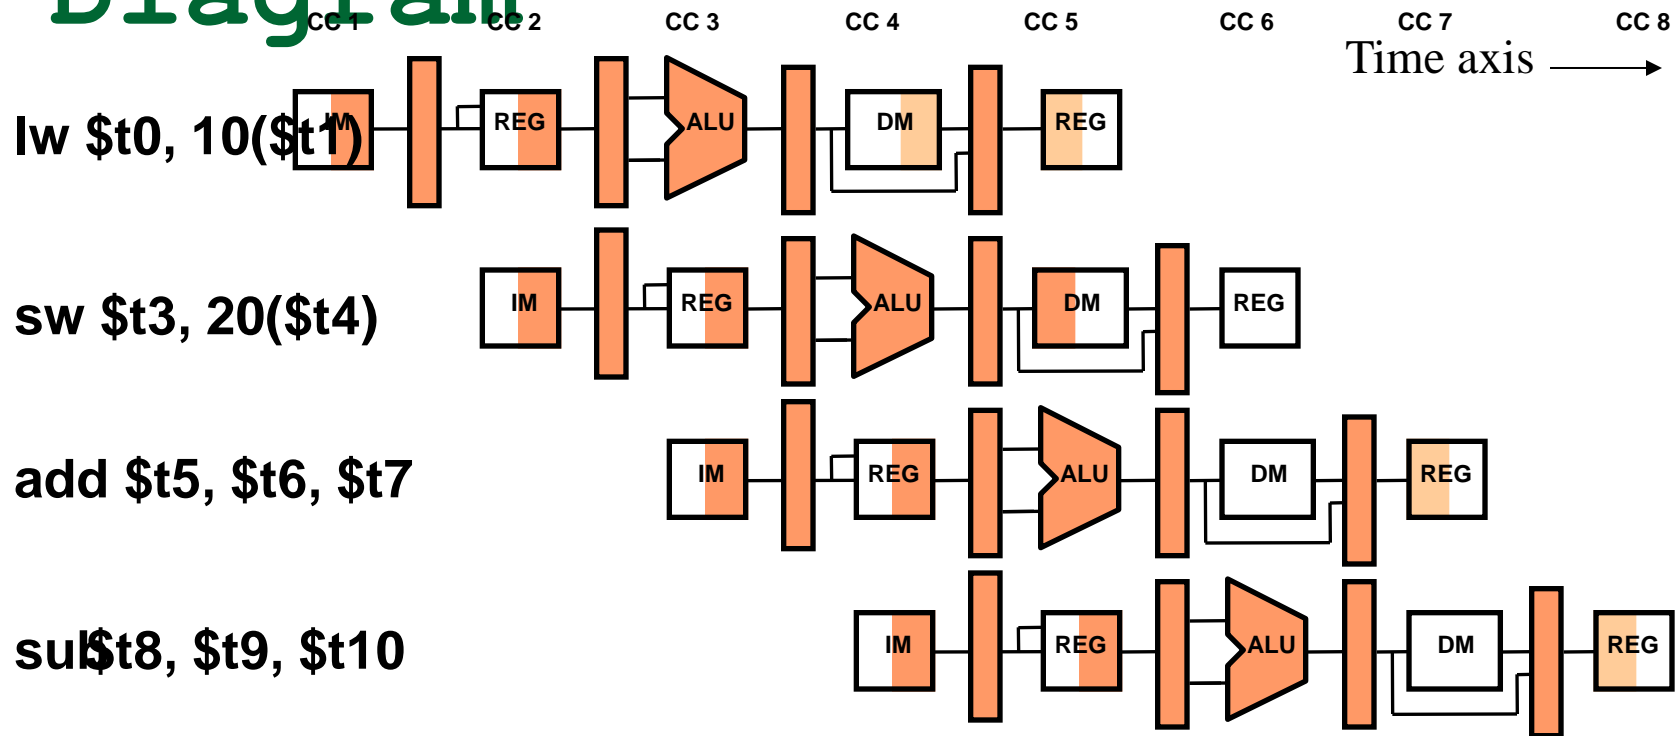


# Single-Clock-Cycle Diagram: Clock Cycle 8

SUB  
↔



# Alternative View – Multiple-Clock-Cycle Diagram



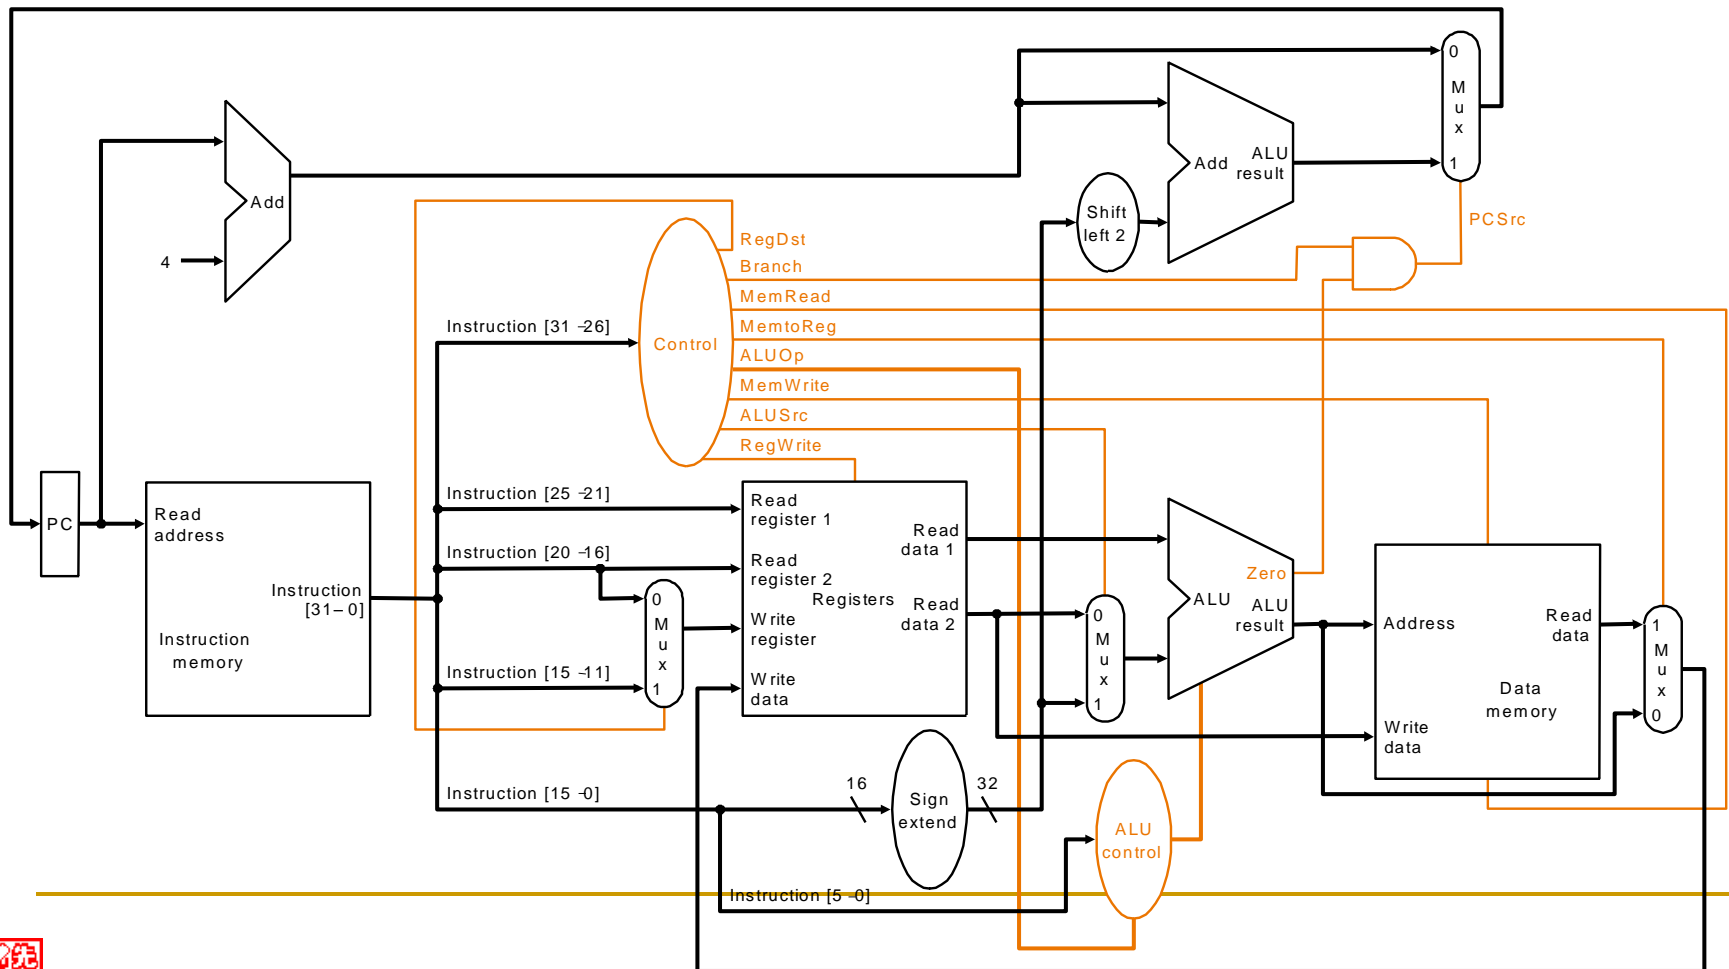


# Notes

- One significant difference in the execution of an R-type instruction between multicycle and pipelined implementations:
  - register write-back for the R-type instruction is the 5<sup>th</sup> (the last write-back) pipeline stage vs. the 4<sup>th</sup> stage for the multicycle implementation. *Why?*
  - think of *structural hazards* when writing to the register file...
- Worth repeating: the *essential difference* between the pipeline and multicycle implementations is the insertion of pipeline registers to *decouple the 5 stages*
- The CPI of an *ideal pipeline* (no stalls) is 1. *Why?*
- The RaVi Architecture Visualization Project of Dortmund U. has pipeline simulations – see link in our Additional Resources page
- As we develop control for the pipeline keep in mind that the text *does not consider jump* – should not be too hard to implement!



# Recall Single-Cycle Control – the Datapath





# Recall Single-Cycle - ALU

Control  
opcode  
ALU action

input

operation

LW	00	load word	xxxxxx	add	010
SW	00	store word	xxxxxx	add	010
Branch eq	01	branch eq	xxxxxx	subtract	110
R-type	10	add	100000	add	010
R-type	10	subtract	100010	subtract	110
R-type	10	AND	100100	and	000
R-type					001
R-type					less 111

ALUOp		Func field						Operation
ALUOp1	ALUOp0	F5	F4	F3	F2	F1	F0	
0	0	X	X	X	X	X	X	010
0	1	X	X	X	X	X	X	110
1	X	X	X	0	0	0	0	010
1	X	X	X	0	0	1	0	110
1	X	X	X	0	1	0	0	000
1	X	X	X	0	1	0	1	001
1	X	X	X	1	0	1	0	111

Truth table for ALU control bits





# Control Signals

Effect of control bits

Signal Name	Effect when deasserted	Effect when asserted
RegDst	The register destination number for the	The register
	destination number for the	
	Write register comes from the rd field (bits 15-11)	Write register
	from the rt field (bits 20-16)	
RegWrite	None	
	register on the Write register input is written	
ALUSrc	The second ALU operand comes from the	with the value on the Write data input
	the sign-extended,	The second ALU operand
	second register file output (Read data 2)	lower 16
	the instruction	
PCSrc	The PC is replaced by the output of the adder	The PC is replaced
	the output of the adder	
	that computes the value of PC + 4	that computes the
	target	

MemRead	None	Data memory contents designated								
address										
Determine	Instruction	RegDst	ALUSrc	Memto-Reg	RegWrite	MemRead	MemWrite	Branch	ALUOp1	ALUOp0
Miner	R-format	None	0	0	1	0	0	0	1	0
control	lw	0	1	1	1	1	0	0	0	0
bits data	sw	X	1	X	0	0	1	0	0	0
input	MemtoReg	The value fed to the register	0	0	0	0	1	1	The value fed to the	
08:51:24	register Write data input									



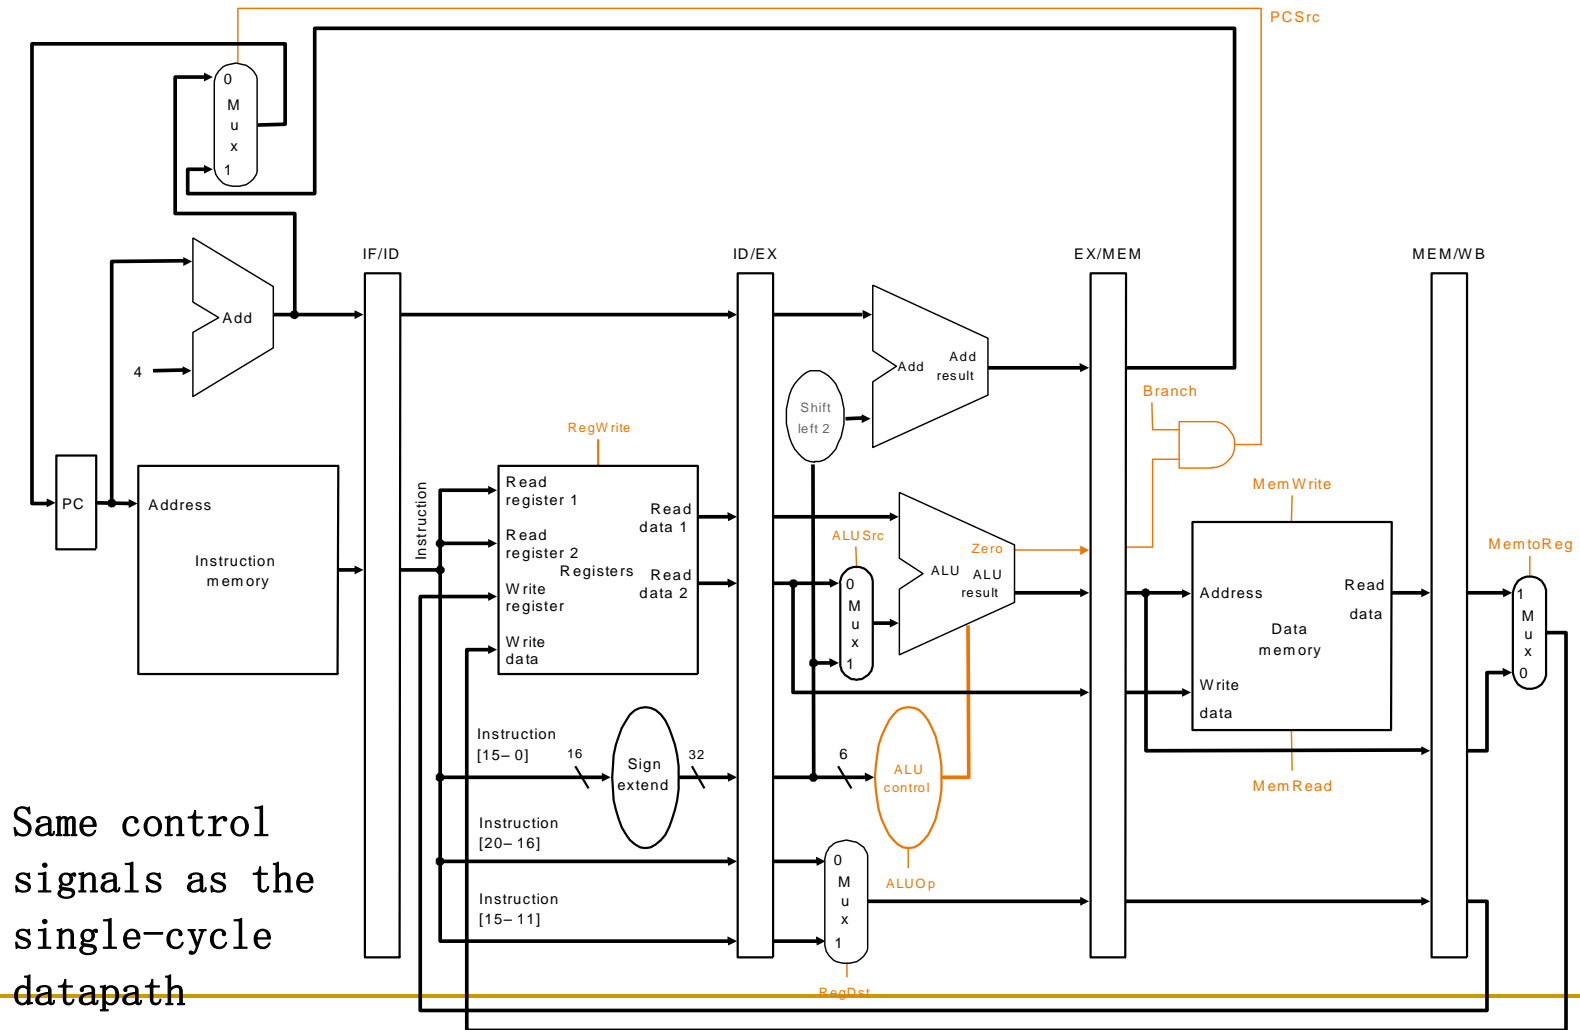
# Pipeline Control

- Initial design - *motivated by single-cycle datapath control* - use the same control signals
- Observe:
  - No separate write signal for the PC as it is written every cycle
  - No separate write signals for the pipeline registers as they are written every cycle
  - No separate read signal for instruction memory as it is read every clock cycle
  - No separate read signal for register file as it is read every clock cycle
- Need to set control signals during each pipeline stage
- Since control signals are associated with components active during a single pipeline stage, ~~can group control lines into five groups~~ according to pipeline stage

} Will be modified by hazard detection unit!!



# Pipelined Datapath with Control I



# Pipeline Control Signals

- There are five stages in the pipeline


- instruction fetch / PC increment*
- instruction decode / register fetch*
- execution / address calculation*
- memory access*
- write back*

Nothing to control as ins  
read and PC write are alw

Instruction	Execution/Address Calculation stage control lines				Memory access stage control lines			stage control lines	
	Reg Dst	ALU Op1	ALU Op0	ALU Src	Branch	Mem Read	Mem Write	Reg write	Mem to Reg
R-format	1	1	0	0	0	0	0	1	0
lw	0	0	0	1	0	1	0	1	1
sw	X	0	0	1	0	0	1	0	X
beq	X	0	1	0	1	0	0	0	X


# Pipeline Control Implementation

- Pass control signals along just like the data - extend each pipeline register to hold needed control bits for succeeding stages




- Note: The 6-bit *funct field* of the instruction required in the EX stage to generate ALU control can be retrieved as the 6 least significant bits of the immediate field which is sign-extended and passed from the IF/ID register to the ID/EX register

# Pipelined Datapath with Control II



Control signals  
emanate from  
the control  
portions of the  
pipeline registers



```
lw    $10, 20($1)
sub   $11, $2, $3
and   $12, $4, $7
or    $13, $6, $7
add   $14, $8, $9
```

**IF: lw \$10, 20(\$1)**

**ID: before<1>**

**EX: before<2>**

**MEM: before<3>**

**WB: before<4>**

**Clock cycle 1**

**Clock 1**

先生印

# Pipelined



# Execution

Clock cycle 3  
Clock 3

## Instruction sequence:

**Control**

```
lw $t0, 20($t1)
sub $t1, $t2, $t3
and $t2, $t4, $t7
or $t3, $t6, $t7
add $t4, $t8, $t9
```



Clock cycle 4  
Clock 4


# Pipelined Execution and Control

## Instruction sequence:

```
lw $10, 20($1)
sub $11, $2, $3
and $12, $4, $7
or $13, $6, $7
add $14, $8, $9
```

Label “after<i>” means  
i th instruction after add



Clock cycle 6



# Pipelined Execution

## Instruction sequence:

```
lw $t0, 20($t1)
sub $t1, $t2, $t3
and $t2, $t4, $t7
or $t3, $t6, $t7
add $t4, $t8, $t9
```






# Pipelined Execution and Control

## ■ Instruction sequence:

```
lw  $10, 20($1)
sub $11, $2, $3
and $12, $4, $7
or  $13, $6, $7
add $14, $8, $9
```






# Revisiting Hazards

- So far our datapath and control have ignored hazards
- We shall revisit *data hazards* and *control hazards* and enhance our datapath and control to handle them in *hardware*...

# Data Hazards and Forwarding

- Problem with starting an instruction before previous are finished:
  - data dependencies that go backward in time - called *data hazards*



# Software Solution

- Have compiler guarantee *never* any data hazards!
  - *by rearranging instructions to insert independent instructions between instructions that would otherwise have a data hazard between them,*
  - *or, if such rearrangement is not possible, insert*

sub	<i>nops</i>	\$2, \$1, \$3		sub	\$2, \$1, \$3
lw		\$10, 40 (\$3)		nop	
slt		\$5, \$6, \$7		nop	
and		\$12, \$2, \$5	or	and	\$12, \$2, \$5
or		\$13, \$6, \$2		or	\$13, \$6, \$2
add		\$14, \$2, \$2		add	\$14, \$2, \$2
sw		\$15, 100 (\$2)		sw	\$15, 100 (\$2)

- Such compiler solutions may not always be possible, and ~~nops slow the machine down~~

MIPS: nop = "no operation" = 00...0 (32b)


# Hardware Solution:

## Forwarding

Idea: use intermediate data, do not wait for result to be finally written to the destination register. Two steps:

1. *Detect data hazard*
2. *Forward intermediate data to resolve hazard*

# Pipelined Datapath with Control II (as before)



# Hazard Detection

## ■ Hazard conditions:

1a.  $EX/MEM.RegisterRd = ID/EX.RegisterRs$

1b.  $EX/MEM.RegisterRd = ID/EX.RegisterRt$

2a.  $MEM/WB.RegisterRd = ID/EX.RegisterRs$

2b.  $MEM/WB.RegisterRd = ID/EX.RegisterRt$

- Eg., in the earlier example, first hazard between sub \$2, \$1, \$3 and and \$12, \$2, \$5 is detected when the and is in EX stage and the sub is in MEM stage because

- $EX/MEM.RegisterRd = ID/EX.RegisterRs = \$2$  (1a)

## ■ Whether to forward also depends on:

- if the later instruction is going to write a register - if not, no need to forward, even if there is register number match as in conditions above
- if the destination register of the later instruction is \$0 - in which case there is no need to forward value (\$0 is always 0 and never overwritten)

# Data Forwarding

## ■ Plan:

- *allow inputs to the ALU not just from ID/EX, but also later pipeline registers, and*
- *use multiplexors and control signals to choose appropriate inputs to ALU*

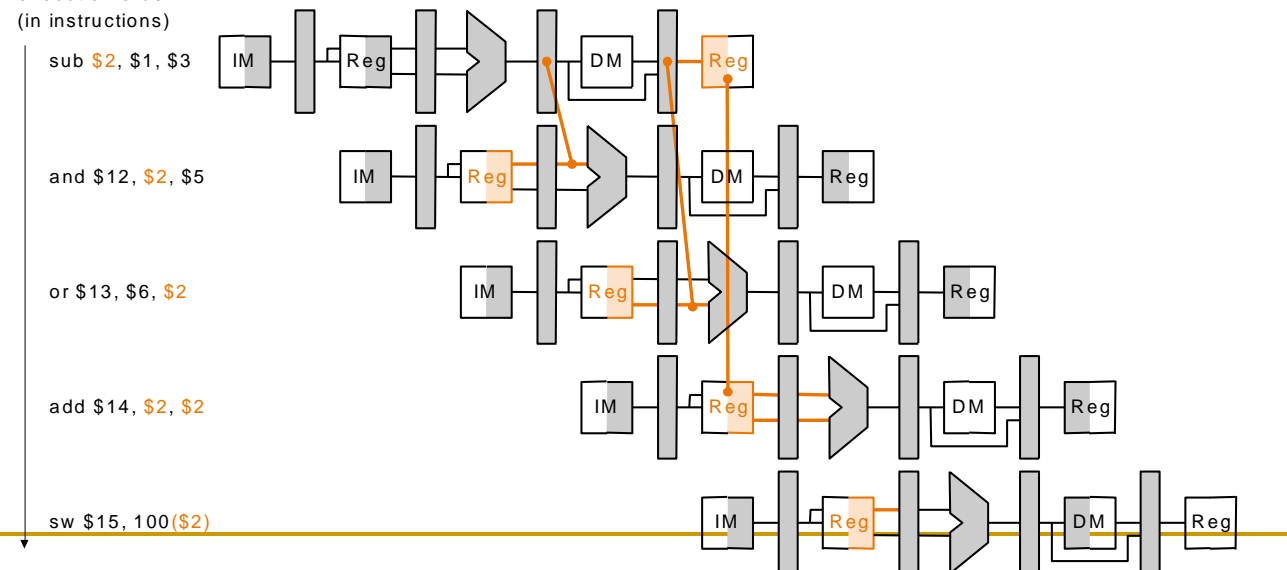
Time (in clock cycles) →

	CC 1	CC 2	CC 3	CC 4	CC 5	CC 6	CC 7	CC 8	CC 9
Value of register \$2 :	10	10	10	10	10/- 20	- 20	- 20	- 20	- 20
Value of EX/MEM :	X	X	X	- 20	X	X	X	X	X
Value of MEM/WB :	X	X	X	X	- 20	X	X	X	X

```

sub $2, $1, $3
and $12, $2, $5
or $13, $6, $2
add $14, $2, $2
sw $15, 100($2)
  
```

Program  
execution order  
(in instructions)

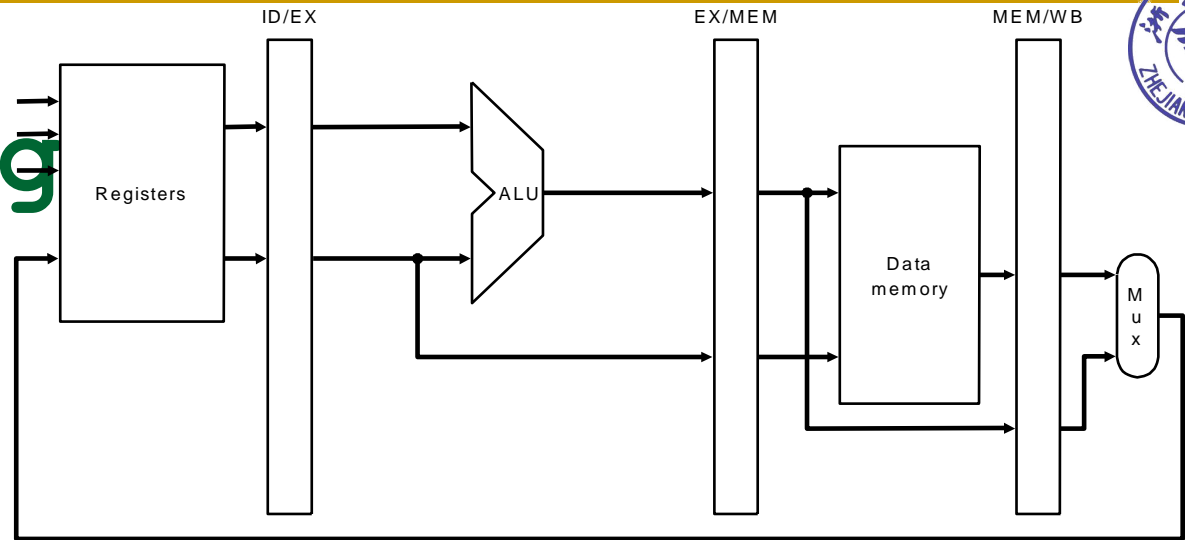


Dependencies between pipelines move forward in time

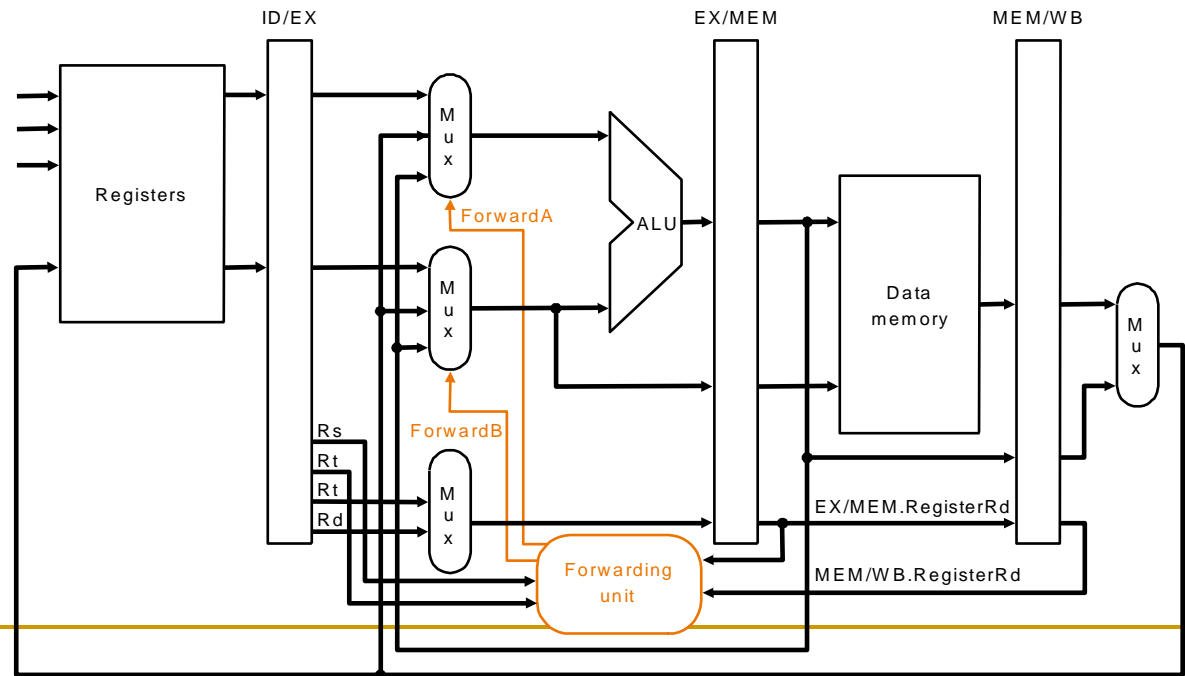




# Forwarding Hardware



a. No forwarding Datapath before adding forwarding hardware



b. With forwarding Datapath after adding forwarding hardware



# Forwarding Hardware: Multiplexor Control

Mux control	Source	Explanation
ForwardA = 00 the register file	ID/EX	The first ALU operand comes from
ForwardA = 10 from prior ALU result	EX/MEM	The first ALU operand is forwarded
ForwardA = 01 from data memory	MEM/WB	The first ALU operand is forwarded or an earlier ALU
result		
ForwardB = 00 the register file	ID/EX	The second ALU operand comes from
ForwardB = 10 from prior ALU result	EX/MEM	The second ALU operand is forwarded
ForwardB = 01 from data memory	MEM/WB	Depending on the ALU operand is forwarded most r (see datapath with control diagram) or an earlier ALU

# Data Hazard: Detection and Forwarding

- Forwarding unit determines multiplexor control according to the following rules:

```
1.  EX hazard
    if (          EX/MEM.RegWrite
// if there is a write...
        and ( EX/MEM.RegisterRd  $\neq$  0 )
// to a non-$0 register...
        and ( EX/MEM.RegisterRd = ID/EX.RegisterRs ) )
// which matches, then...    ForwardA = 10

    if (          EX/MEM.RegWrite
// if there is a write...
        and ( EX/MEM.RegisterRd  $\neq$  0 )
// to a non-$0 register...
        and ( EX/MEM.RegisterRd = ID/EX.RegisterRt ) )
// which matches, then... ForwardB = 10
```



# Data Hazard: Detection and Forwarding

2.

MEM hazard

```

    if (          MEM/WB.RegWrite
// if there is a write...
        and ( MEM/WB.RegisterRd  $\neq$  0 )
// to a non-$0 register...
        and ( EX/MEM.RegisterRd  $\neq$  ID/EX.RegisterRs )      // and
not already a register match

// with earlier pipeline register...
        and ( MEM/WB.RegisterRd = ID/EX.RegisterRs ) ) // but
match with later pipeline

```

register, then...

ForwardA = 01

```

    if (          MEM/WB.RegWrite
// if there is a write...
        and ( MEM/WB.RegisterRd  $\neq$  0 )
// to a non-$0 register...
        and ( EX/MEM.RegisterRd  $\neq$  ID/EX.RegisterRt )      // and
not already a register match earlier pipeline (EX/MEM) register has

```

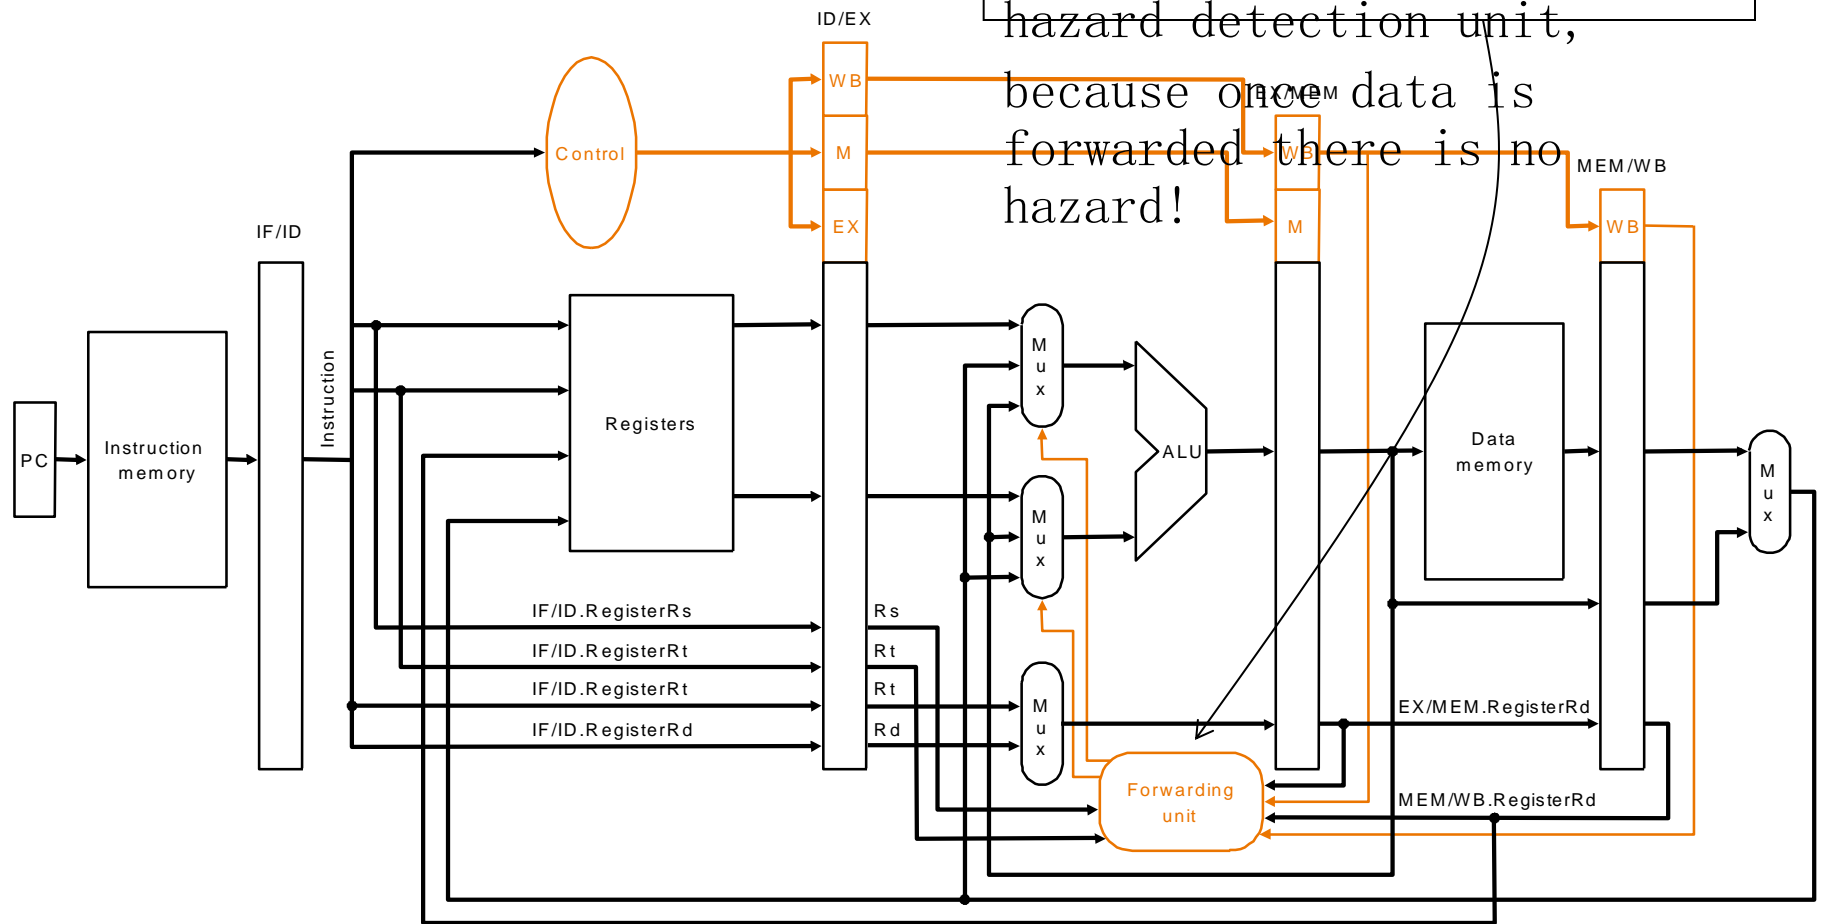
This check is necessary, e.g., for sequences such as add \$1, \$1, \$2;



# Forwarding Hardware with Control

Called forwarding unit, not  
hazard detection unit,

because once data is  
forwarded there is no  
hazard!



Datapath with forwarding hardware and control wires - certain details, e.g., branching hardware, are omitted to simplify the drawing

Note: so far we have only handled forwarding to R-type instructions...

# Forwarding



Clock cycle 3


## Execution

example:

```
sub $2, $1, $3
and $4, $2, $5
or $4, $4, $2
add $9, $4, $2
```

Clock cycle 4


Clock 4



08:51:25

Clock cycle 5

sub \$2, \$1, \$3  
and \$4, \$2, \$5  
or \$4, \$4, \$2  
add \$9, \$4, \$2



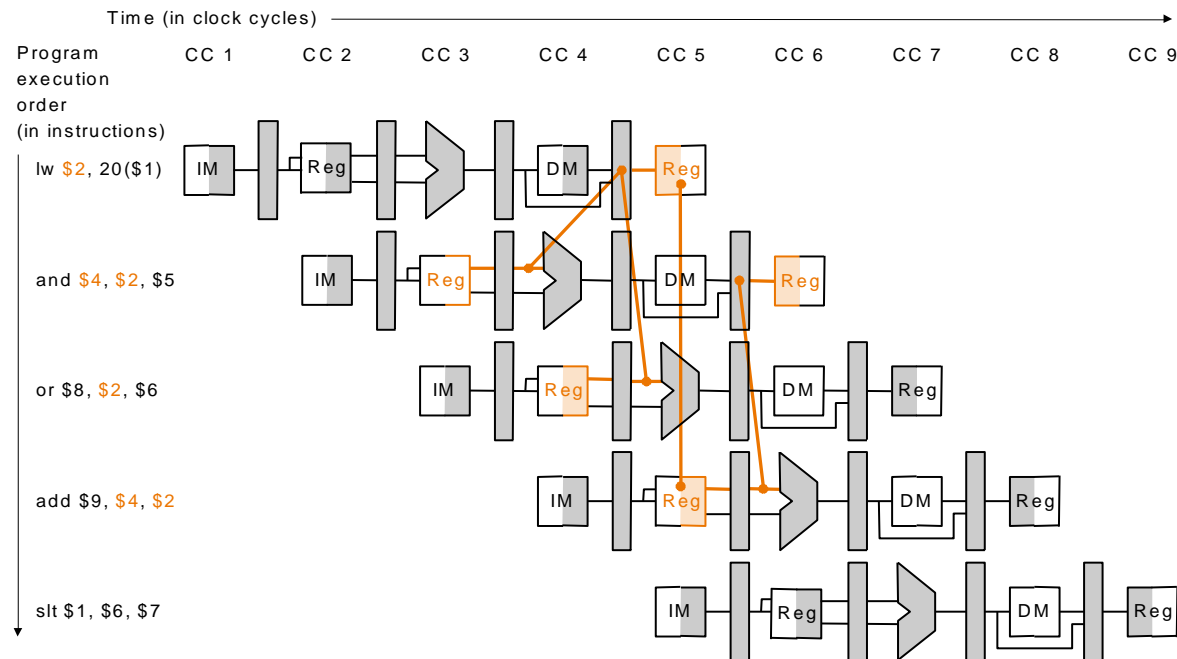
# Data Hazards and Stalls



- Load word can still cause a hazard:
  - an instruction tries to read a register following a load instruction that writes to the same register

```
lw $2, 20($1)
and $4, $2, $5
or $8, $2, $6
add $9, $4, $2
slt $1, $6, $7
```

As even a pipeline dependency goes backward in time forwarding will not solve the hazard

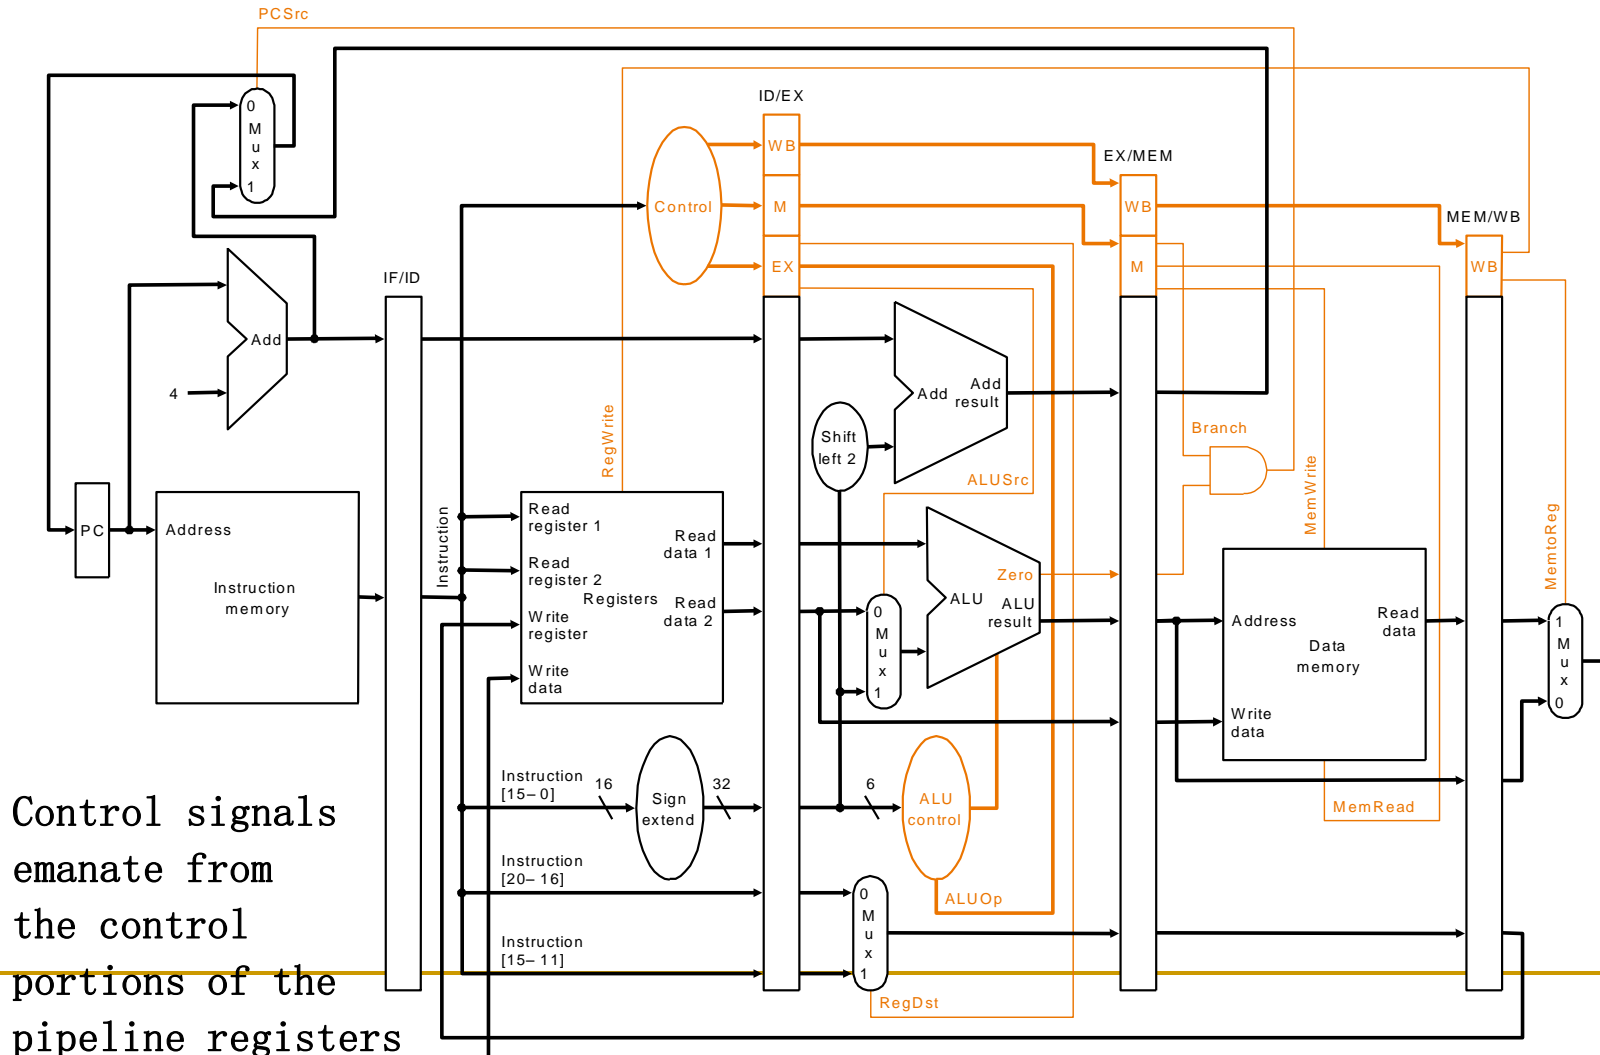


- therefore, we need a hazard detection unit to stall the pipeline after the load instruction





# Pipelined Datapath with Control II (as before)





# Hazard Detection Logic to Stall

- Hazard detection unit implements the following check if to stall


```
if ( ID/EX.MemRead                                // if the
    instruction in the EX stage is a load...
    and ( ( ID/EX.RegisterRt = IF/ID.RegisterRs )      // and
        the destination register
        or ( ID/EX.RegisterRt = IF/ID.RegisterRt ) ) ) //
    matches either source register

//

of the instruction in the ID stage, then...
stall the pipeline
```

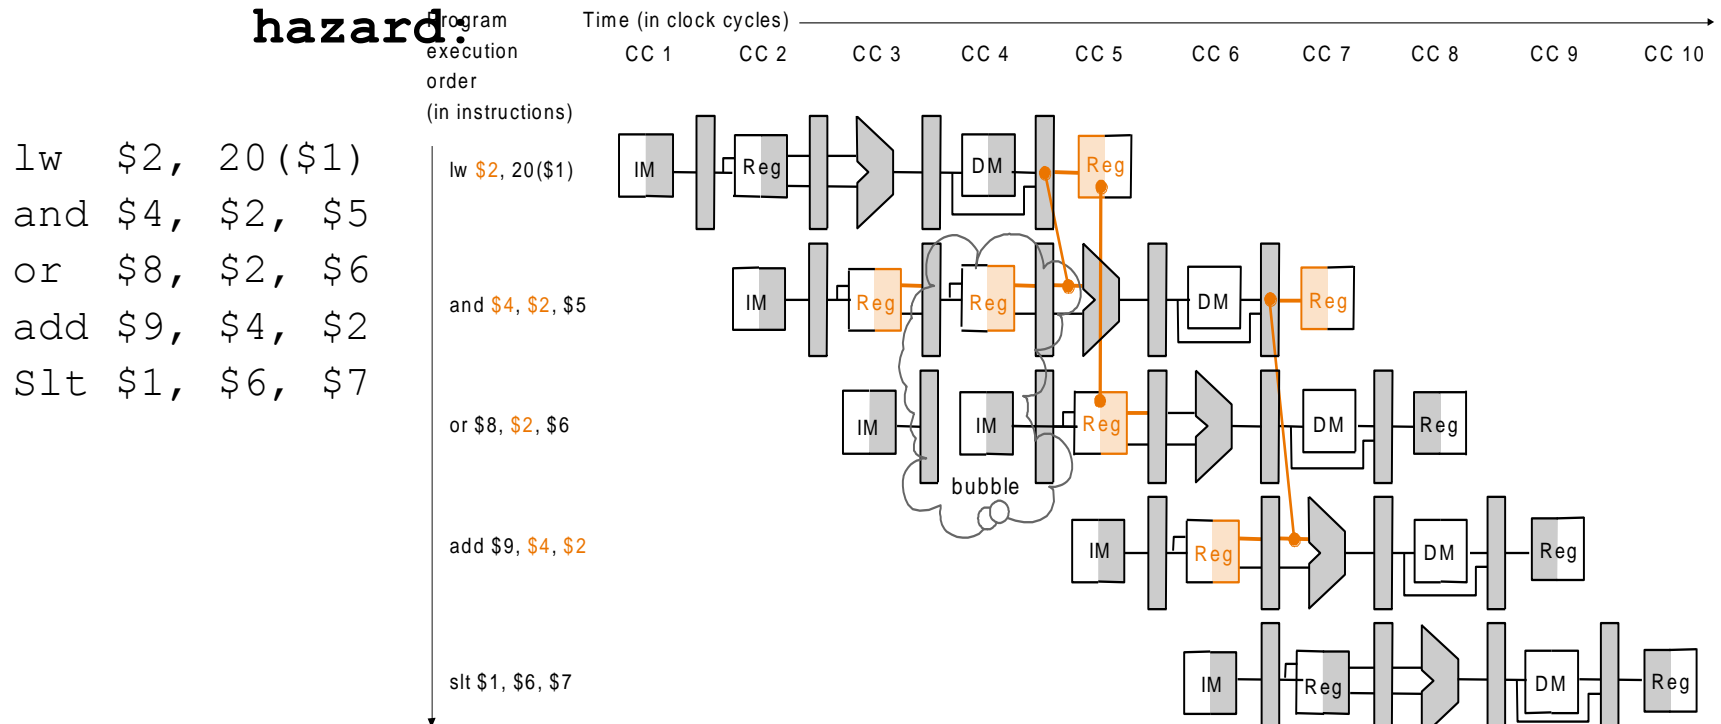
# Mechanics of Stalling

- If the check to stall verifies, then the *pipeline needs to stall only 1 clock cycle* after the load as after that the forwarding unit can resolve the dependency
- What the hardware does to stall the pipeline 1 cycle:
  - *does not let the IF/ID register change (disable write!) – this will cause the instruction in the ID stage to repeat, i.e., stall*
  - *therefore, the instruction, just behind, in the IF stage must be stalled as well – so hardware does not let the PC change (disable write!) – this will cause the instruction in the IF stage to repeat, i.e., stall*
  - *changes all the EX, MEM and WB control fields in the ID/EX pipeline register to 0, so effectively the instruction just behind the load becomes a nop – a bubble is said to have been inserted into the pipeline*



# Stalling Resolves a Hazard

- Same instruction sequence as before for which forwarding by itself could not resolve the hazard:



Hazard detection unit inserts a 1-cycle bubble in the pipeline, after which all pipeline register dependencies go forward so then the forwarding unit can handle them and there are no more hazards




# Stalling


## Execution example:

```
lw $2, 20($1)
and $4, $2, $5
or $4, $4, $2
add $9, $4, $2
```

Clock cycle 2  
Clock 2



Clock cycle 3  
Clock 3




# Stalling

## Execution example (cont.):

```
lw $2, 20($1)
and $4, $2, $5
or $4, $4, $2
add $9, $4, $2
```

Clock cycle 5

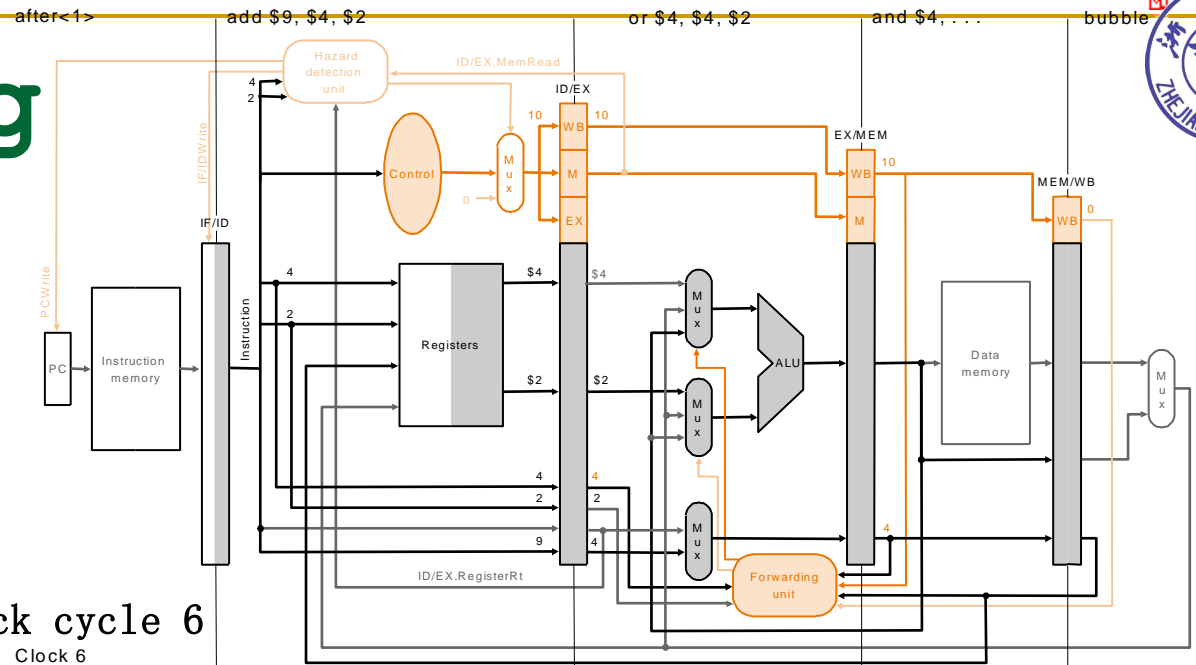


# Stalling


## Execution example (cont.):

```
lw $2, 20($1)
and $4, $2, $5
or $4, $4, $2
add $9, $4, $2
```

Clock cycle 6  
Clock 6



Clock cycle 7  
Clock 7





# Control (or Branch) Hazards




- Problem with branches in the pipeline we have so far is that the *branch decision is not made till the MEM stage* – so what instructions, if at all, should we insert into the pipeline following the branch instructions?
- Possible solution: *stall* the pipeline till branch decision is known
  - *not efficient, slow the pipeline significantly!*
- Another solution: *predict* the branch outcome
  - *e.g., always predict branch-not-taken – continue with next sequential instructions*
  - *if the prediction is wrong have to flush the pipeline behind the branch – discard instructions already fetched or decoded – and continue execution at the branch target*



# Predicting Branch-not-taken:

## Misprediction delay



The outcome of branch taken (prediction wrong) is decided only when **beq** is in the MEM stage, so the following three sequential instructions already in the pipeline have to be flushed and execution resumes at **lw**

# Optimizing the Pipeline to Reduce Branch Delay

- *Move the branch decision from the MEM stage (as in our current pipeline) earlier to the ID stage*
  - *calculating the branch target address involves moving the branch adder from the MEM stage to the ID stage - inputs to this adder, the PC value and the immediate fields are already available in the IF/ID pipeline register*
  - *calculating the branch decision is efficiently done, e.g., for equality test, by XORing respective bits and then ORing all the results and inverting, rather than using the ALU to subtract and then test for zero (when there is a carry delay)*
- ~~with the more efficient equality test we can put it in the ID stage without significantly lengthening this stage - remember an objective of pipeline~~

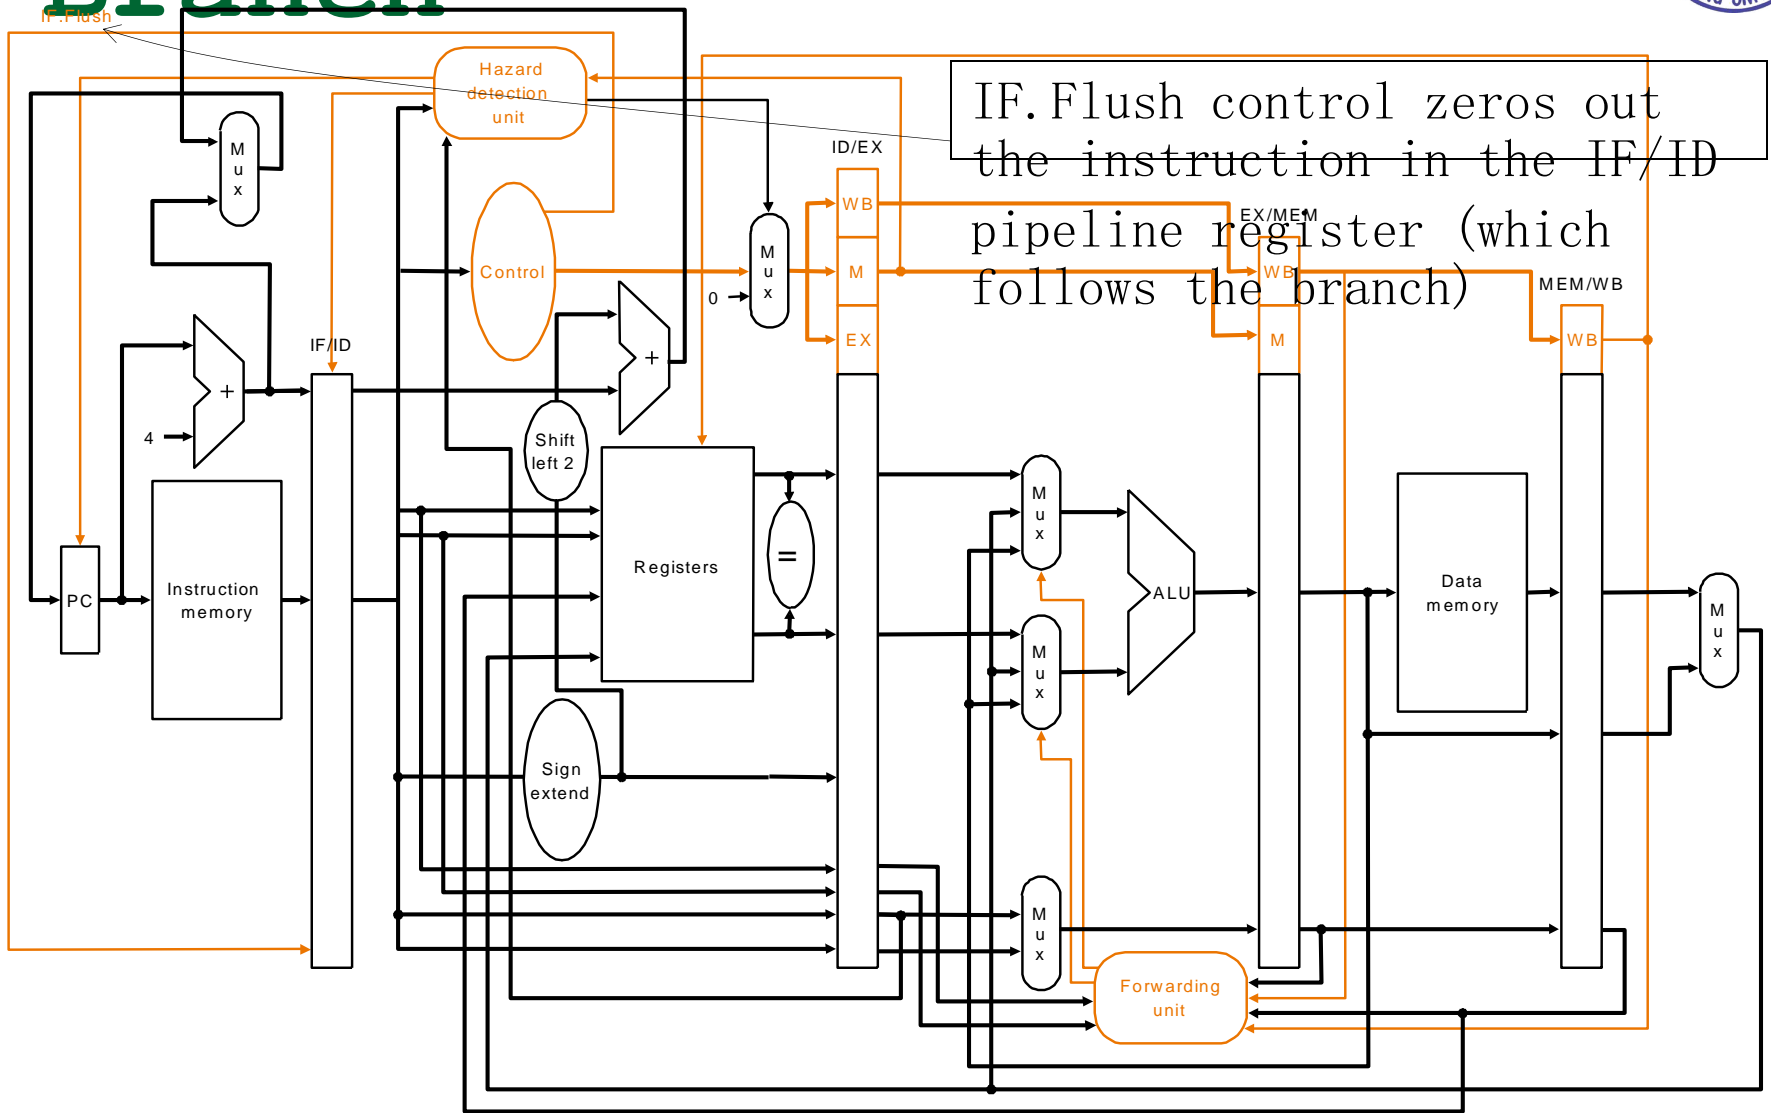
# Flushing on Misprediction



- Same strategy as for stalling on load-use data hazard...
- Zero out all the control values (or the instruction itself) in pipeline registers for the instructions following the branch that are already in the pipeline – effectively turning them into nops – so they are flushed
  - in the optimized pipeline, with branch decision made in the ID stage, we have to flush only one instruction in the IF stage – the branch delay penalty is then only one clock cycle



## Branch



Branch decision is moved from the MEM stage to the ID stage - simplified drawing not showing enhancements to the forwarding and hazard detection units

# Pipelined Branch

## Execution example:

```

36 sub $10, $4, $8
40 beq $1, $3, 7
44 and $12, $2, $5
48 or $13, $2, $6
52 add $14, $4, $2
56 slt $15, $6, $7


```

...


```
72 lw $4, 50($7)
```

Optimized pipeline with only one bubble as a result of the taken branch

Clock cycle 3  
Clock 3



Clock cycle 4  
Clock 4



# Simple Example: Comparing Performance



- *Compare performance for single-cycle, multicycle, and pipelined datapaths using the gcc instruction mix*
  - assume 2 ns for memory access, 2 ns for ALU operation, 1 ns for register read or write
  - assume gcc instruction mix 23% loads, 13% stores, 19% branches, 2% jumps, 43% ALU
  - for pipelined execution assume
    - 50% of the loads are followed immediately by an instruction that uses the result of the load
    - 25% of branches are mispredicted
    - branch delay on misprediction is 1 clock cycle
    - jumps always incur 1 clock cycle delay so their average time is 2 clock cycles



# Simple Example: Comparing Performance

- *Single-cycle* (p. 373): average instruction time 8 ns
- *Multicycle* (p. 397): average instruction time 8.04 ns
- *Pipelined*:
  - loads use 1 cc (clock cycle) when no load-use dependency and 2 cc when there is dependency – given 50% of loads are followed by dependency the average cc per load is 1.5
  - stores use 1 cc each
  - branches use 1 cc when predicted correctly and 2 cc when not – given 25% misprediction average cc per branch is 1.25
  - jumps use 2 cc each
  - ~~ALU instructions use 1 cc each~~
  - therefore, average CPI is

$$1.5 \times 23\% + 1 \times 13\% + 1.25 \times 19\% + 2 \times 2\% + 1 \times 43\% =$$







# Thank You!

非常天空

浙江大学计算机学院

